

Operating System Design and Implementation

Getting started with kernel and kernel debugging

Charles Tsao

Outline

- How to develop an operating system
 - hardware, compile, assembler, linker
- How to compile kernel codes
 - make
- How to debug kernel codes
 - gdb, kgdb, ICE
- How to maintain kernel codes
 - Version control, CVS, SVN, GIT
 - patch
- A case study of Linux

Hardware

- How to develop an operating system for a new processor
 - Simulator vs. emulator vs. virtual machine
- Simulator
 - A program to reproduce the behavior of a computer system based on an abstract model
- Emulator
 - Hardware or software or both that duplicates (or emulates) the functions of one computer system (the guest) in another computer system (the host), different from the first one, so that the emulated behavior closely resembles the behavior of the real system (the guest)¹
- Virtual machine
 - A virtual machine (VM) is a software implementation of a machine (i.e. a computer) that executes programs like a physical machine²

^{1,2} <http://en.wikipedia.org/>

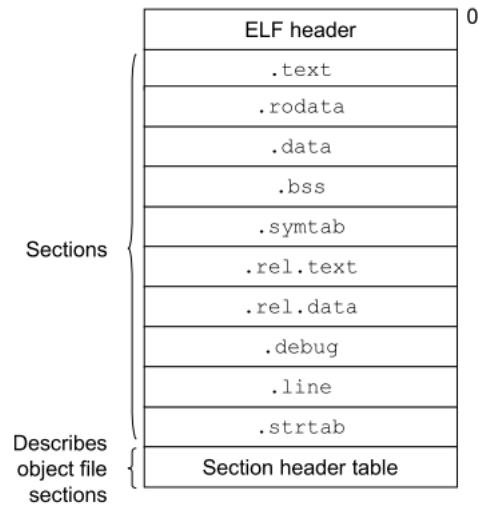
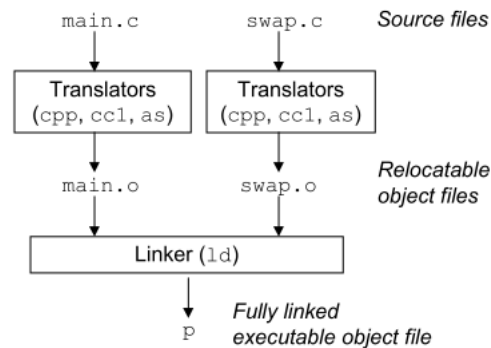
Compiler

- Compiler vs. cross compiler
 - How to develop a compiler for new processor?
- Assembler
 - How to develop an assembler?
- Linker
 - Why?

Linking

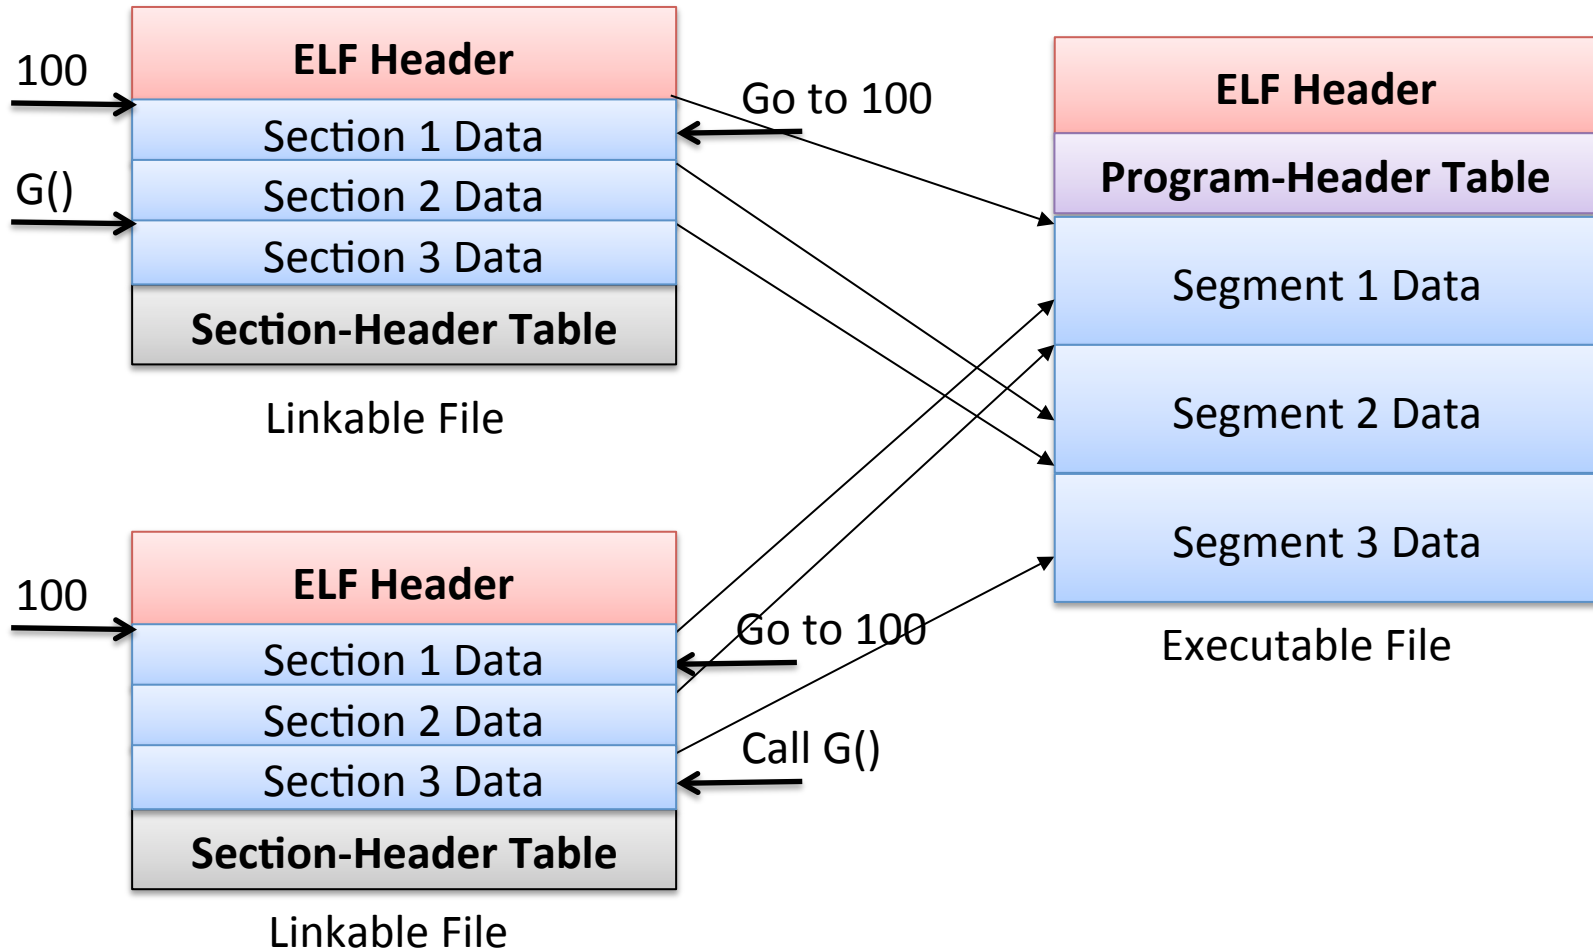
<hr/> <div><i>code/link/main.c</i></div> <pre>1 /* main.c */ 2 void swap(); 3 4 int buf[2] = {1, 2}; 5 6 int main() 7 { 8 swap(); 9 return 0; 10 }</pre> <hr/> <div><i>code/link/main.c</i></div>	<hr/> <div><i>code/link/swap.c</i></div> <pre>1 /* swap.c */ 2 extern int buf[]; 3 4 int *bufp0 = &buf[0]; 5 int *bufp1; 6 7 void swap() 8 { 9 int temp; 10 11 bufp1 = &buf[1]; 12 temp = *bufp0; 13 *bufp0 = *bufp1; 14 *bufp1 = temp; 15 }</pre> <hr/> <div><i>code/link/swap.c</i></div>
(a) main.c	(b) swap.c

Linking



Num:	Value	Size	Type	Bind	Ot	Ndx	Name
8:	0	8	OBJECT	GLOBAL	0	3	buf
9:	0	17	FUNC	GLOBAL	0	1	main
10:	0	0	NOTYPE	GLOBAL	0	UND	swap

Linking

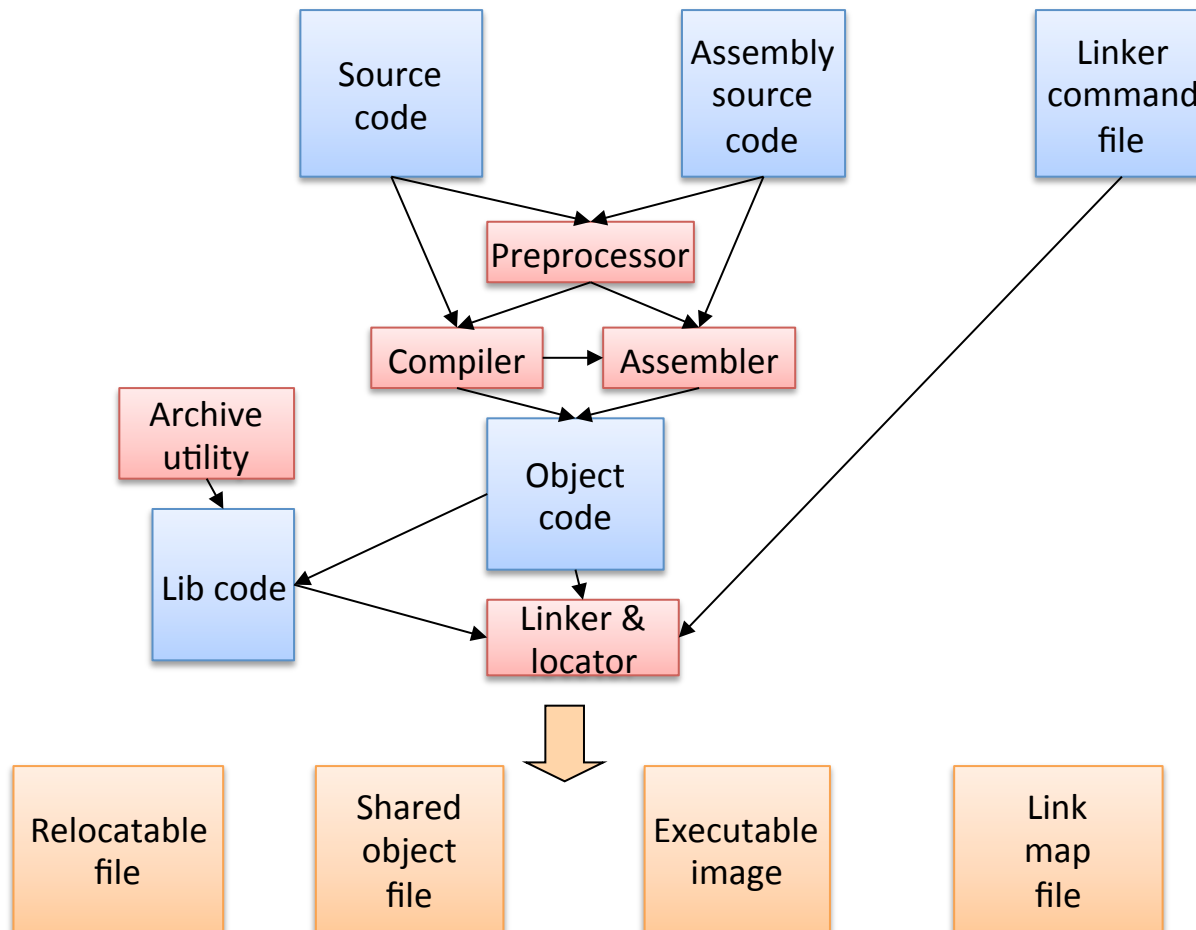


Linking

- Please review system programming and compiler if you are not familiar with below terms
 - Static linking
 - Dynamic linking
 - Relocations
 - Symbol table
 - Share library
 - Linking and loading

How to compile kernel codes

- Creating an executable image



How to compile kernel codes

- 17,090 – The number of files in Linux 2.6.11
- 37,626 – The number of files in Linux 3.2
- How can I find and compile my network interface card driver among 100 network interface card drivers?
- Shall I recompile again when I modify one file?
- How can we produce kernel image ?

Make and Makefile

- Make: utility to provide a convenient facility to build, install, and uninstall projects
- Makefile: script file for make to compile and link programs

Make and Makefile

```
target ... : prerequisites ...  
    recipe  
    ...  
    ...
```

```
edit : main.o kbd.o command.o display.o \  
      insert.o search.o files.o utils.o  
      cc -o edit main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o  
  
main.o : main.c defs.h  
      cc -c main.c  
kbd.o : kbd.c defs.h command.h  
      cc -c kbd.c  
command.o : command.c defs.h command.h  
      cc -c command.c  
display.o : display.c defs.h buffer.h  
      cc -c display.c  
insert.o : insert.c defs.h buffer.h  
      cc -c insert.c  
search.o : search.c defs.h buffer.h  
      cc -c search.c  
files.o : files.c defs.h buffer.h command.h  
      cc -c files.c  
utils.o : utils.c defs.h  
      cc -c utils.c  
clean :  
      rm edit main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o
```

Make and Makefile

```
objects = main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o  
  
edit : $(objects)  
      cc -o edit $(objects)  
main.o : main.c defs.h  
      cc -c main.c  
kbd.o : kbd.c defs.h command.h  
      cc -c kbd.c  
command.o : command.c defs.h command.h  
      cc -c command.c  
display.o : display.c defs.h buffer.h  
      cc -c display.c  
insert.o : insert.c defs.h buffer.h  
      cc -c insert.c  
search.o : search.c defs.h buffer.h  
      cc -c search.c  
files.o : files.c defs.h buffer.h command.h  
      cc -c files.c  
utils.o : utils.c defs.h  
      cc -c utils.c  
clean :  
      rm edit $(objects)
```

```
objects = main.o kbd.o command.o display.o \  
          insert.o search.o files.o utils.o  
  
edit : $(objects)  
      cc -o edit $(objects)  
  
main.o : defs.h  
kbd.o : defs.h command.h  
command.o : defs.h command.h  
display.o : defs.h buffer.h  
insert.o : defs.h buffer.h  
search.o : defs.h buffer.h  
files.o : defs.h buffer.h command.h  
utils.o : defs.h  
  
.PHONY : clean  
clean :  
      rm edit $(objects)
```

Make and Makefile

- Variables and settings
 - make config
 - make menuconfig
- Phony targets
 - make all
 - make clean
 - make depend
 - make install
 - make uninstall

The portability problem

- Hardware differences
- OS differences
- Compiler differences

```
your source files --> [autoscan*] --> [configure.scan] --> configure.ac
```

```
configure.ac --.  
              | .-----> autoconf* -----> configure  
[aclocal.m4] --+---+  
              | '-----> [autoheader*] --> [config.h.in]  
[acsite.m4] ---'
```

```
Makefile.in
```

```
[acinclude.m4] --.  
              |  
[local macros] --+---> aclocal* --> aclocal.m4  
              |  
configure.ac ----'
```

The portability problem

```
configure.ac --.  
               +--> automake* --> Makefile.in  
Makefile.am ---'  
  
               .-----> [config.cache]  
configure* -----+-----> config.log  
                  |  
[config.h.in] -.      v      .-> [config.h] -.  
               +--> config.status* -+      +--> make*  
Makefile.in ---'               '-> Makefile ---'
```


How to debug kernel codes

- Kernel logs (discontinues logs)
 - Printk
 - Oops and Kallsyms
- Kernel debug supports
 - Kexec, kdump, SysRq
- Kernel hacking options
- Kernel debug tools
 - gdb, kgdb, kdb
- Profile
 - OProfile
- Trace
 - KFT, LTT/LTTng
 - Gprof
- Lock detection
 - Lockmeter
- Memory leaking
- Test equipment

Debugging and profiling device drivers

- `printk()`
 - Loglevels

Loglevel	Description
KERN_EMERG	An emergency condition; the system is probably dead
KERN_ALERT	A problem that requires immediate attention
KERN_CRIT	A critical condition
KERN_ERR	An error
KERN_WARNING	A warning
KERN_NOTICE	A normal, but perhaps noteworthy, condition
KERN_INFO	An informational message
KERN_DEBUG	A debug message typically superfluous

printk()

- Log buffer
- Klogd
 - /proc/kmsg or syslog()
- Syslogd
 - Appends all the messages it receives to a file
 - /var/log/messages
 - /etc/syslog.conf
- Not for booting stage debugger (early_printk())
- Not for debugging non-console case (via serial port)
- Not easy to detect racing condition

Kernel debugging options

- Turn on in kernel hacking/linuxconfig
- BUG()/ BUG_ON() case oops (stack trace, error message dump to kernel)
- Panic() prints error messages and halts the kernel
- dump_stack()

ksymoops

```
NIP: C013A7F0 LR: C013A7F0 SP: C0685E00 REGS: c0905d10 TRAP: 0700
Not tainted
MSR: 00089037 EE: 1 PR: 0 FP: 0 ME: 1 IR/DR: 11
TASK = c0712530[0] 'swapper' Last syscall: 120
GPR00: C013A7C0 C0295E00 C0231530 0000002F 00000001 C0380CB8 C0291B80 C02D0000
GPR08: 000012A0 00000000 00000000 C0292AA0 4020A088 00000000 00000000 00000000
GPR16: 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000
GPR24: 00000000 00000005 00000000 00001032 C3F7C000 00000032 FFFFFFFF C3F7C1C0
Call trace: [c013ab30] [c0020744] [c001b864] [c0007e80] [c00061c4]
[c0007b84] [c0007bf8] [c0003ae8]
```

- Ksymoops + system.map+module information
- Linux 2.6.X uses kallsyms

ksymoops

Oops: Exception in kernel mode, sig: 4
Unable to handle kernel NULL pointer dereference at virtual address 00000001

NIP: C013A7F0 LR: C013A7F0 SP: C0685E00 REGS: c0905d10 TRAP: 0700

Not tainted

MSR: 00089037 EE: 1 PR: 0 FP: 0 ME: 1 IR/DR: 11

TASK = c0712530[0] 'swapper' Last syscall: 120

GPR00: C013A7C0 C0295E00 C0231530 0000002F 00000001 C0380CB8 C0291B80 C02D0000

GPR08: 000012A0 00000000 00000000 C0292AA0 4020A088 00000000 00000000 00000000

GPR16: 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000

GPR24: 00000000 00000005 00000000 00001032 C3F7C000 00000032 FFFFFFFF C3F7C1C0

Call trace:

[c013ab30] tulip_timer+0x128/0x1c4

[c0020744] run_timer_softirq+0x10c/0x164

[c001b864] do_softirq+0x88/0x104

[c0007e80] timer_interrupt+0x284/0x298

[c00033c4] ret_from_except+0x0/0x34

[c0007b84] default_idle+0x20/0x60

[c0007bf8] cpu_idle+0x34/0x38

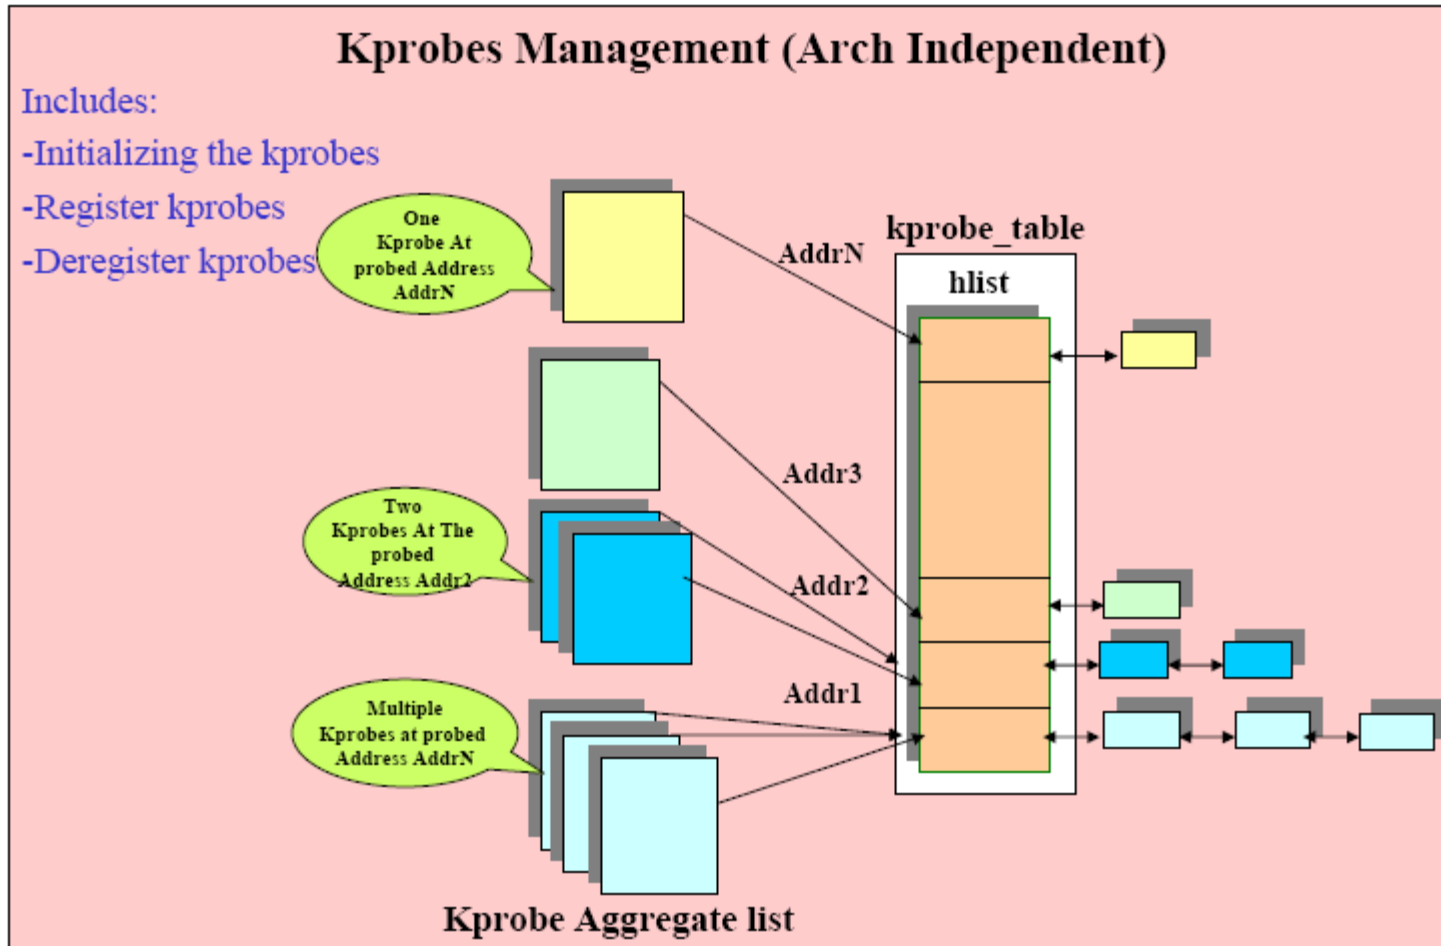
[c0003ae8] rest_init+0x24/0x34

SysRq

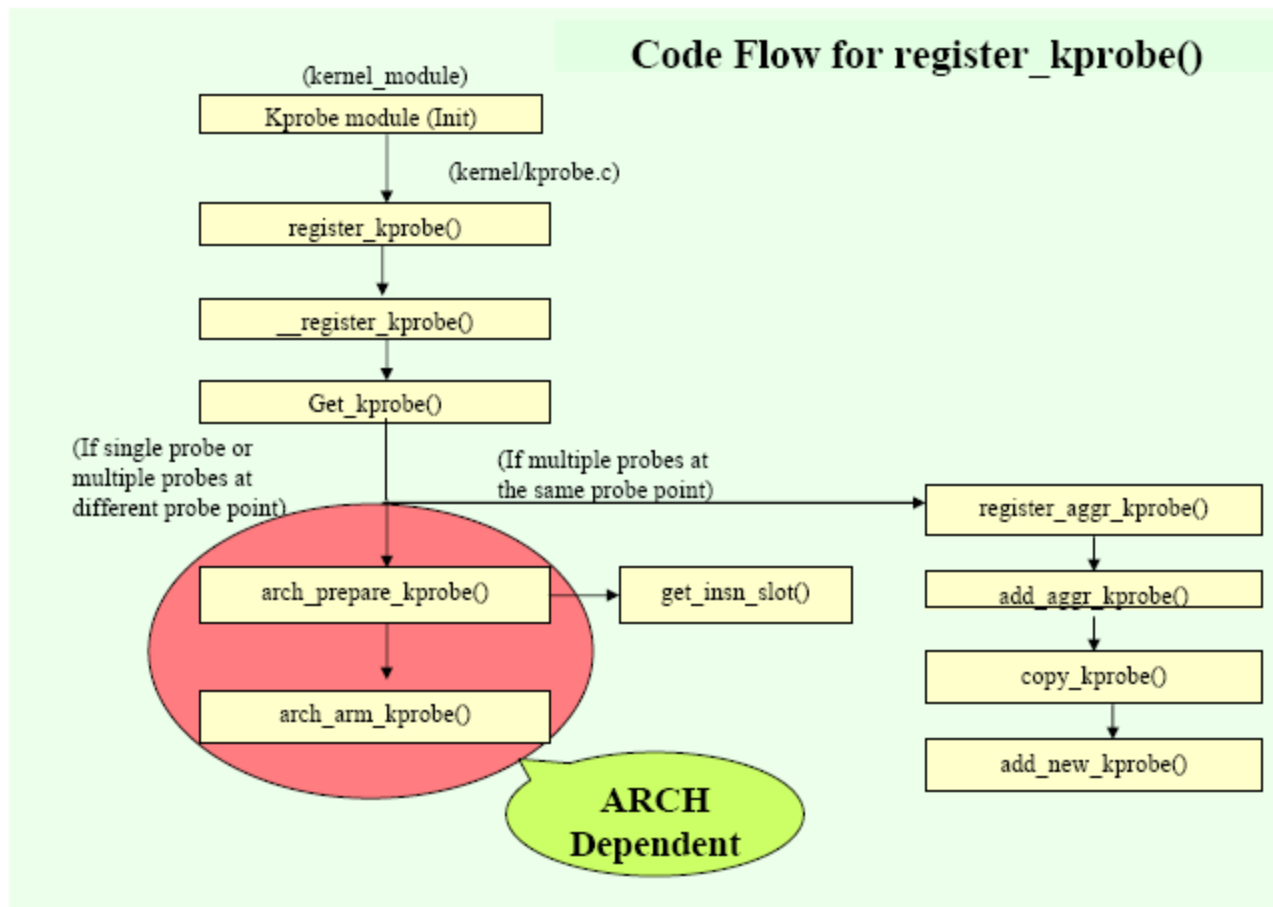
Key Command	Description
SysRq-b	Reboot the machine
SysRq-e	Send a SIGTERM to all processes except init
SysRq-h	Display SysRq help on the console
SysRq-i	Send a SIGKILL to all processes except init
SysRq-k	Secure Access Key: kill all programs on this console
SysRq-l	Send a SIGKILL to all processes including init
SysRq-m	Dump memory information to console
SysRq-o	Shut down the machine
SysRq-p	Dump registers to console
SysRq-r	Turn off keyboard raw mode
SysRq-s	Sync all mounted file systems to disk
SysRq-t	Dump task information to console
SysRq-u	Unmount all mounted file systems

Kprobes

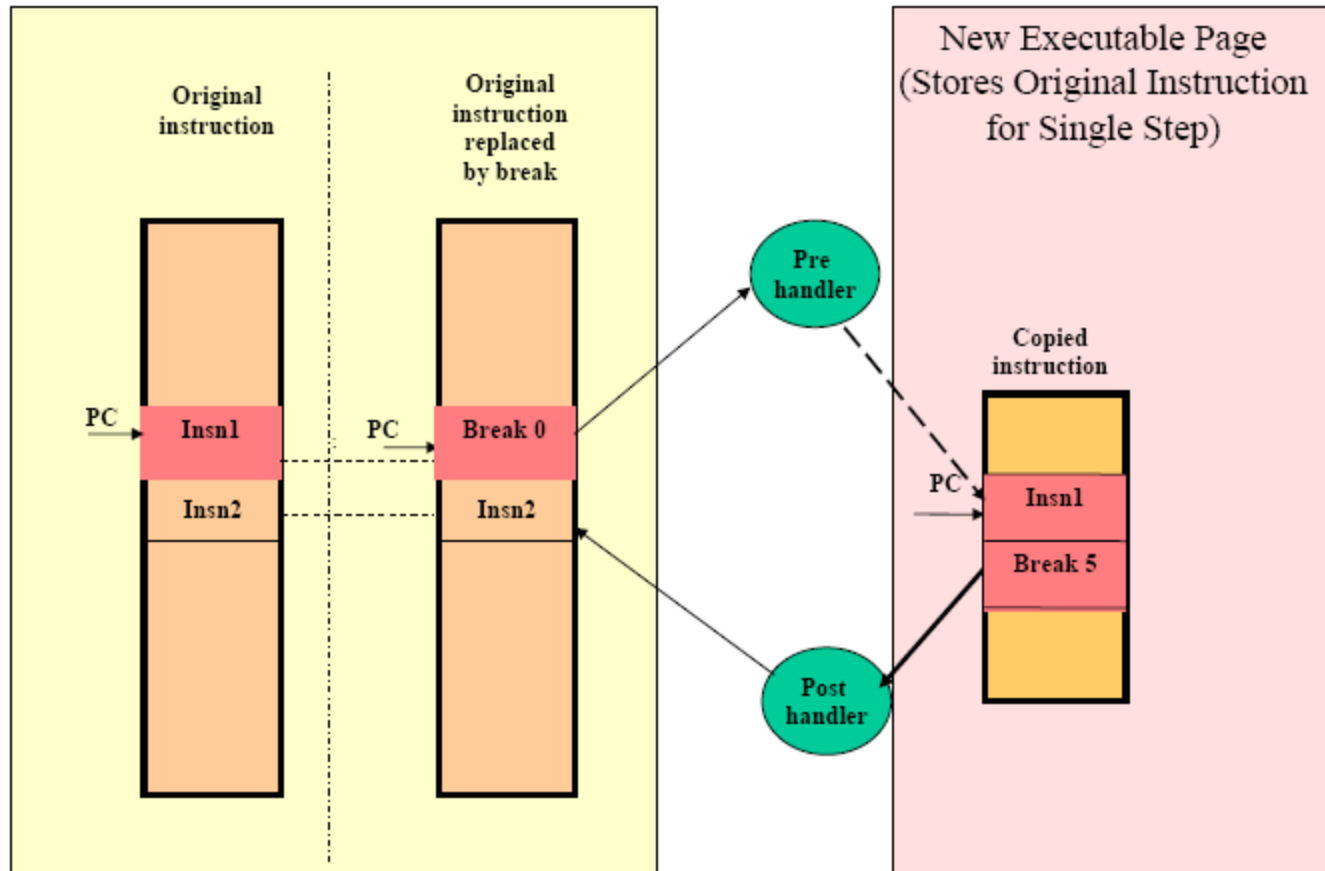
- Turn on CONFIG_KPROBES (Instrumentation Support Kprobes) in the kernel configuration menu



Kprobes



Kprobes




Kexec and Kdump

- Kexec uses `exec()` to spawn a new kernel over a running kernel without the overhead of boot
- save several seconds of reboot time
- capturing a dump after a kernel crash
- `CONFIG_KEXEC` (Processor Type and Features Kexec System Call) in the kernel configuration menu

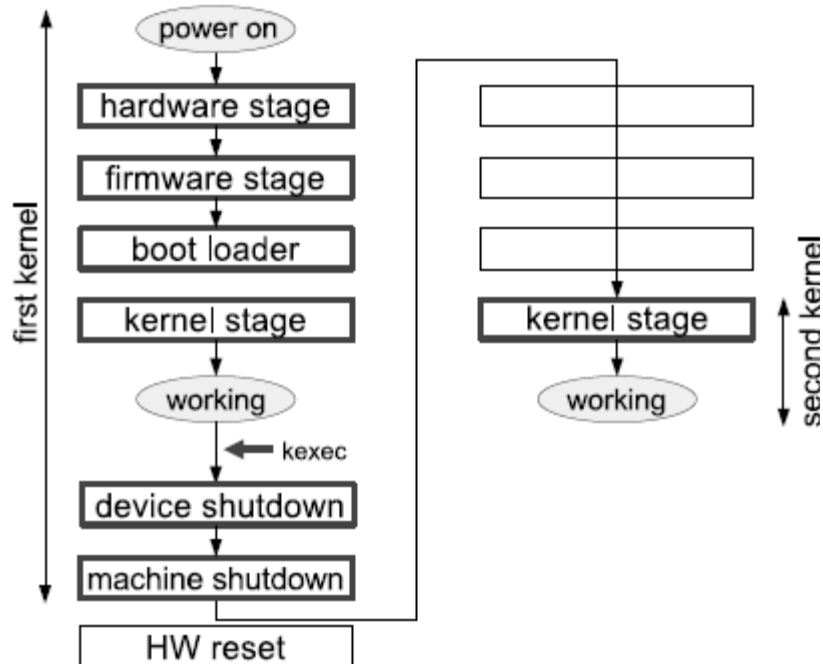
Kexec and Kdump

- Normal Linux boot

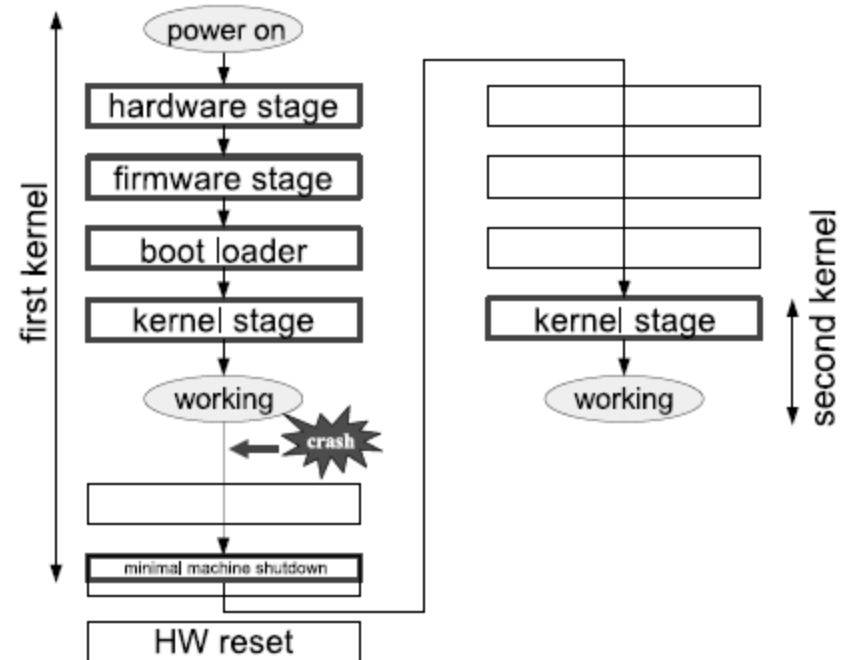


Kexec and Kdump

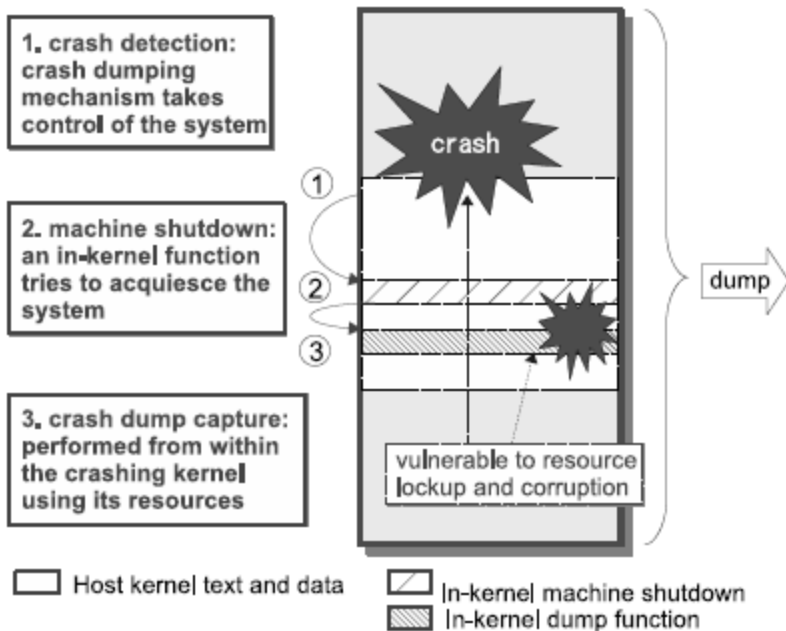
- Kexec boot



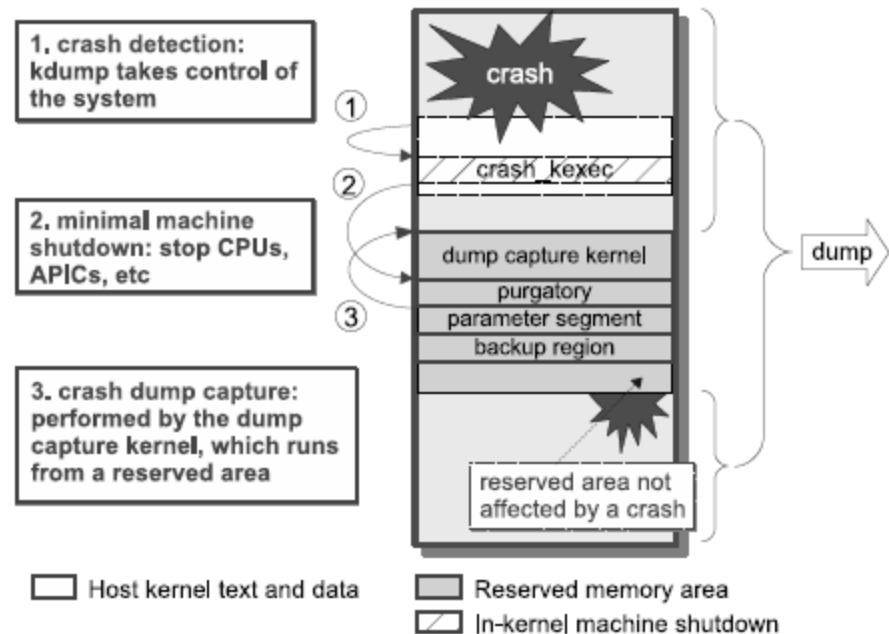
kdump boot



Kdump



(a) In-kernel crash dumping



(b) Kdump-based crash dumping

Kernel hacking options

- Some kernel hacking options are architecture-dependent
- CONFIG_PRINTK_TIME: show Timing information on printk
- CONFIG_DEBUG_SLAB: debug slab memory allocations
- CONFIG_DEBUG_SPINLOCK: finds lock-related problems
- CONFIG_MAGIC_SYSRQ: Magic SysRq key
- CONFIG_DETECT_SOFTLOCKUP: detect tight loops in kernel code that last for more than 10 seconds

Kernel hacking options


- `CONFIG_DEBUG_SLAB/CONFIG_DEBUG_HIMEM/CONFIG_DEBUG_PAGE_ALLOC` : help debug memory management problems
- `CONFIG_DEBUG_STACKOVERFLOW`: warnings if the available stack space falls below a threshold
- `CONFIG_DEBUG_STACK_USAGE`): adds stack space instrumentation to the magic Sysrq key output
- `CONFIG_DEBUG_BUGVERBOSE`: verbose `BUG()` reporting
- `CONFIG_KALLSYMS`: debug an "oops" message

`gdb`

- Compile kernel with `-g` flag
- `gdb vmlinux /proc/kcore`
- Cannot modify the kernel data
- Cannot single-step
- Cannot set breakpoint

kgdb

- <http://kgdb.linsyssoft.com/>
- Remote debug
- Kernel patched + gdb (over serial line)
- Full gdb functions



kdb

- <http://oss.sgi.com/projects/kdb/>
- built-in kernel debugger (not remote debugger)
- kernel patch
- support variable modification, breakpoints, and single-stepping, ...

kdb vs kgdb

	<i>KDB</i>	<i>KGDB</i>
Debugger environment	It is a debugger that needs to be built inside the kernel. All it requires is a console using which commands can be entered and output displayed on the console.	It requires a development machine to run the debugger as a normal process that communicates with the target using the GDB protocol over a serial cable. Recent versions of KGDB support the Ethernet interface.
Kernel support/patches required	KDB requires two patches: a common kernel patch that implements the architecture-independent functionality and an architecture-dependent patch.	KGDB makes use of a single patch that has three components: <ul style="list-style-type: none"> ■ GDB stub that implements the GDB protocol on the target side, ■ Changes to the serial (or Ethernet) driver for sending and receiving the messages between the target and the development machine, ■ The changes to the exception handlers for giving control to the debugger when an exception happens.
Support for source-level debugging	No support for source-level debugging	Support for source-level debugging provided the kernel is compiled with the <code>-g</code> flag on the development machine and the kernel source tree is available. On the development machine where the debugger application runs, <code>-g</code> option tells gcc to generate debugging information while compiling, which in conjunction with source files provides source-level debugging.

	<i>KDB</i>	<i>KGDB</i>
Debugging features offered	The most commonly used debugging features of KDB are: <ul style="list-style-type: none"> ■ displaying and modifying memory and registers ■ applying breakpoints ■ stack backtrace Along with the user-applied breakpoints, KDB is invoked when the kernel hits an irrecoverable error condition such as panic or OOPS. The user can use the output of KDB to diagnose the problem.	Supports GDB execution control commands, stack trace, and KGDB-specific watchpoints among a host of other features such as thread analysis.
Kernel module debugging	KDB provides support for kernel module debugging.	Debugging modules using KGDB is tricky because the module is loaded on the target machine and the debugger (GDB) runs on a different machine; so the KGDB debugger needs to be informed of the module load address. KGDB 1.9 is accompanied by a special GDB that can automatically detect module loading and unloading. For KGDB versions equal to or less than 1.8, the developer has to make use of an explicit GDB command <code>add-symbol-file</code> to load the module object into GDB's memory along with the module load address.
Web sites for download	http://oss.sgi.com/projects/kdb/	http://kgdb.linsyssoft.com/

Some tricks

- Debug process-related kernel code
 - Use UID
- Rate kernel print

```
static unsigned long prev_jiffy = jiffies;    /* rate limiting */

if (time_after(jiffies, prev_jiffy + 2*HZ)) {
    prev_jiffy = jiffies;
    printk(KERN_ERR "blah blah blah\n");
}


static unsigned long limit = 0;

if (limit < 5) {
    limit++;
    printk(KERN_ERR "blah blah blah\n");
}
```

Linux Profile

- Linux-built-in vs. 3rd party package
- Instrument vs. non-instrument
- Trace-based vs. Counting-based vs. Sampling based
- Kernel profiling vs. AP profiling

System Load Monitoring



System Load Monitoring

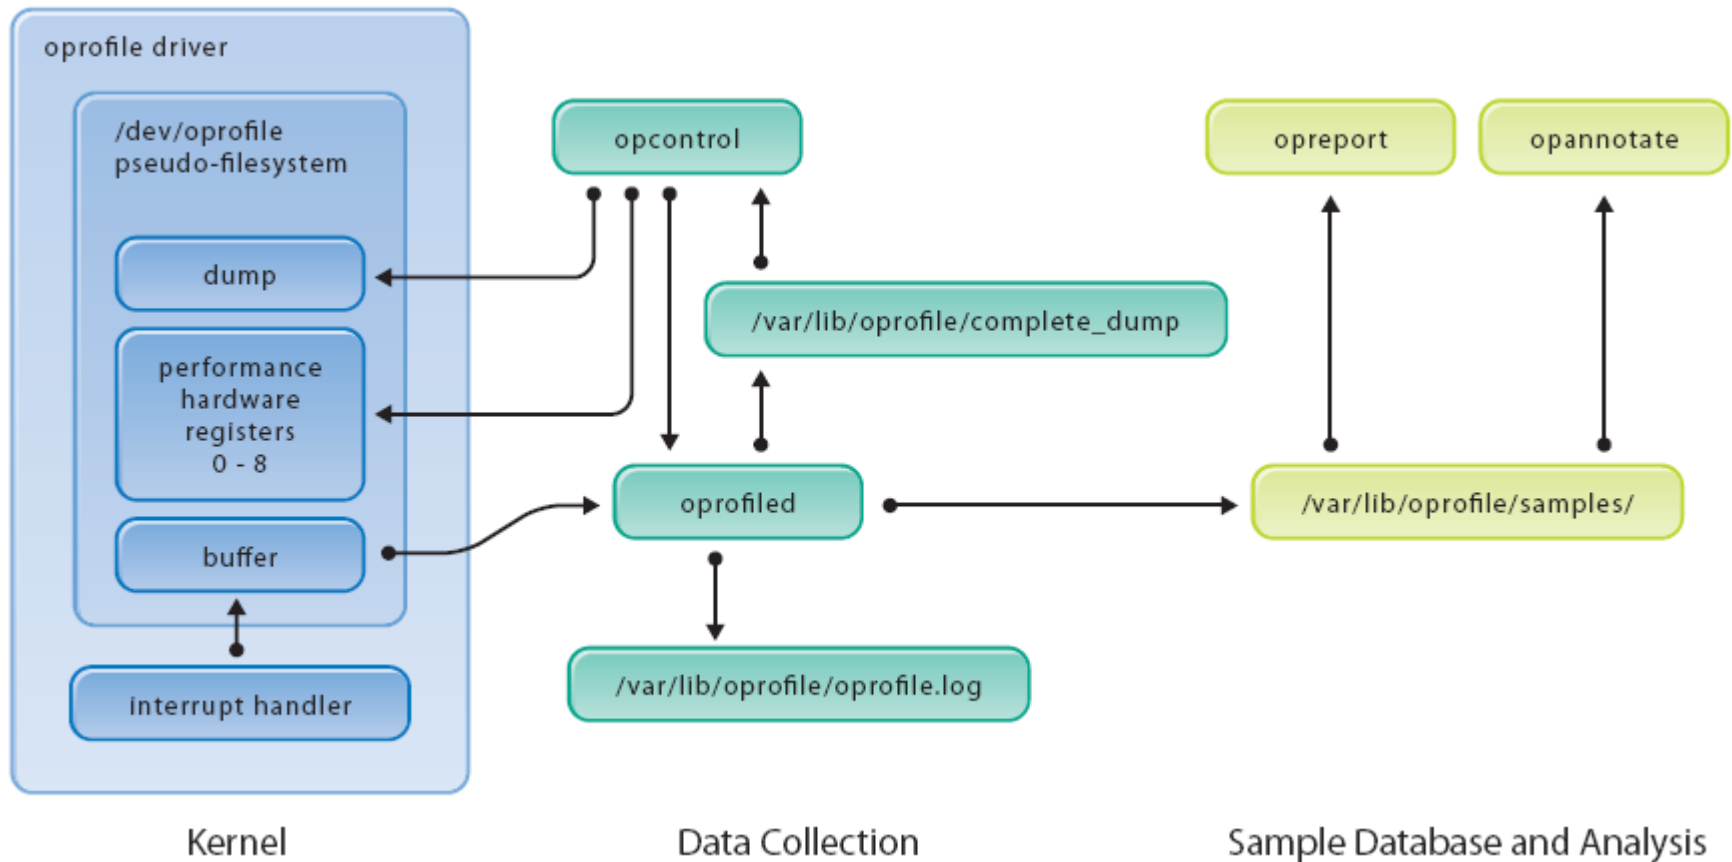
- /proc/interrupts

```
          CPU0
0: 80448940      XT-PIC  timer
1:  174412      XT-PIC  keyboard
2:      0       XT-PIC  cascade
8:      1       XT-PIC  rtc
10:  410964     XT-PIC  eth0
12:   60330     XT-PIC  PS/2 Mouse
14: 1314121     XT-PIC  ide0
15: 5195422     XT-PIC  ide1
NMI:      0
ERR:      0
```

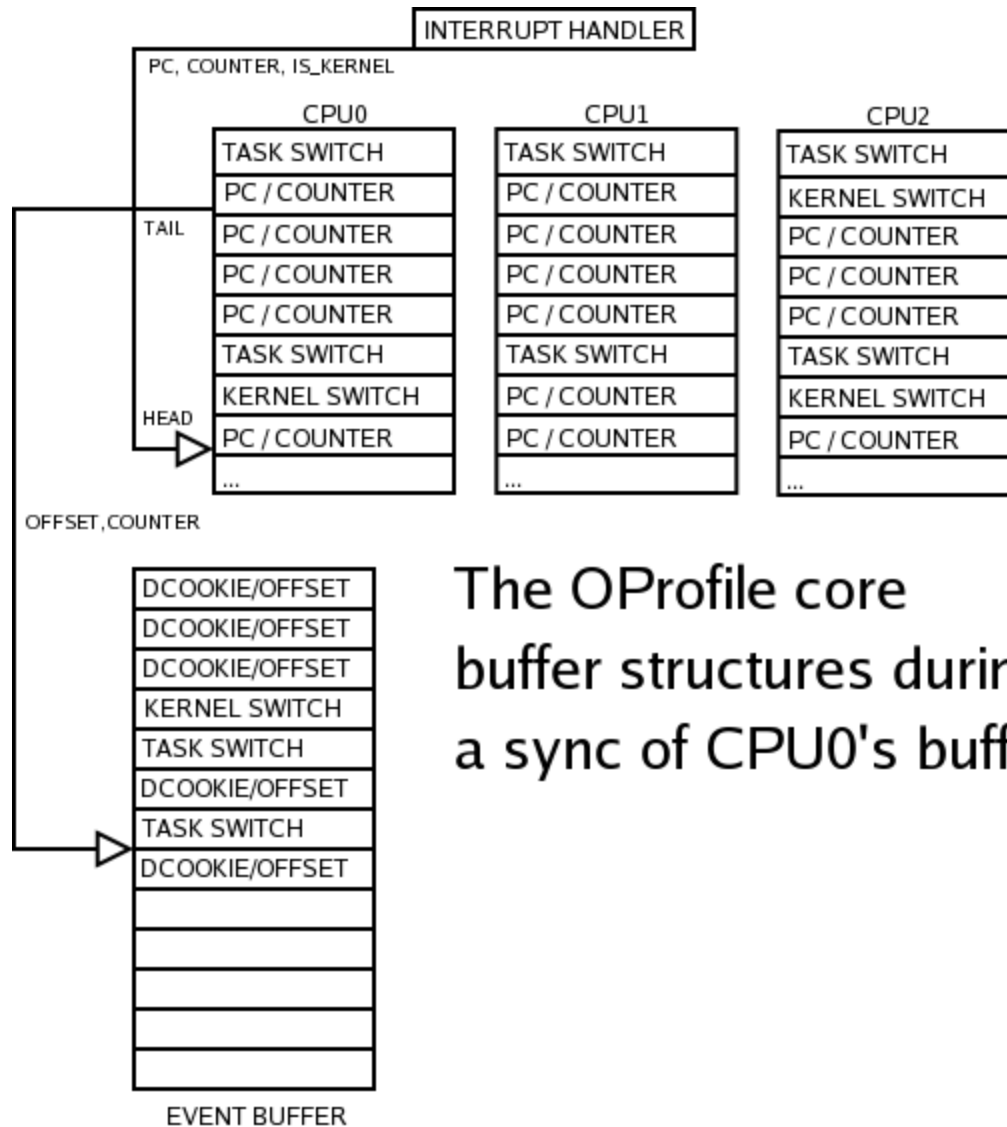
For a multi-processor machine, this file may look slightly different:

```
          CPU0          CPU1
0: 1366814704          0      XT-PIC  timer
1:    128          340  IO-APIC-edge keyboard
2:      0           0      XT-PIC  cascade
8:      0           1  IO-APIC-edge  rtc
12:   5323          5793  IO-APIC-edge PS/2 Mouse
13:      1           0      XT-PIC  fpu
16: 11184294 15940594  IO-APIC-level Intel EtherExpress Pro 10/100 Ethernet
20:  8450043 11120093  IO-APIC-level megaraid
30:   10432   10722   IO-APIC-level aic7xxx
31:      23      22   IO-APIC-level aic7xxx
NMI:      0
ERR:      0
```


OProfile Architecture



Oprofile Buffer



The OProfile core
buffer structures during
a sync of CPU0's buffer

System-wide binary image summary

```
$ opreport --exclude-dependent
CPU: PIII, speed 863.195 MHz (estimated)
Counted CPU_CLK_UNHALTED events (clocks processor is not halted) with a unit mask of 0x00 (No
 450385 75.6634 cclplus
  60213 10.1156 lyx
  29313  4.9245 XFree86
 11633  1.9543 as
 10204  1.7142 oprofiled
  7289  1.2245 vmlinux
  7066  1.1871 bash
  6417  1.0780 oprofile
  6397  1.0747 vim
  3027  0.5085 wineserver
  1165  0.1957 kdeinit
   832  0.1398 wine
...
```

Function breakdown

```
/usr/bin/opreport image:/home/wcohen/dcraw/dcraw_1 \  
-l --threshold 1
```

CPU: P4 / Xeon, speed 1495.19 MHz (estimated) Counted
GLOBAL_POWER_EVENTS events (time during which
processor is not stopped) with a unit mask of 0x01
(count cycles when processor is active) count 750000

vma	samples	%	image name	symbol name
0804d338	25428	66.3830	dcraw_1	vng_interpolate
00c4b8e0	3499	9.1346	libm-2.3.2.so	__ieee754_pow
080514a0	2997	7.8240	dcraw_1	convert_to_rgb
080517ac	1654	4.3180	dcraw_1	write_ppm
08049260	1279	3.3390	dcraw_1	decompress
00c507d0	744	1.9423	libm-2.3.2.so	__isnan
00c4e0d0	698	1.8222	libm-2.3.2.so	__pow
0804cd6c	546	1.4254	dcraw_1	scale_colors
00b6f0c0	494	1.2896	libc-2.3.2.so	getc
00c50800	386	1.0077	libm-2.3.2.so	__GI___finite

*Listing 2. Per function breakdown of samples for initial
program*

KFT

- Be careful of interpreting results
 - Duration
 - Do not subtract interrupts and thread switching
 - Delta
 - The problem may be caused by child functions
- Good for debugging straight-line code
 - No block or lock by mutex/semaphores

How to implement it

```
static inline void
__noinstrument do_func_entry(struct kft_run* run, void *this_fn,
                             void *call_site)
{
    * check for log full condition
    * acquire lock on trace log so that multiple CPUs are serialized
    * allocate space for the new entry
    * unlock the log
}

static inline void
__noinstrument do_func_exit(struct kft_run* run, void *this_fn,
                             void *call_site)
{
    collect the pid
    get the lock on trace log
    find matching entry in log - searching backwards from current log end
    when we find the entry point calculate the runtime (delta)
    check if it fits the filter criteria (if not don't log it)
    record CPU and PID
    and unlock the log file
}
```

How to implement it

Entry	Delta	PID	Function	Caller
1	9	-1	0xc02797c0	0xc02e6e70
2	2	-1	0xc02d2500	0xc027983c
5	4	-1	0xc02d0a70	0xc0279851
5	1	-1	0xc02ce590	0xc02d0ac9
6	3	-1	0xc02d0a10	0xc02d0ae4
6	2	-1	0xc02d0930	0xc02d0a37
7	0	-1	0xc014e390	0xc02d09cd
8	0	-1	0xc014e270	0xc02d0a48

How to install it

```
Kernel Hacking --->  
[*] Kernel Function Trace  
[*] Static function tracing configuration
```

- Static: kernel/kftstatic.conf
- Dynamic: cat /trace.config > /proc/kft
- Save system.map

How to configure the trace run

```
begin
  trigger start entry start_kernel
  trigger stop entry to_userspace
  filter mintime 500
end
```

- Triggers
 - trigger start entry start_kernel
 - trigger stop exit do_fork
 - trigger start time 100000000
 - trigger stop time 5000

```
trigger:
  either "start" or "stop", and then one of:
    entry <funcname>
    exit <funcname>
    time <time-in-usecs>

syntax:
trigger start|stop entry|exit|time <arg>
```

Filters

- Filters
 - filter mintime 100
 - filter maxtime 5000000
 - filter noints
 - filter onlyints
 - filter funclist do_fork sys_read fend

```
filters
    maxtime <max-time>
    mintime <min-time>
    noints
    onlyints
    funclist <func1> <func2> fend

syntax:
filter noints|onlyints|maxtime|mintime|funclist <args> fend
```

Watch

```
watches
    stack <low-water-threshold>
    worst-stack <starting-low-water-threshold>

syntax:
watch stack|worst-stack <threshold>
```

How to configure the trace run

- Static kernel/kftstatic.conf
- Dynamic
 - Edit trace.config
 - Sym2addr trace.config system.map > trace.config2
 - Cat trace.config2 > /proc/kft

```
new
begin
  trigger start entry 0xc001d804
  trigger stop time 5000000
  filter mintime 500
  filter maxtime 0
  filter noints
end
```

```
new
begin
  trigger start entry do_fork
  trigger stop exit do_fork
  filter mintime 10
  filter maxtime 400
  filter noints
  logentries 500
end
```

```
new
begin
  trigger start time 5000000
  trigger stop time 5000
  filter onlyints
end
```

How to read and process the trace results

- `Cat /proc/kft_data > kft.log`
- `addr2sym kft.log -m system.map > kft.lst`

Entry	Delta	PID	Function	Called At
23662	1333	0	con_init	console_init+0x78
25375	209045	0	calibrate_delay	start_kernel+0xf0
234425	106067	0	mem_init	start_kernel+0x130
234432	105278	0	free_all_bootmem_node	mem_init+0xc8
234435	105270	0	free_all_bootmem_core	free_all_bootmem_node+0x28
340498	4005	0	kmem_cache_sizes_init	start_kernel+0x134

Post processing by kd

- [show all functions sorted by time]
- `$./kd kftsample.lst | less`

- [show only 10 top time-consuming functions]
- `$./kd -n 10 kftsample.lst`

- [show only functions lasting longer than 100 milliseconds]
- `$./kd -t 100000 kftsample.lst`

- [show each function's most time-consuming child, and the number of times it was called.]
- `$./kd -f Fcatlmn kftsample.lst`

- [show call traces]
- `$./kd -c kftsample.lst`

- [show call traces with timing data, and functions interlaced]
- `$./kd -c -l -i kftsample.lst`

Entry	Delta	PID	Function	Called At
1	0	0	start_kernel	L6+0x0
14	8687	0	setup_arch	start_kernel+0x35
39	891	0	setup_memory	setup_arch+0x2a8
53	872	0	register_bootmem_low_pages	setup_memory+0x8f
54	871	0	free_bootmem	register_bootmem_low_pages +0x95
54	871	0	free_bootmem_core	free_bootmem+0x34
930	7432	0	paging_init	setup_arch+0x2af
935	7427	0	zone_sizes_init	paging_init+0x4e
935	7427	0	free_area_init	zone_sizes_init+0x83
935	7427	0	free_area_init_node	free_area_init+0x4b
935	3759	0	__alloc_bootmem_node	free_area_init_node+0xc5
935	3759	0	__alloc_bootmem_core	__alloc_bootmem_node+0x43
4694	3668	0	free_area_init_core	free_area_init_node+0x75
4817	3535	0	memmap_init_zone	free_area_init_core+0x2bd
8807	266911	0	time_init	start_kernel+0xb6
8807	261404	0	get_cmos_time	time_init+0x1c
270211	5507	0	select_timer	time_init+0x41
270211	5507	0	init_tsc	select_timer+0x45
270211	5507	0	calibrate_tsc	init_tsc+0x6c
275718	1638	0	console_init	start_kernel+0xbb
275718	1638	0	con_init	console_init+0x59
275954	733	0	vgacon_save_screen	con_init+0x288
277376	6730	0	mem_init	start_kernel+0xf8
277376	1691	0	free_all_bootmem	mem_init+0x52
277376	1691	0	free_all_bootmem_core	free_all_bootmem+0x24
284118	25027	0	calibrate_delay	start_kernel+0x10f
293860	770	0	__delay	calibrate_delay+0x62
293860	770	0	delay_tsc	__delay+0x26
294951	1534	0	__delay	calibrate_delay+0x62
294951	1534	0	delay_tsc	__delay+0x26

297134	1149	0	__delay	calibrate_delay+0xbe
297134	1149	0	delay_tsc	__delay+0x26
.				
.				
.				
1638605	0	145	filemap_nopage	do_no_page+0xef
1638605	0	145	__lock_page	filemap_nopage+0x286
1638605	0	145	io_schedule	__lock_page+0x95
1638605	0	145	schedule	io_schedule+0x24
1638605	0	5	schedule	worker_thread+0x217
1638605	0	1	to_userspace	init+0xa6

```
$ ~/work/kft/kft/kd -n 30 kftboot-9.1st
```

Function	Count	Time	Average	Local
schedule	192	5173790	26946	5173790
do_basic_setup	1	1159270	1159270	14
do_initcalls	1	1159256	1159256	627
delay	156	619322	3970	0
delay_tsc	156	619322	3970	619322
__const_udelay	146	608427	4167	0
probe_hwif	8	553972	69246	126
do_probe	31	553025	17839	68
ide_delay_50ms	103	552588	5364	0
isapnp_init	1	383138	383138	18
isapnp_isolate	1	383120	383120	311629
ide_init	1	339778	339778	22
probe_for_hwifs	1	339756	339756	103
ide_scan_pcibus	1	339653	339653	13
init_setup_piix	2	339640	169820	0
ide_scan_pcidv	2	339640	169820	0
piix_init_one	2	339640	169820	0
ide_setup_pci_device	2	339640	169820	242
probe_hwif_init	4	339398	84849	40

Entry	Duration	Local	Pid	Trace
4	20428	209	33	do_fork
7	6	6	33	alloc_pidmap
18	2643	84	33	copy_process
21	114	19	33	dup_task_struct
24	8	6	33	prepare_to_copy
27	2	2	33	sub_preempt_count
35	22	9	33	kmem_cache_alloc
38	2	2	33	__might_sleep
43	11	9	33	cache_alloc_refill
49	2	2	33	sub_preempt_count
60	65	6	33	__get_free_pages
63	59	14	33	__alloc_pages
65	3	3	33	__might_sleep
71	3	3	33	zone_watermark_ok
77	37	17	33	buffered_rmqueue
80	4	4	33	__rmqueue
86	3	3	33	sub_preempt_count
92	3	3	33	bad_range
98	2	2	33	__mod_page_state
103	8	5	33	prep_new_page
106	3	3	33	set_page_refs
117	2	2	33	zone_statistics
141	25	4	33	do_posix_clock_monotonic_gettime
143	21	6	33	do_posix_clock_monotonic_get
146	15	6	33	
do_posix_clock_monotonic_gettime_parts				
149	9	6	33	getnstimeofday
152	3	3	33	do_gettimeofday
169	3	3	33	copy_semundo
174	41	17	33	copy_files

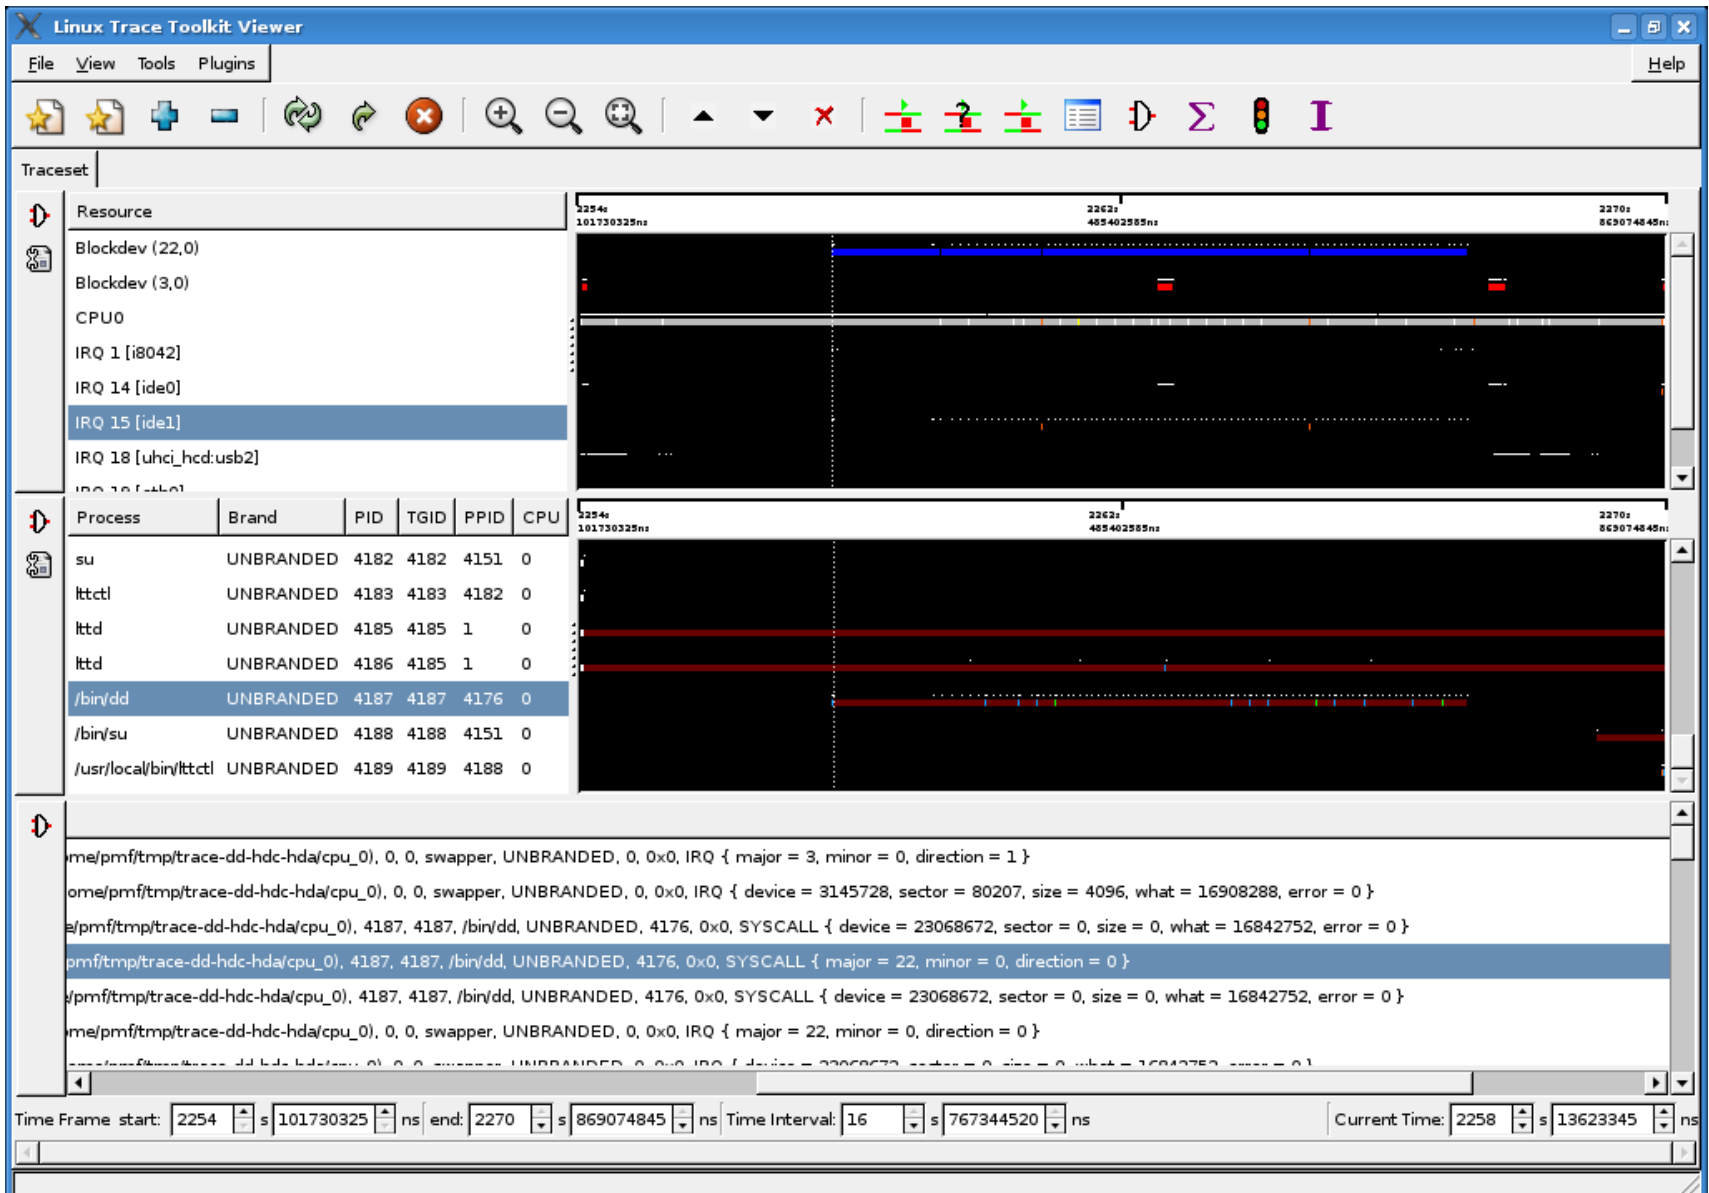
Gprof

- `cc -g -c myprog.c utils.c -pg`
- `cc -o myprog myprog.o utils.o -pg`

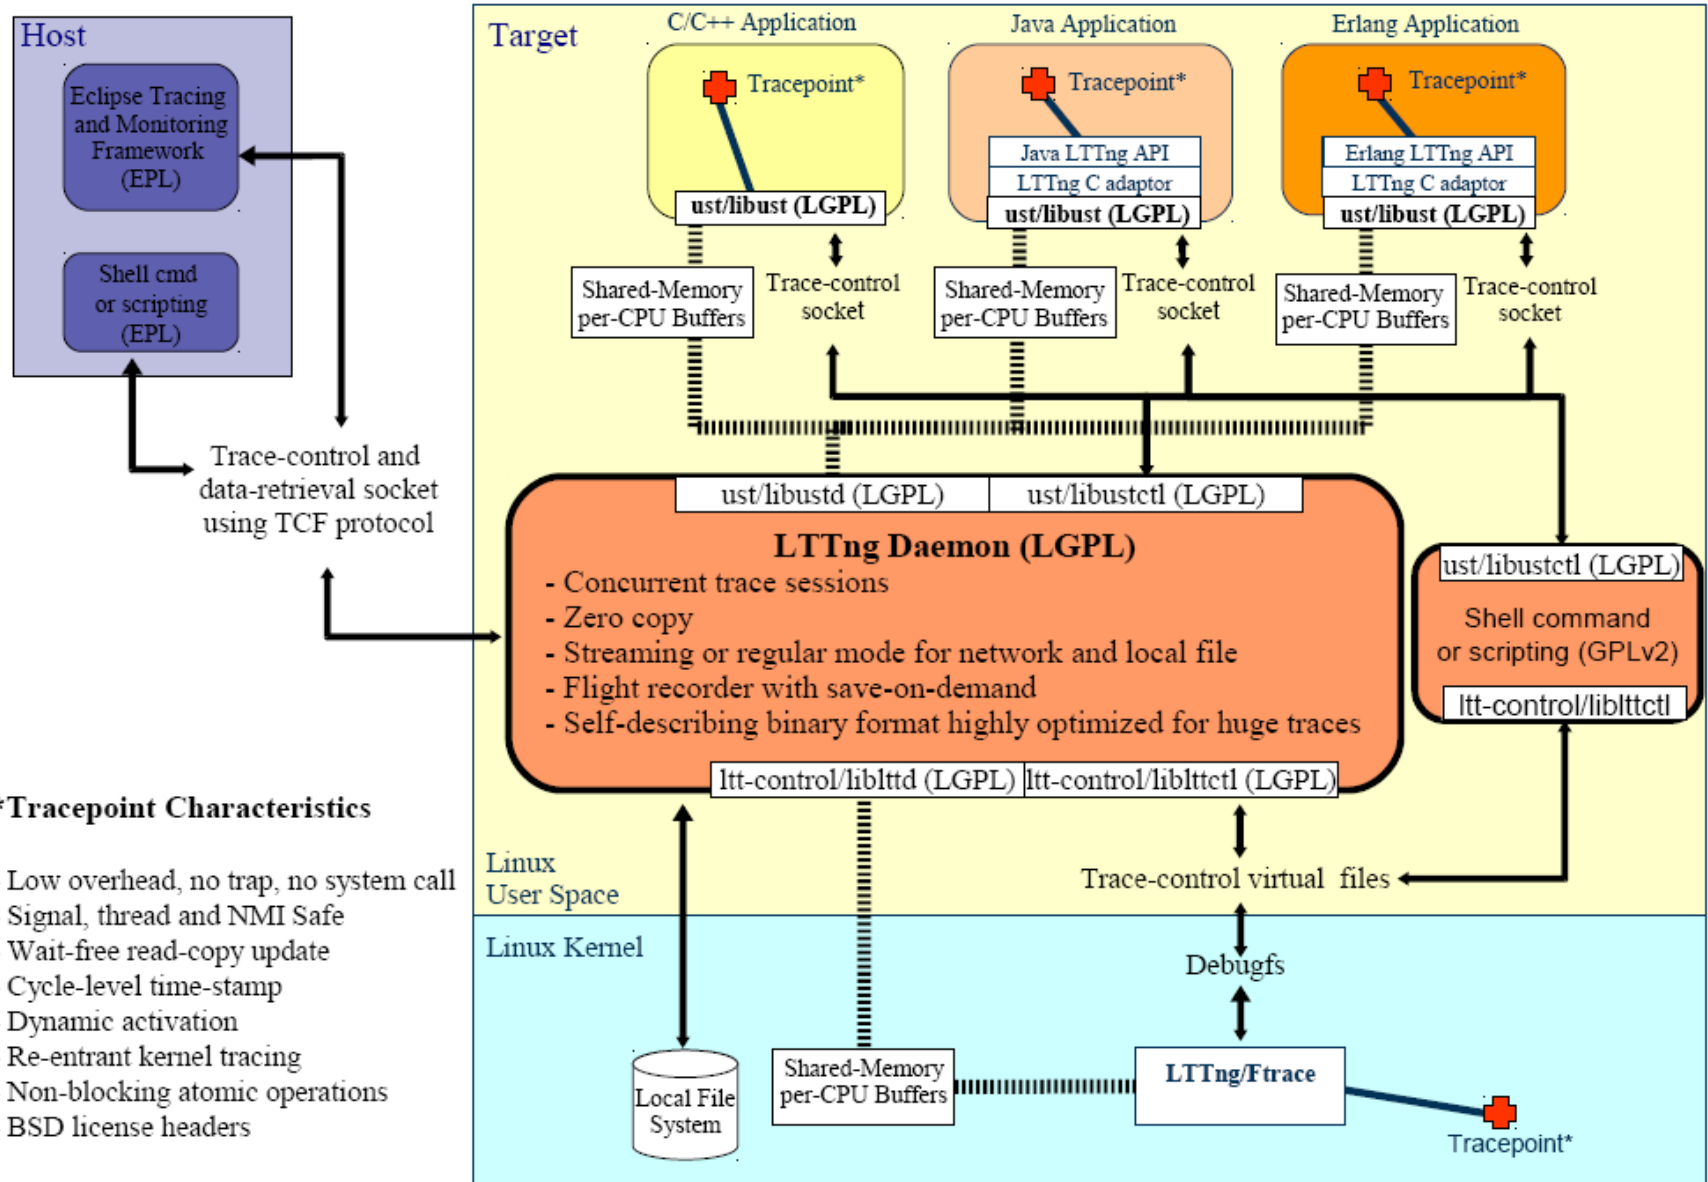
Each sample counts as 0.01 seconds.

	% cumulative	self		self	total	
time	seconds	seconds	calls	ms/call	ms/call	name
33.34	0.02	0.02	7208	0.00	0.00	open
16.67	0.03	0.01	244	0.04	0.12	offtime
16.67	0.04	0.01	8	1.25	1.25	memcpy
16.67	0.05	0.01	7	1.43	1.43	write
16.67	0.06	0.01				mcount
0.00	0.06	0.00	236	0.00	0.00	tzset
0.00	0.06	0.00	192	0.00	0.00	tolower
0.00	0.06	0.00	47	0.00	0.00	strlen
0.00	0.06	0.00	45	0.00	0.00	strchr
0.00	0.06	0.00	1	0.00	50.00	main
0.00	0.06	0.00	1	0.00	0.00	memcpy
0.00	0.06	0.00	1	0.00	10.11	print
0.00	0.06	0.00	1	0.00	0.00	profil
0.00	0.06	0.00	1	0.00	50.00	report

Linux Trace Toolkit



LTTng Low-Overhead Tracing Architecture



*Tracepoint Characteristics

- Low overhead, no trap, no system call
- Signal, thread and NMI Safe
- Wait-free read-copy update
- Cycle-level time-stamp
- Dynamic activation
- Re-entrant kernel tracing
- Non-blocking atomic operations
- BSD license headers

Other useful information

- Linux Test Project
 - <http://ltp.sourceforge.net/>
 - a suite consisting of around 3,000 tests designed to exercise different parts of the kernel
- User Mode Linux
 - <http://user-mode-linux.sourceforge.net/>
 - lets you debug the kernel without "oops"ing the machine

lockmeter

- Lockmeter is a tool for instrumenting the spin locks in a multiprocessor Linux kernel
- <http://oss.sgi.com/projects/lockmeter/>

RWLOCK	READERS	HOLD	MAX	RDR	BUSY	PERIOD	WAIT					
UTIL	CON	MEAN	READERS	MEAN	(MAX)	MEAN	(MAX)	TOTAL NAME
52.91%	0.00%	105.8us	5	114.8us	(8274.8us)	0.0us	(3.3us)	1402747	tasklist_lock		←Note 4
	0.00%						0us		28	count_active_tasks+0x10		
	0.23%						0.0us	(2.3us)	429	exit_notify+0x1c		
	0.00%						0us		5	exit_notify+0xb8		
	0.00%						0us		576	get_pid_list+0x18		
	0.00%						0.0us	(3.0us)	1224079	kill_something_info+0xb8		
	0.00%						0us		11002	proc_pid_lookup+0x4c		
	0.00%						0us		7	proc_root_lookup+0x30		
	0.00%						0us		165306	schedule+0x6d0		
	0.00%						0us		25	session_of_pgrp+0x14		
	1.12%						0.0us	(3.3us)	801	setscheduler+0x78		
	0.00%						0us		18	sys_setpgid+0x38		
	0.00%						0us		1	sys_setsid+0x10		
	0.00%						0us		461	sys_wait4+0x158		
	0.00%						0us		9	will_become_orphaned_pgrp+0x14		
0.33%	0.06%	0.5us	2	0.5us	(6.5us)	0.0us	(528us)	856780	xtime_lock		
	0.06%						0.0us	(528us)	856780	do_gettimeofday+0x10		

RWLOCK	WRITERS	HOLD	WAIT (ALL)	WAIT (WW)	SPIN	SPIN								
UTIL	CON	MEAN	(MAX)	MEAN	(MAX)	MEAN	(MAX)	TOTAL ALL WW NAME
0.00%	10.53%	0.8us	(2.7us)	9.8us	(1515us)	0.8us	(1.7us)	1691	173	5	tasklist_lock
0.00%	10.68%	1.2us	(2.6us)	7.3us	(1515us)	0.8us	(1.7us)	833	84	5	do_fork+0x8a4
0.00%	2.80%	0.2us	(0.8us)	0.2us	(13us)	0us			429	12	0	exit_notify+0x284
0.00%	17.95%	0.7us	(2.7us)	25us	(486us)	0us			429	77	0	release+0x78
0.12%	1.29%	6.2us	(802us)	0.0us	(7.7us)	0.9us	(4.3us)	28352	281	84	xtime_lock
0.03%	1.68%	2.8us	(802us)	0.0us	(7.7us)	1.0us	(3.1us)	14193	180	59	timer_bh+0x14
0.10%	0.89%	9.6us	(24us)	0.0us	(6.4us)	0.8us	(4.3us)	14159	101	25	timer_interrupt+0x14

Memory leaking

- User space: dmalloc, ...

```
/* dmalloc_test.c */

#include <stdio.h>
#include <stdlib.h>

#ifdef USE_DMALLOC
#include <dmalloc.h>
#endif

int main()
{
    char *test[5];
    unsigned int i;

    for (i=0; i < 5; i++)
    {
        unsigned int size = rand()%1024;
        test[i] = (char *)malloc(size);
        printf ("Allocated memory of size %d\n",size);
    }
    for (i=0; i<2; i++)
        free(test[i*2]);
}
```

```

calling dmalloc malloc
Allocated memory of size 359
calling dmalloc malloc
Allocated memory of size 966
calling dmalloc malloc
Allocated memory of size 105
calling dmalloc malloc
Allocated memory of size 115
calling dmalloc malloc
Allocated memory of size 81
bash>cat dlog
1094293908: 8: Dmalloc version '5.3.0' from 'http://dmalloc.com/'
1094293908: 8: flags = 0x3, logfile 'dlog'
1094293908: 8: interval = 0, addr = 0, seen # = 0, limit = 0
1094293908: 8: starting time = 1094293908
1094293908: 8: process pid = 4709
1094293908: 8: Dumping Chunk Statistics:
1094293908: 8: basic-block 4096 bytes, alignment 8 bytes, heap
grows up
1094293908: 8: heap address range: 0x80c3000 to 0x80ca000, 28672
bytes
1094293908: 8:   user blocks: 3 blocks, 12217 bytes (42%)
1094293908: 8:   admin blocks: 4 blocks, 16384 bytes (57%)
1094293908: 8: external blocks: 0 blocks, 0 bytes (0%)
1094293908: 8:   total blocks: 7 blocks, 28672 bytes
1094293908: 8: heap checked 0
1094293908: 8: alloc calls: malloc 5, calloc 0, realloc 0, free 3
1094293908: 8: alloc calls: recalloc 0, memalign 0, valloc 0
1094293908: 8: alloc calls: new 0, delete 0
1094293908: 8:   current memory in use: 1081 bytes (2 pnts)
1094293908: 8:   total memory allocated: 1626 bytes (5 pnts)
1094293908: 8:   max in use at one time: 1626 bytes (5 pnts)
1094293908: 8: max allocated with 1 call: 966 bytes
1094293908: 8: max unused memory space: 294 bytes (15%)
1094293908: 8: top 10 allocations:
1094293908: 8: total-size count in-use-size count source
1094293908: 8:   1626      5      1081  2  ra-0x8048a46
1094293908: 8:   1626      5      1081  2  Total of 1
1094293908: 8: Dumping Not-Freed Pointers Changed Since Start:
1094293908: 8: not freed: '0x80c6c00|s1' (966 bytes) from
'ra=0x8048a46'
1094293908: 8: not freed: '0x80c8f00|s1' (115 bytes) from
'ra=0x8048a46'
1094293908: 8: total-size count source
1094293908: 8:   1081      2  ra-0x8048a46
1094293908: 8:   1081      2  Total of 1
1094293908: 8: ending time = 1094293908, elapsed since start =
0:00:00

```

Memory leaking

- Kernel: kmemcheck

```
...
static int __init kmemchk_uninitialized_init(void)
{
    char * addr; /* used to store page struct addresses */
    int offset; /* offset to the page */
...
    pages = alloc_pages(GFP_KERNEL,1); /* allocate 2 pages, \
        if __GFP_NOTRACK is specified, no kmemcheck warnings would be issued */

    if(!pages)
        printk("alloc_pages: allocation failed !\n");
    else {
        addr = page_address(pages); /* convert to virt addr */

        offset = 43;
        printk("checkpoint: access mem page: %p offset: %d \n",addr,offset);
        if(*(addr + offset) == 'a' ) /* access uninitialized memory */
            printk("You hit a ramdon char \n");
    }
...
}
...
```

Kmemcheck WARNING

```
...
checkpoint: access mem page: cef52000 offset: 43
WARNING: kmemcheck: Caught 8-bit read from uninitialized memory (cef5202b) --> a
0000000000000000000000000000000000000000000000000000000000000000 --> b
u u u u u u u u u u u u u u u u u u u u u u u u u u u u u u u u u u u --> c
                                     ^                                --> d
Pid: 13017, comm: insmod Tainted: G      D W   (2.6.31.1 #2) V71          --> e
EIP: 0060:[<d09d306a>] EFLAGS: 00010286 CPU: 0
EIP is at 0xd09d306a
EAX: 00000035 EBX: cef52000 ECX: 00000092 EDX: 00885000
ESI: 00000000 EDI: b8018fc0 EBP: cef25f5c ESP: c09e2898
DS: 007b ES: 007b FS: 00d8 GS: 0000 SS: 0068
CR0: 8005003b CR2: cfbd92e0 CR3: 0ef2c000 CR4: 000006d0
DR0: 00000000 DR1: 00000000 DR2: 00000000 DR3: 00000000
DR6: ffff4ff0 DR7: 00000400
[<c0401123>] do_one_initcall+0x23/0x180
[<c0471b71>] sys_init_module+0xb1/0x1f0
[<c0403b14>] sysenter_do_call+0x12/0x28
[<ffffffff>] 0xffffffff
...
```

objdump

```
...
if(!pages)
    printk("alloc_pages: allocation failed !\n");
else {
    addr = page_address(pages); /* convert to virt addr */
4b: e8 fc ff ff ff      call    4c <init_module+0x4c>
50: 89 c3                mov     %eax,%ebx

    offset = 43;
    printk("checkpoint: access mem page: %p offset: %d \n",addr,offset);
52: c7 44 24 08 2b 00 00 movl    $0x2b,0x8(%esp)
59: 00
5a: 89 44 24 04          mov     %eax,0x4(%esp)
5e: c7 04 24 24 00 00 00 movl    $0x24,(%esp)
65: e8 fc ff ff ff      call    66 <init_module+0x66>
    if(*(addr + offset) == 'a' ) /* access uninitialized memory */
6a: 80 7b 2b 61          cmpb    $0x61,0x2b(%ebx)
6e: 75 d3                jne     43 <init_module+0x43>
    printk("You hit a ramdon char \n");
...

```


Test equipment

- JTAG/ICE debugger
- Logical analyzer
- NIC
 - Sniffer
- USB
 - Protocol Analyzer
- ...

How to maintain kernel codes

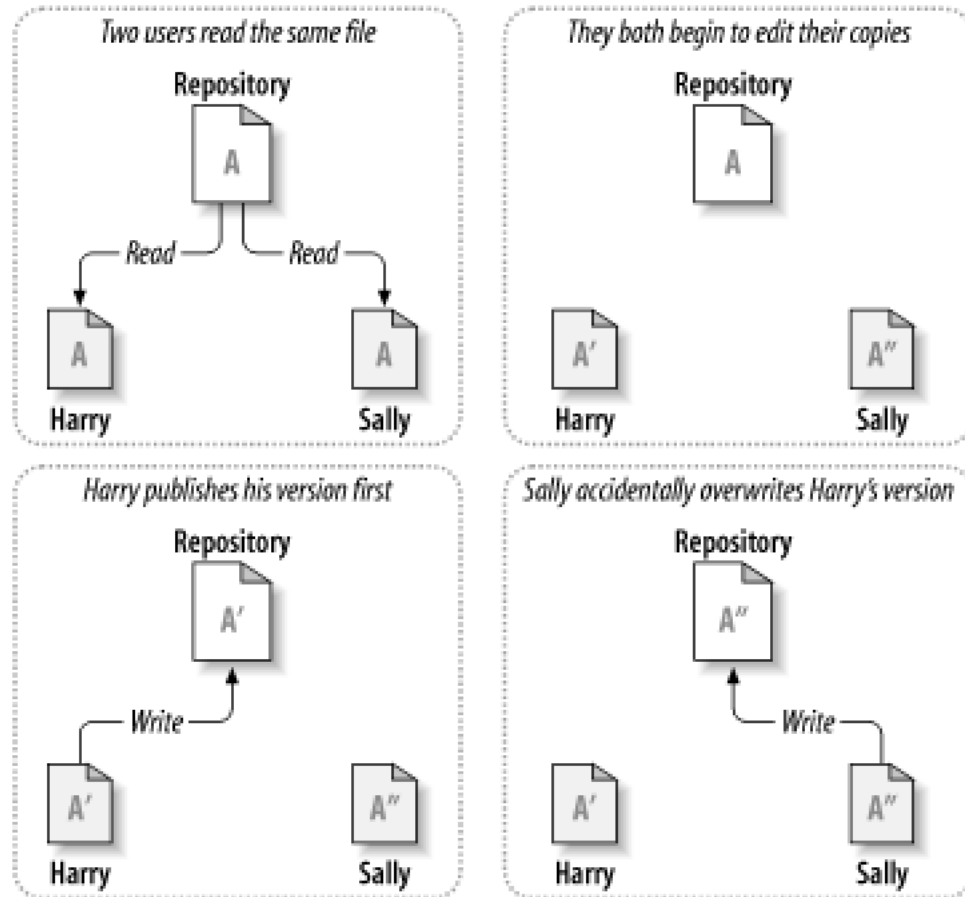
- Why version control?
 - Single developer
 - Multiple developers

How to maintain kernel codes



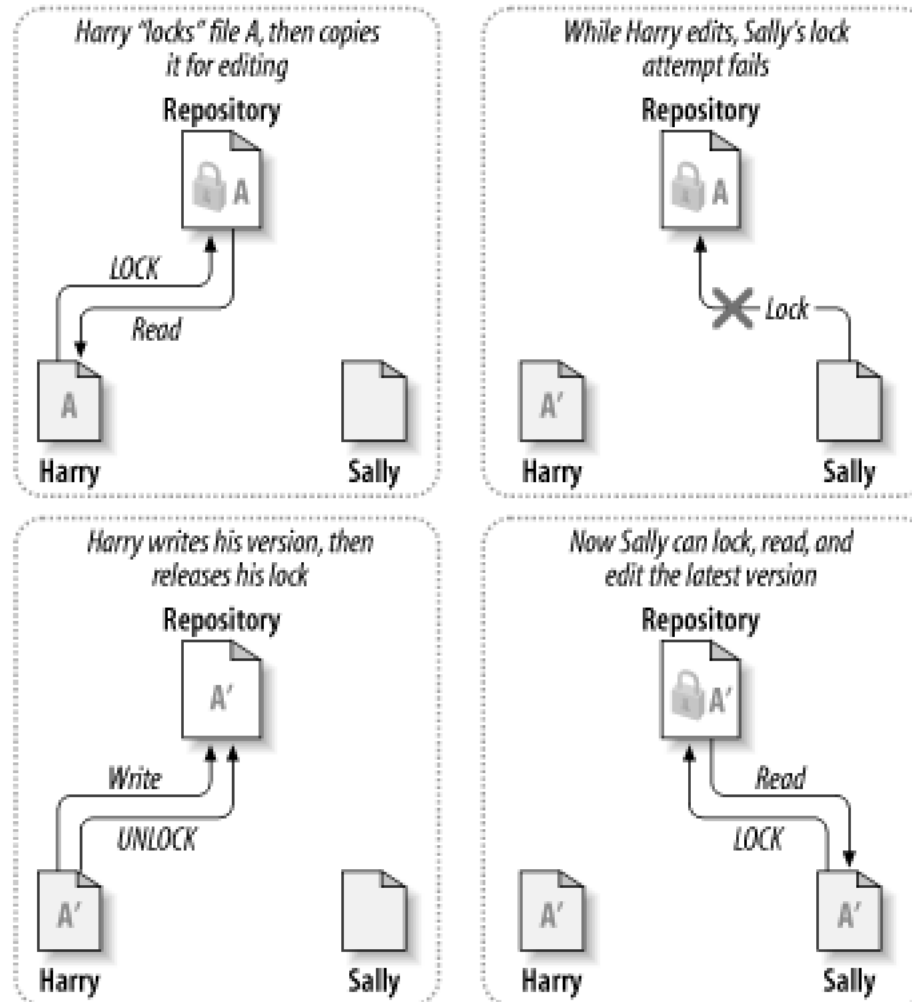
Version Control with Subversion Paperback by [C. Michael Pilato](#), [Ben Collins-Sussman](#), [Brian W. Fitzpatrick](#)

lock-modify-unlock problem



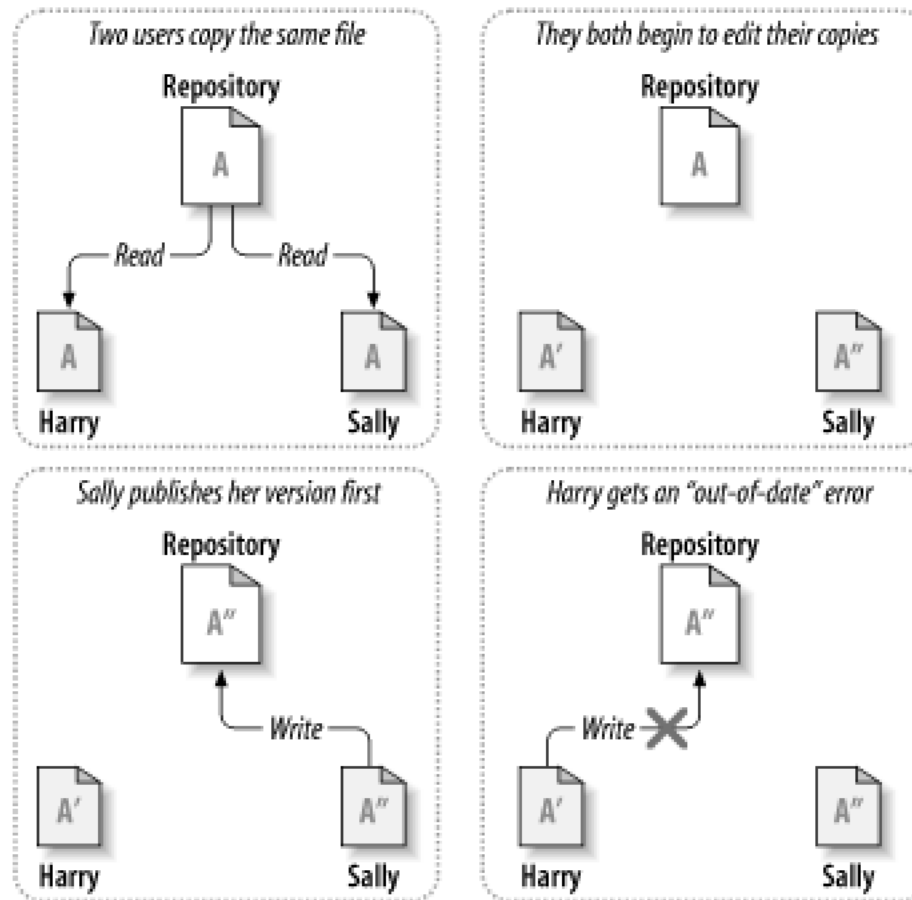
Version Control with Subversion Paperback by [C. Michael Pilato](#), [Ben Collins-Sussman](#), [Brian W. Fitzpatrick](#)

lock-modify-unlock solution

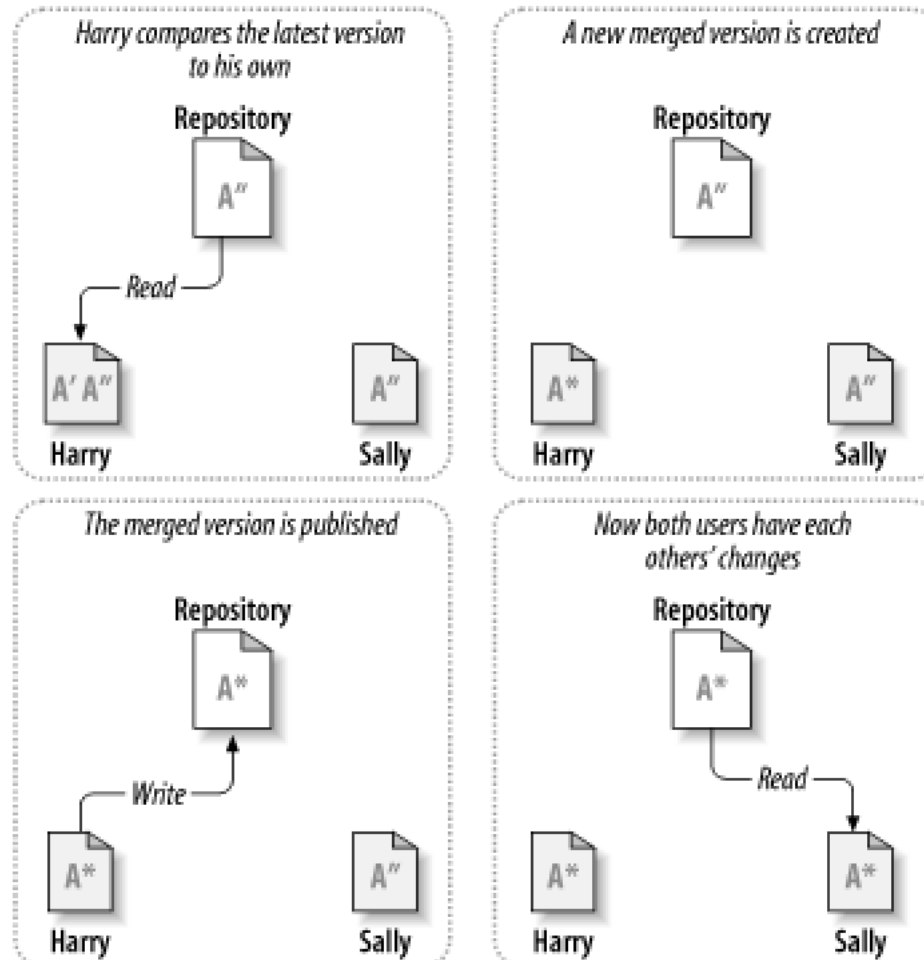


Version Control with Subversion Paperback by [C. Michael Pilato](#), [Ben Collins-Sussman](#), [Brian W. Fitzpatrick](#)

copy-modify-merge problem



copy-modify-merge solution



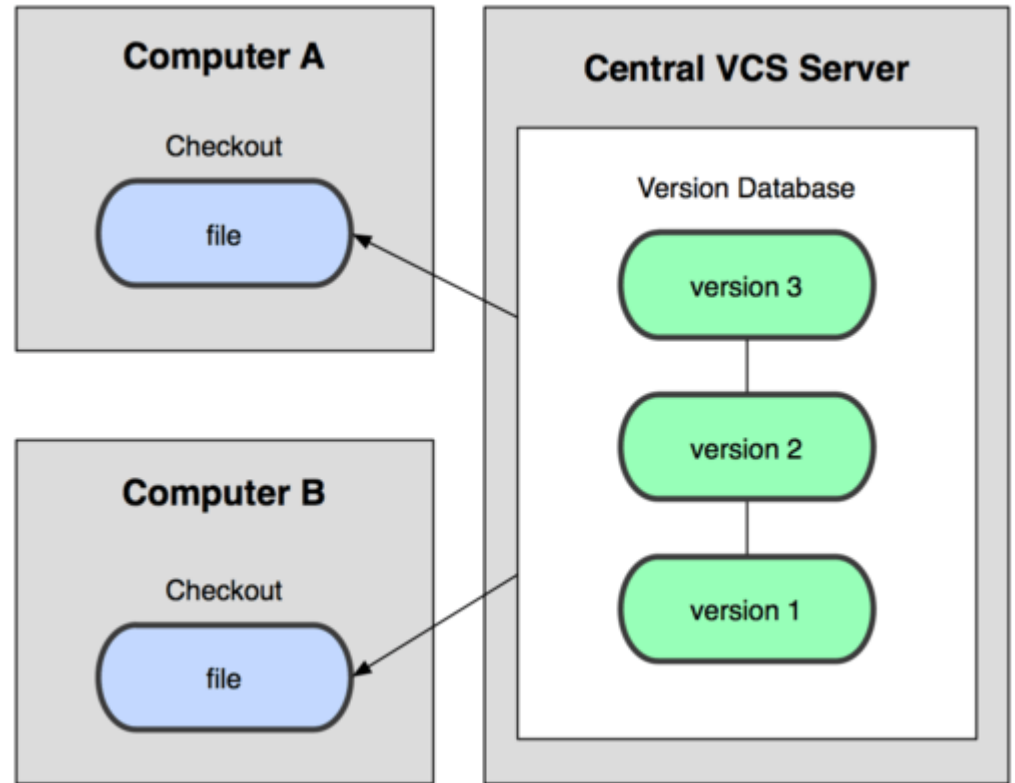
Version Control with Subversion Paperback by [C. Michael Pilato](#), [Ben Collins-Sussman](#), [Brian W. Fitzpatrick](#)

Basic Work Cycle

- Update your working copy (check out your code)
- Make your changes (modify, add, remove, copy, move files/directories) on your working copy
- Review your changes you've made in your working copy
- Fix your mistakes (may start all over from unmodified state)
- Resolve any conflicts (merge others' changes)
- Publish (commit) your changes (lock and commit)
Others can see your work, too!

Centralised Version Control

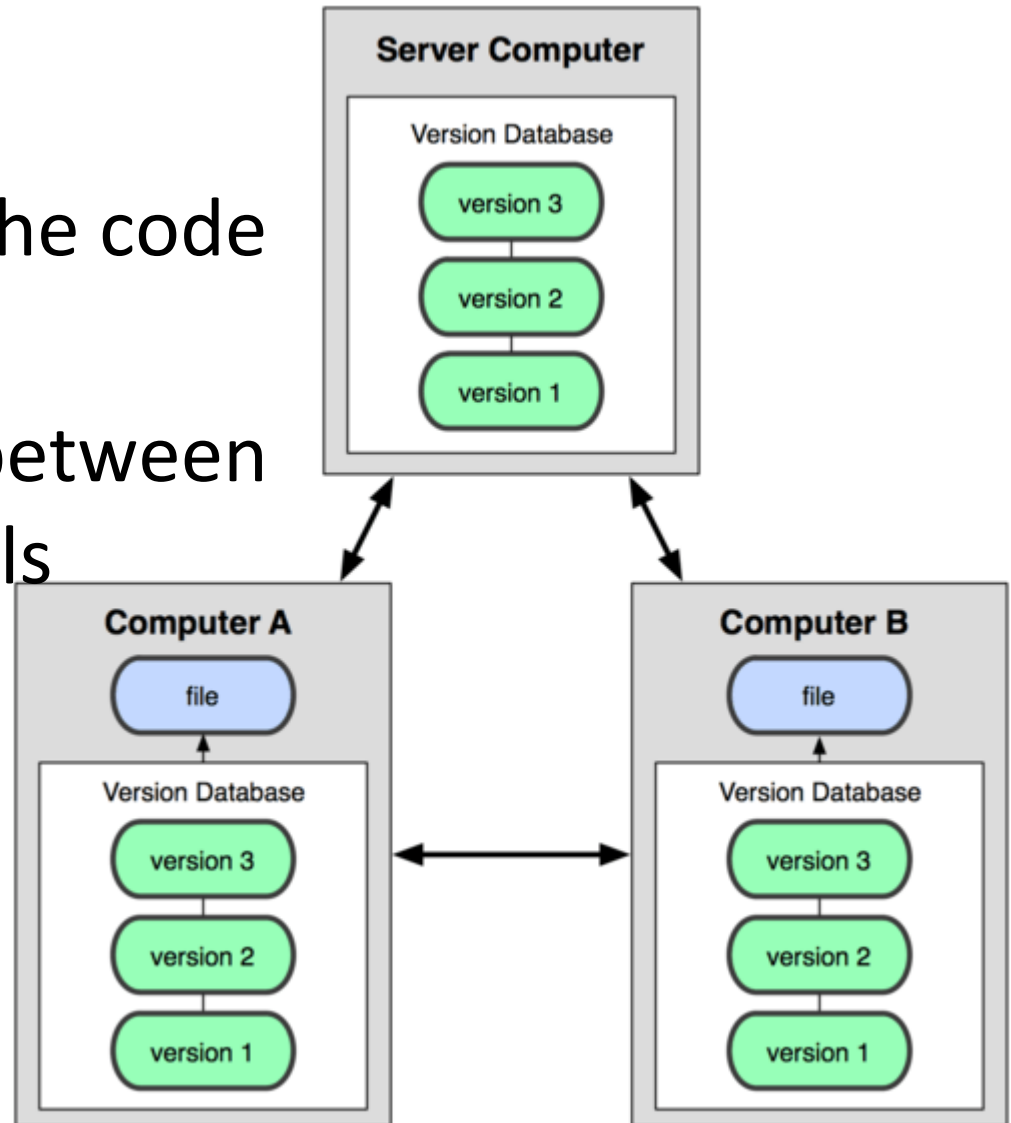
- One server holds the code base
- Clients access the server by check-in/check-outs
- CVS, SVN



What's the problem with SVN/
CVS?

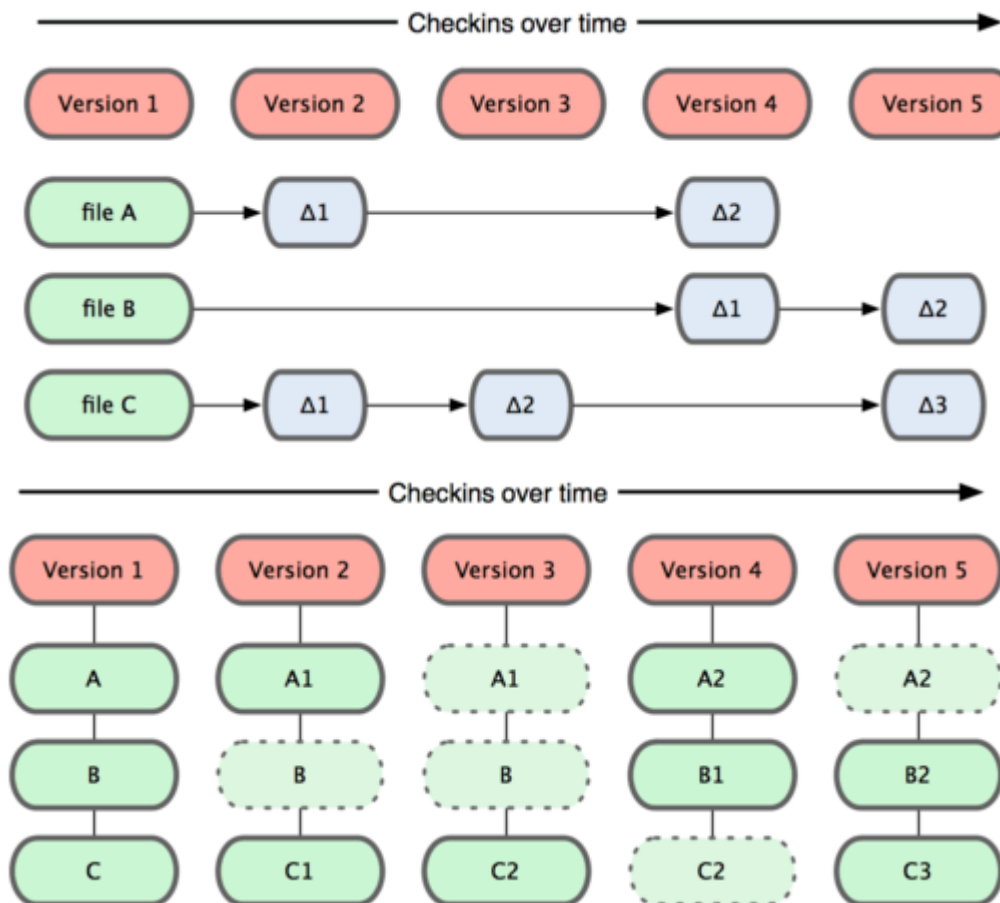
Distributed Version Control

- Each client holds a complete copy of the code base
- Codes are shared between clients by push/pulls
- GIT

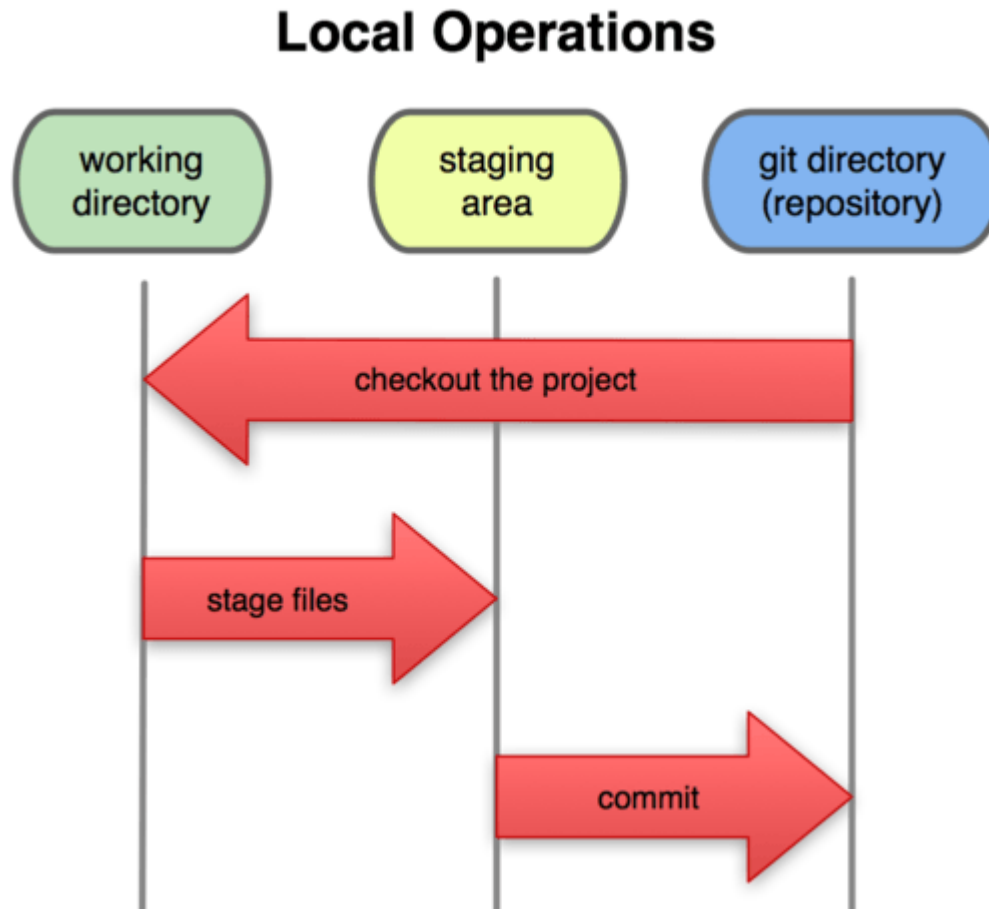


GIT

- Delta form vs. snapshot form

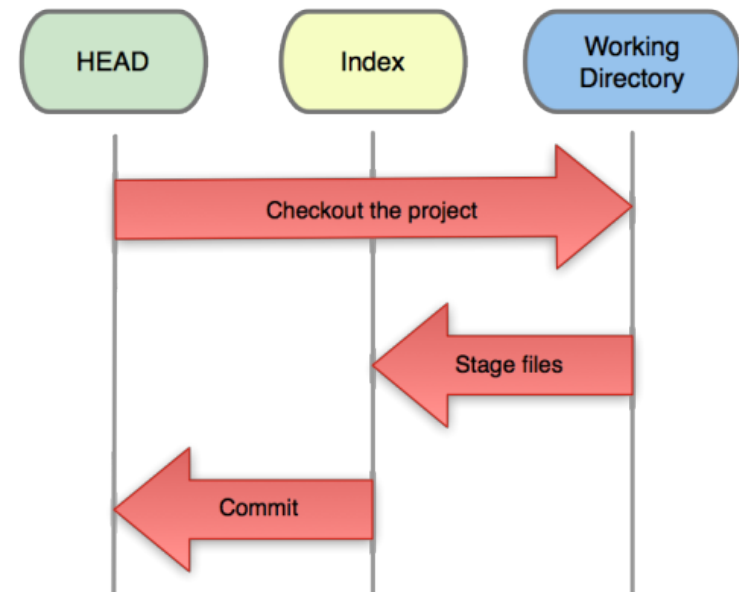


Local operations



Basic GIT workflow

- Init a repo
- Edit files
- Stage the changes
- Review your changes
- Commit the changes

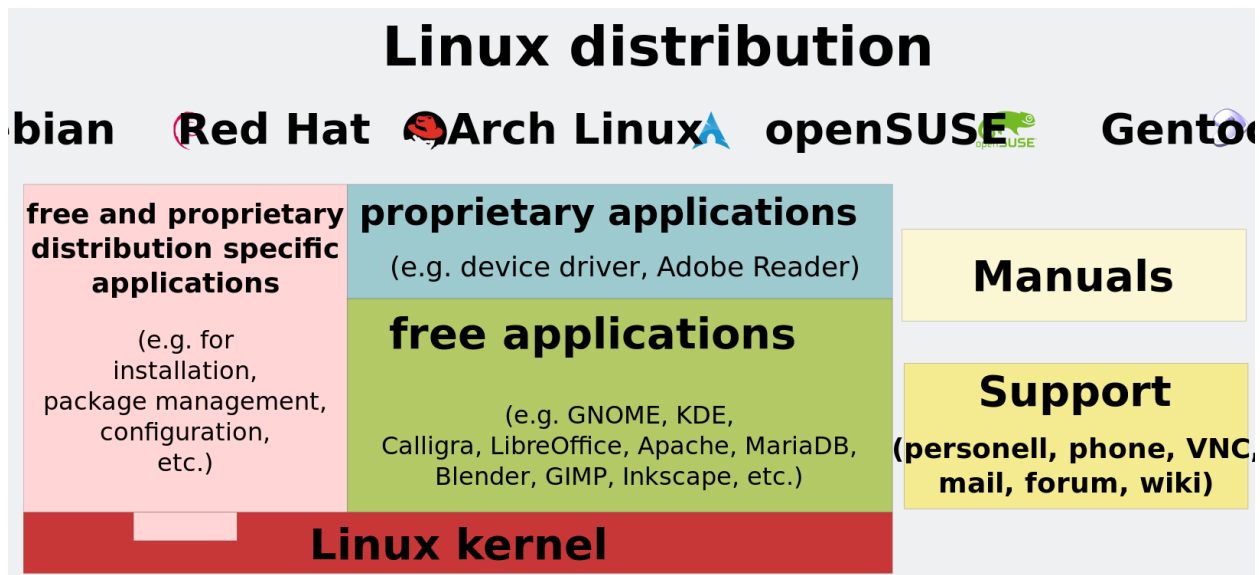


A case study of Linux

- Mainline
 - <http://www.kernel.org/>
- Other Linux Porting
 - ARM Linux (<http://www.arm.linux.org.uk/>)
- Linux versions (A.B.C)
 - A: kernel version
 - B: major revision (even-odd system version numbering system)
 - C: minor revision (every 2–3 months)

Linux Distribution

- Linux Distribution
 - Fedora Core
 - OpenSuSE
 - Debian
 - Ubuntu, Kubuntu
 - Gentoo
 - Slackware



Android

