

Operating System Design and Implementation

Booting process

Charles Tsao

Outline

- Why boot-loader/OS booting are necessary?
- How boot-loader work?
- How does OS boot?
- A case study of PC and Linux 0.11
- Advanced topics

Why boot-loader/OS booting are necessary?

- Bootloader vs. OS booting
 - Bootloader: The program (boot-loader) loads OS into memory
 - OS booting: initialize and prepare for providing services

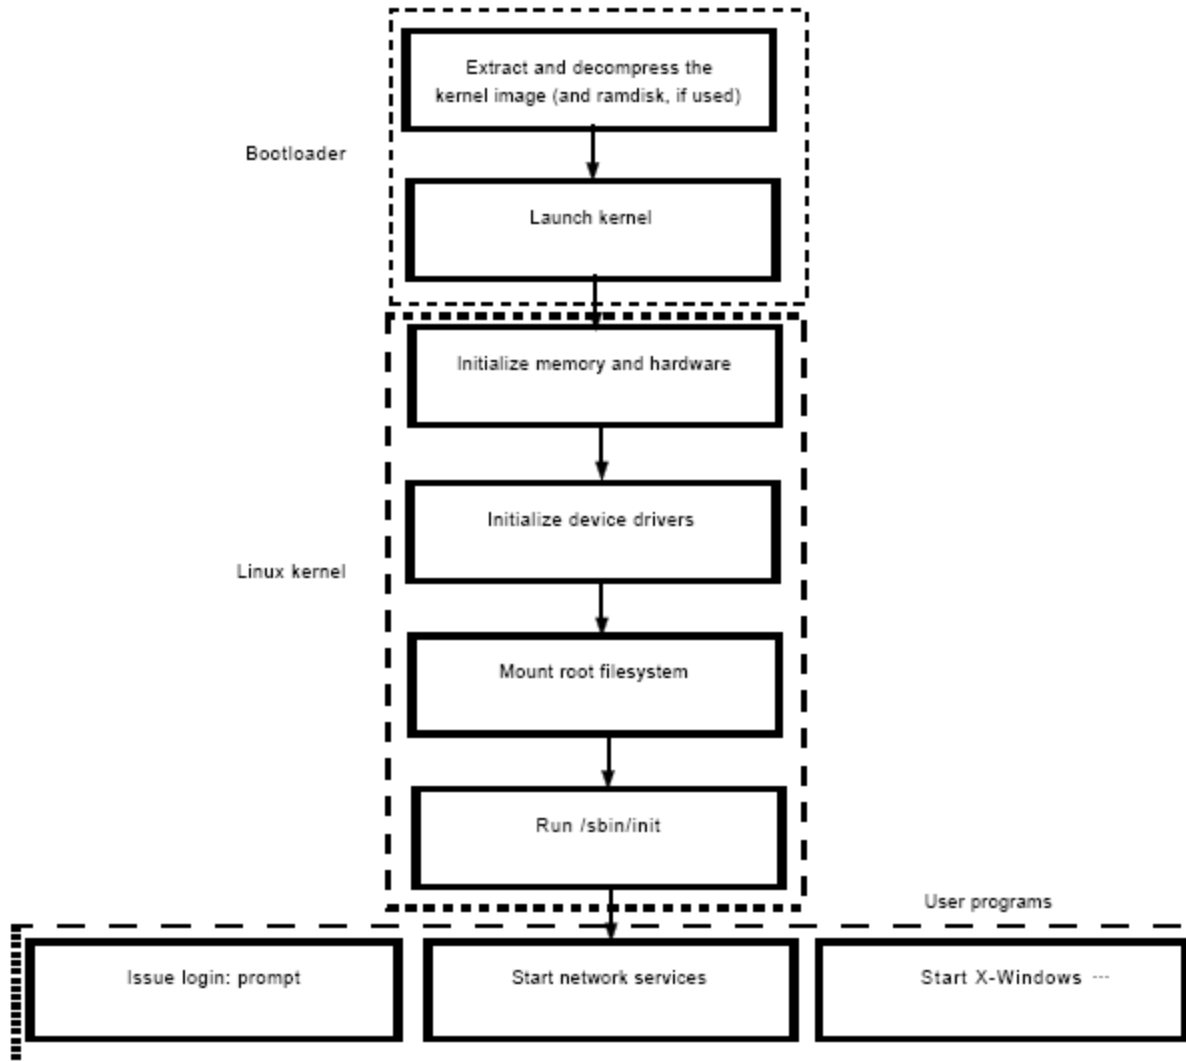
Why a boot-loader is necessary?

- Where is operating system?
- Where is the boot-loader?
 - Permanent storage
 - Addressable
 - ROM, PROM, EPROM, EEPROM, NOR Flash, NV-RAM, ...
- Can we store OS on permanent and addressable memory?
 - Why and why not?

Why OS booting is necessary?

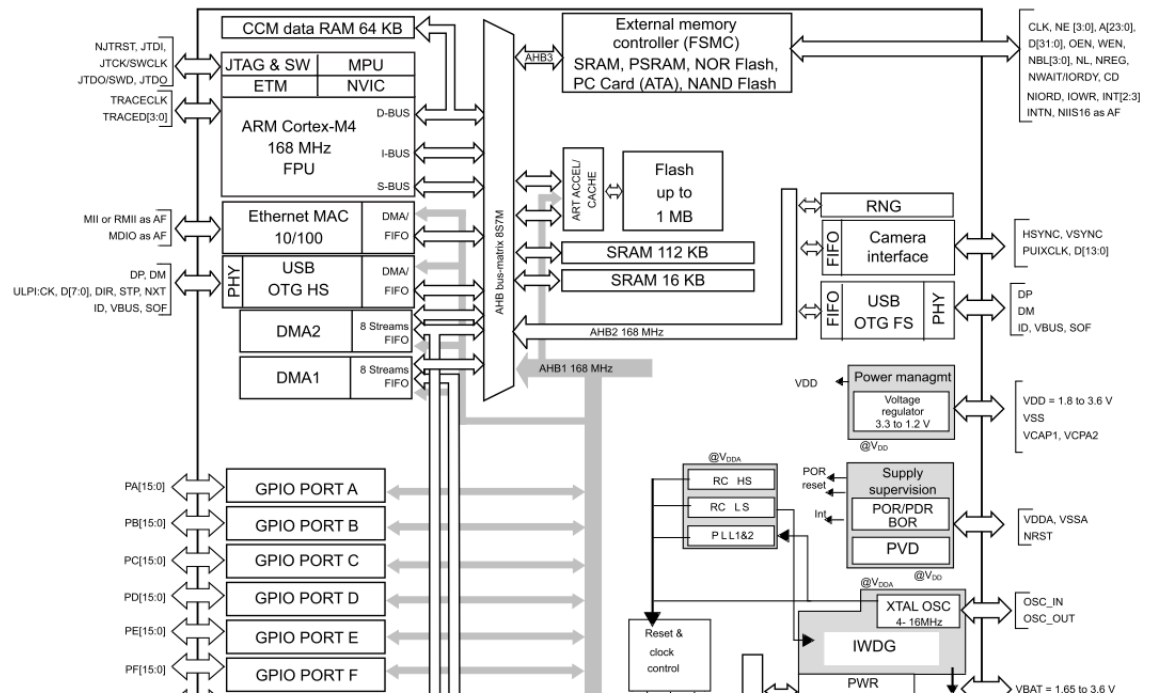
- Load OS/modules itself to correct memory address
- Initialize hardware
- Initialize kernel structure
- Load necessary default process
- ...

Bootloader and OS

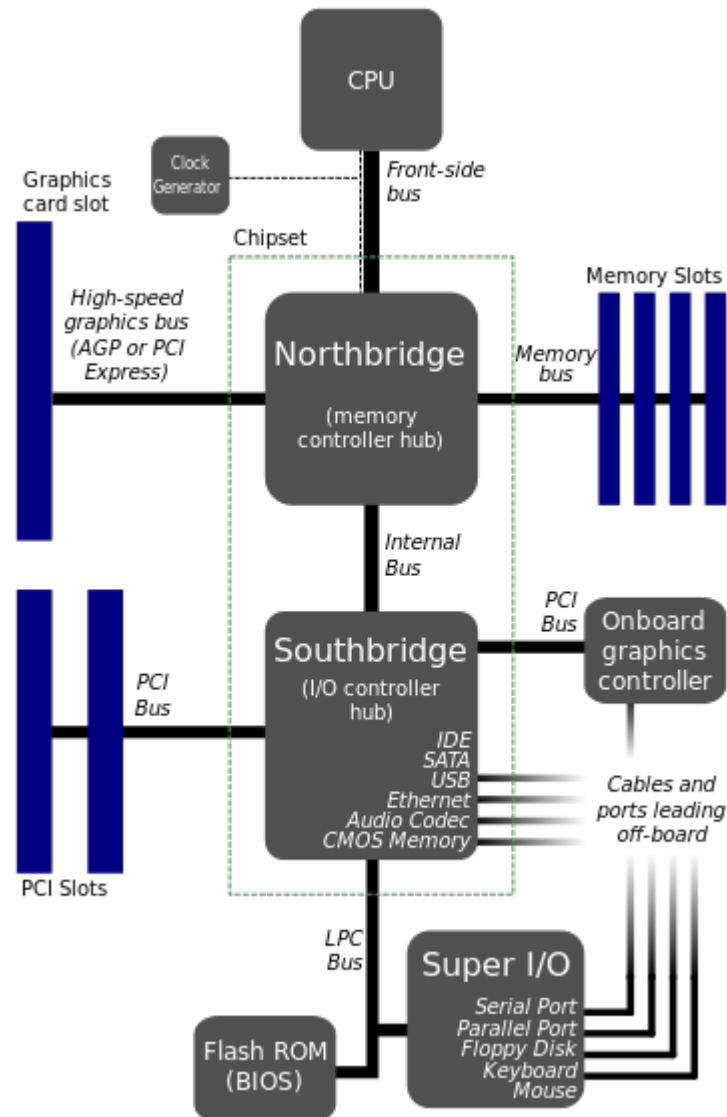


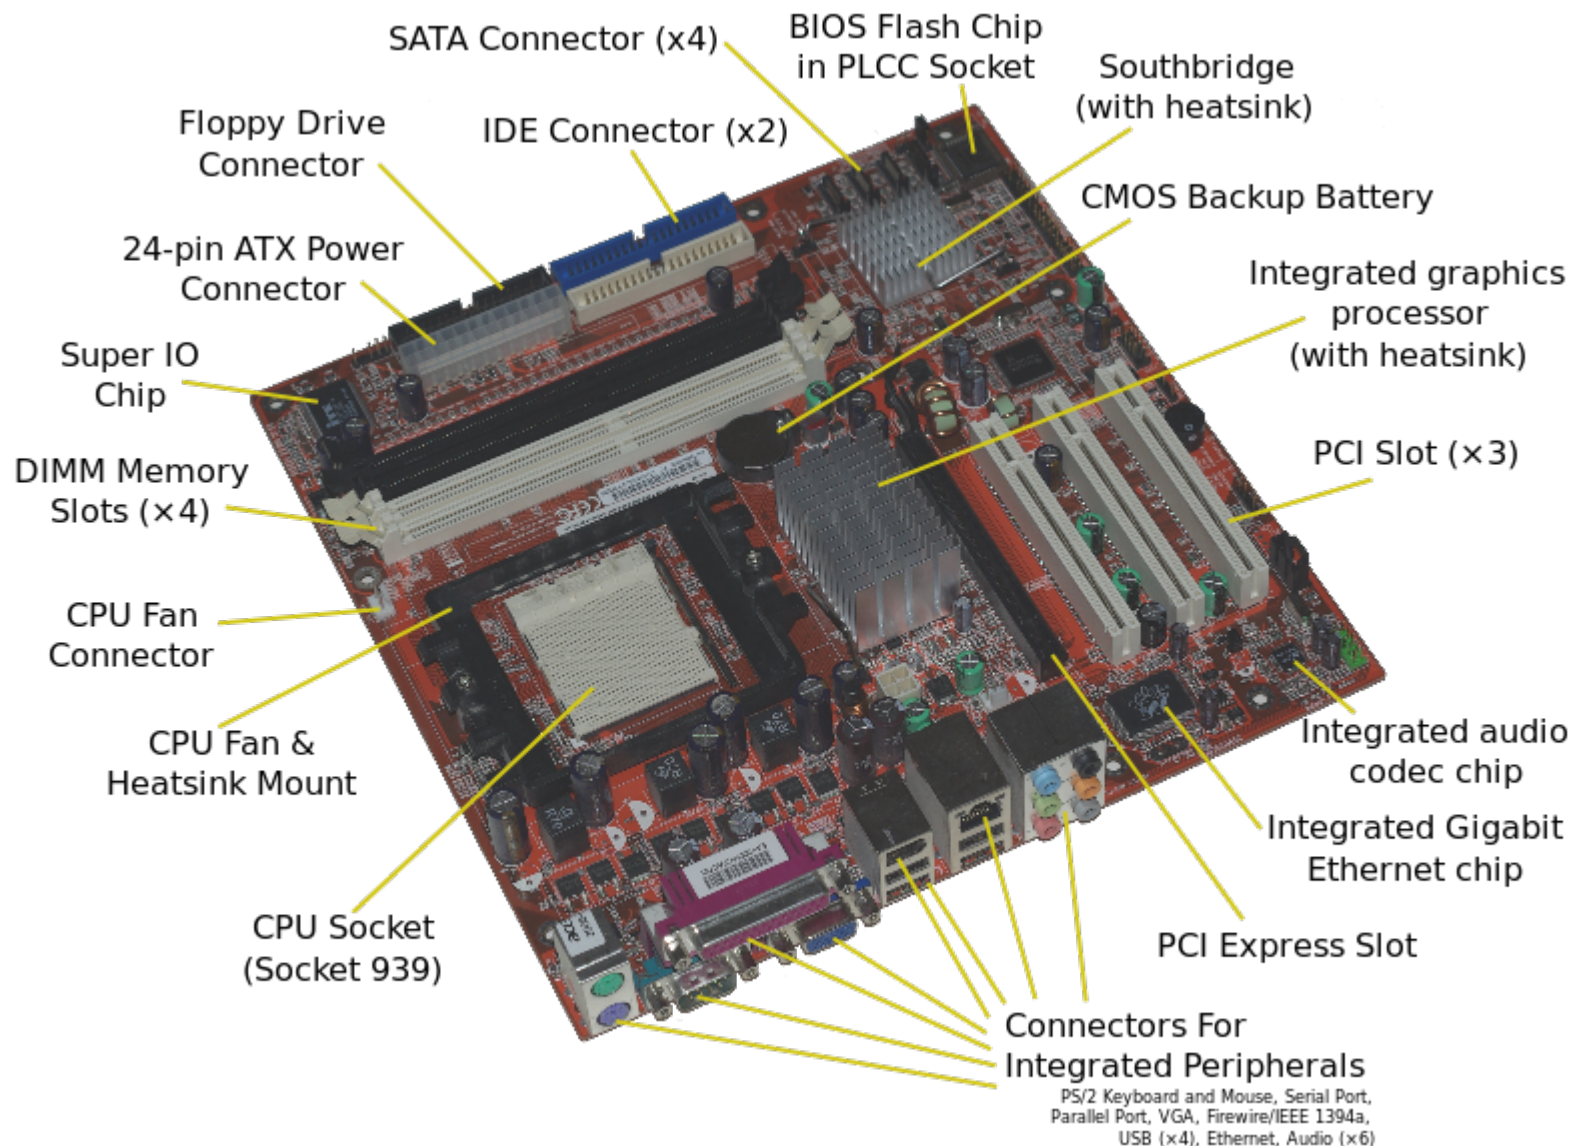
How boot-loader work?

- SoC: e.g. STM32F4

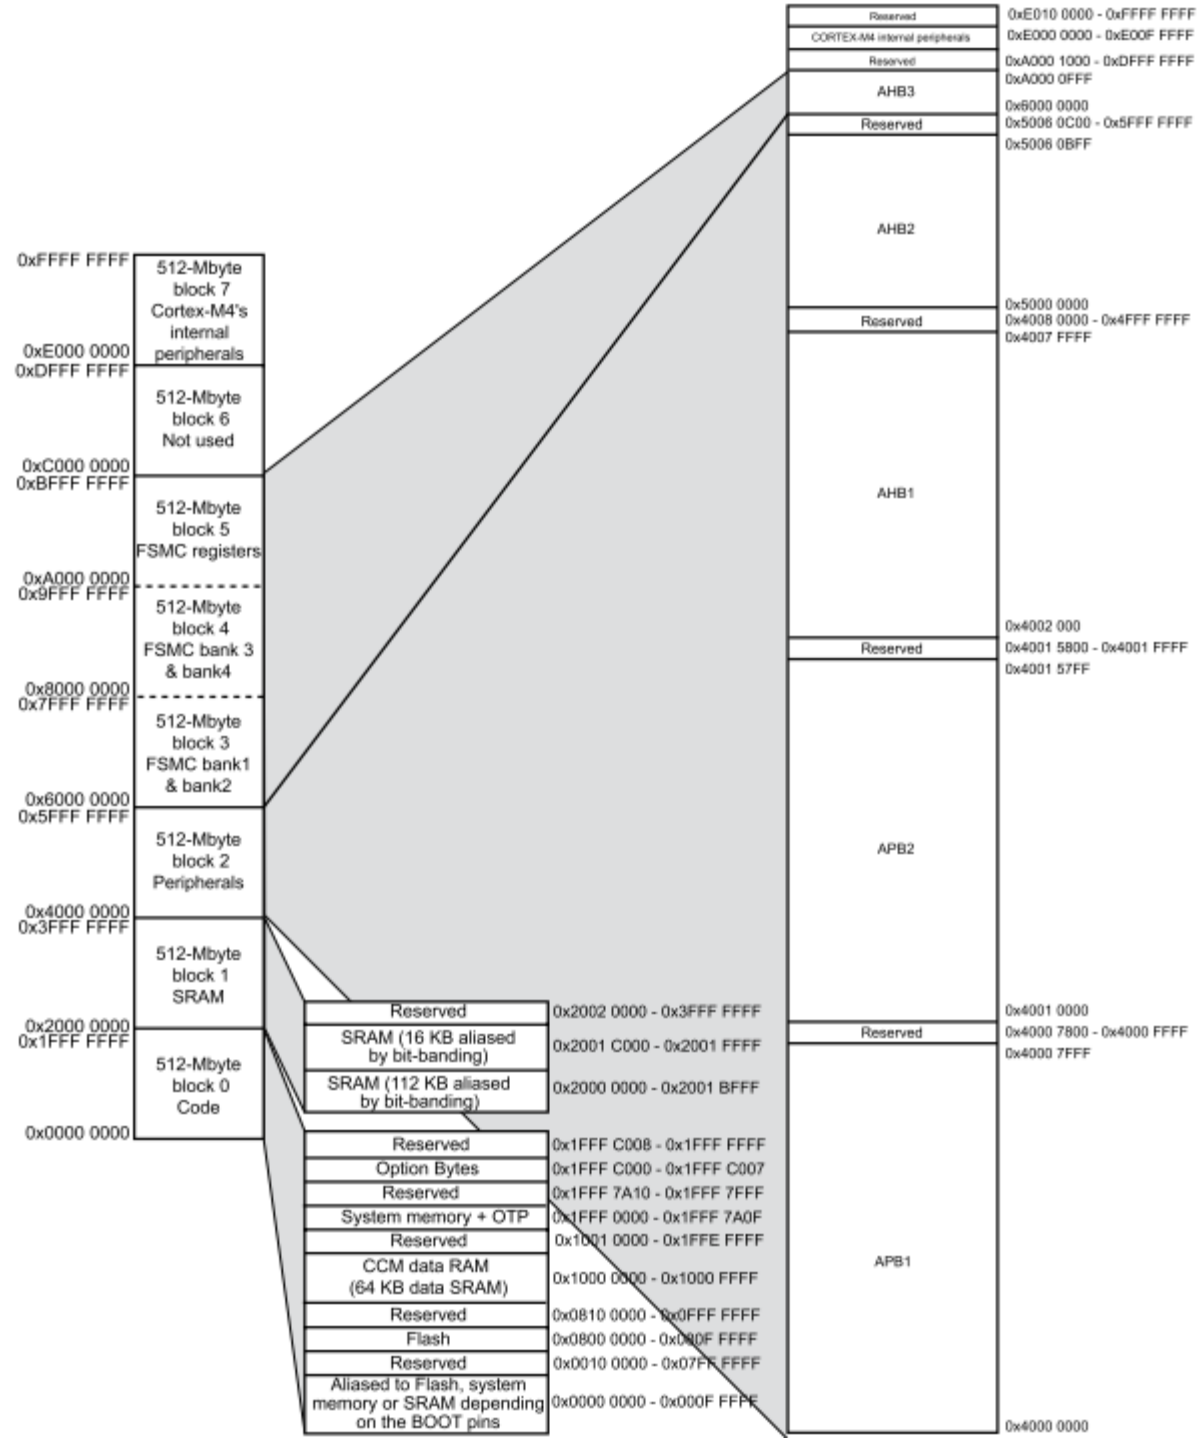


x86

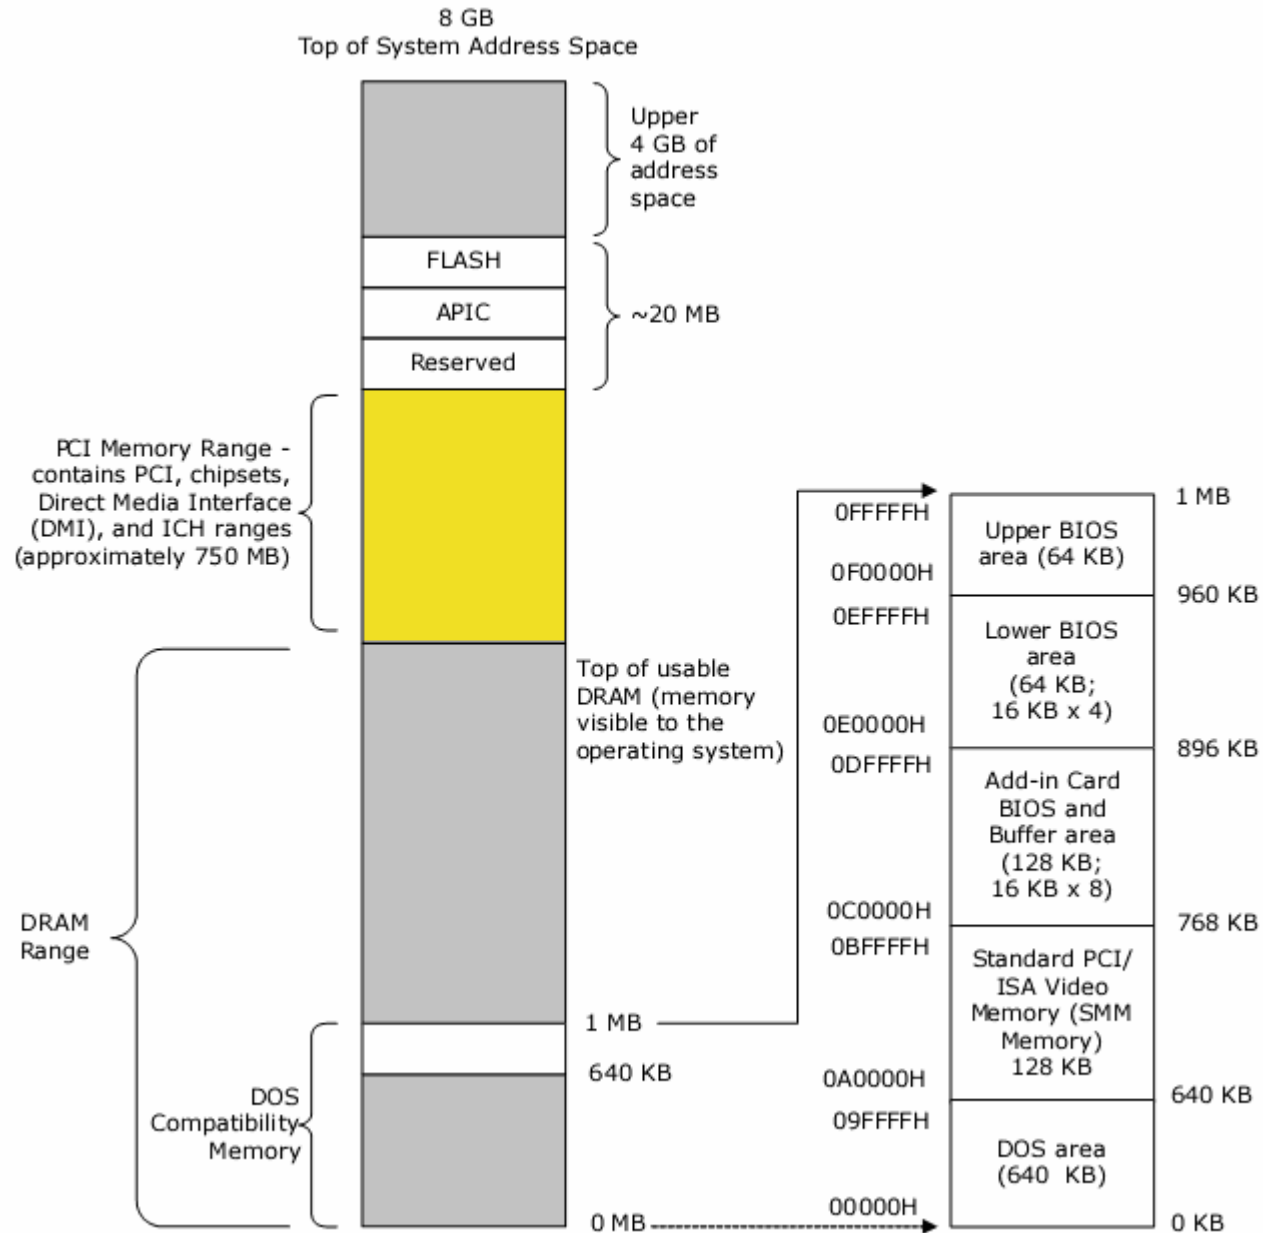




STM32



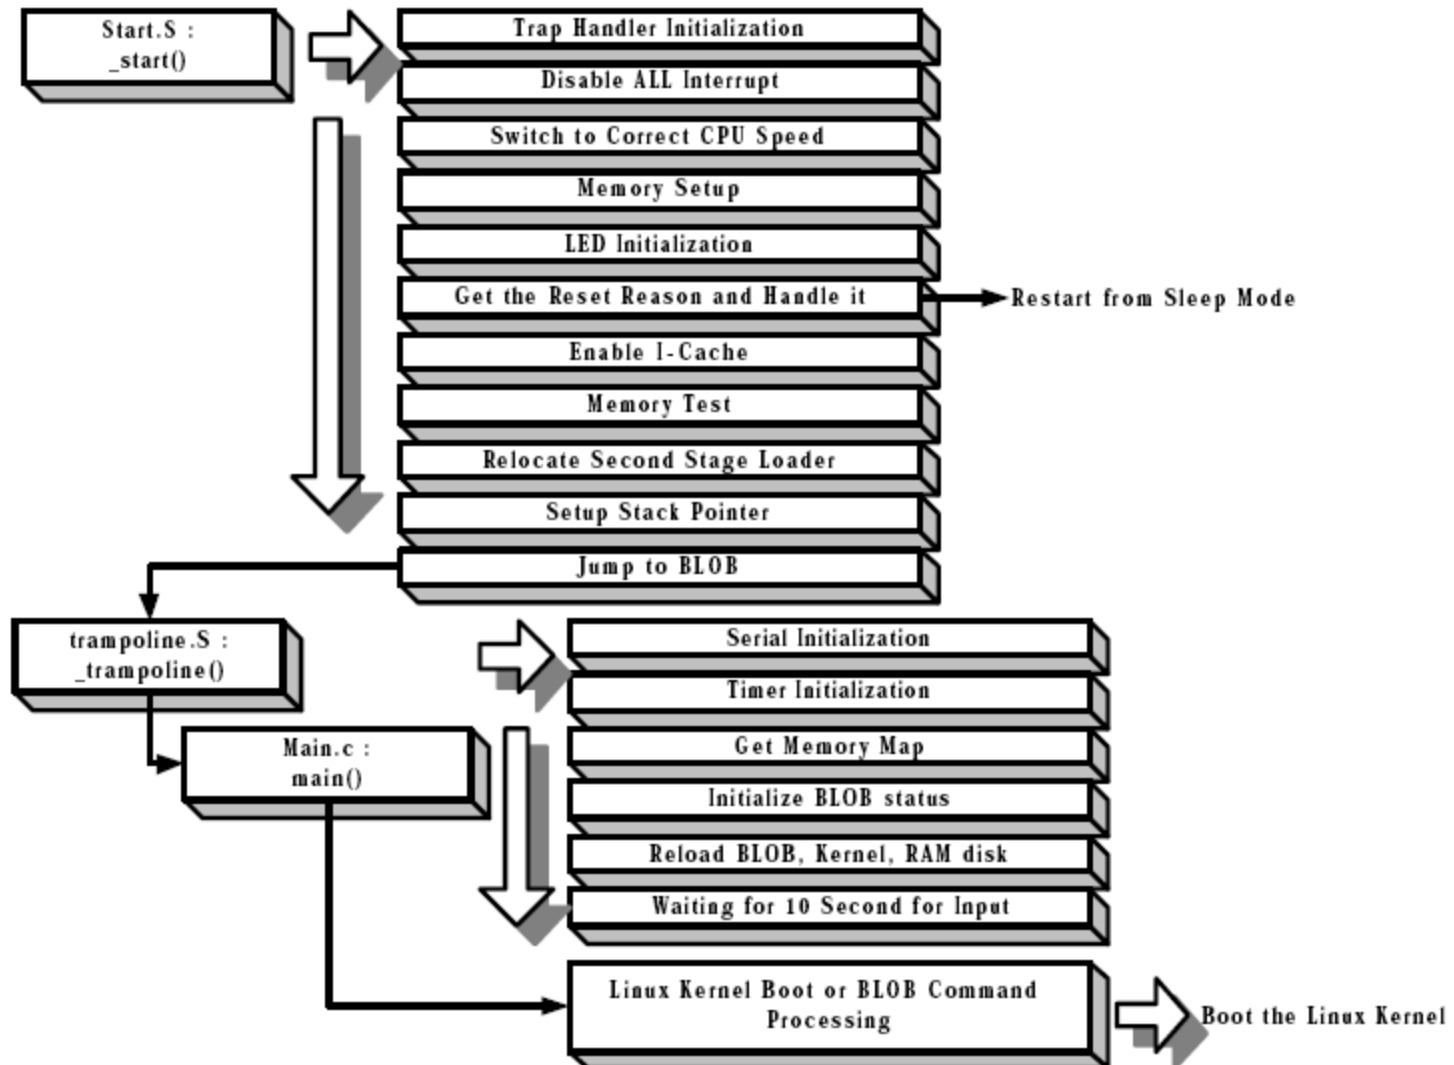
x86



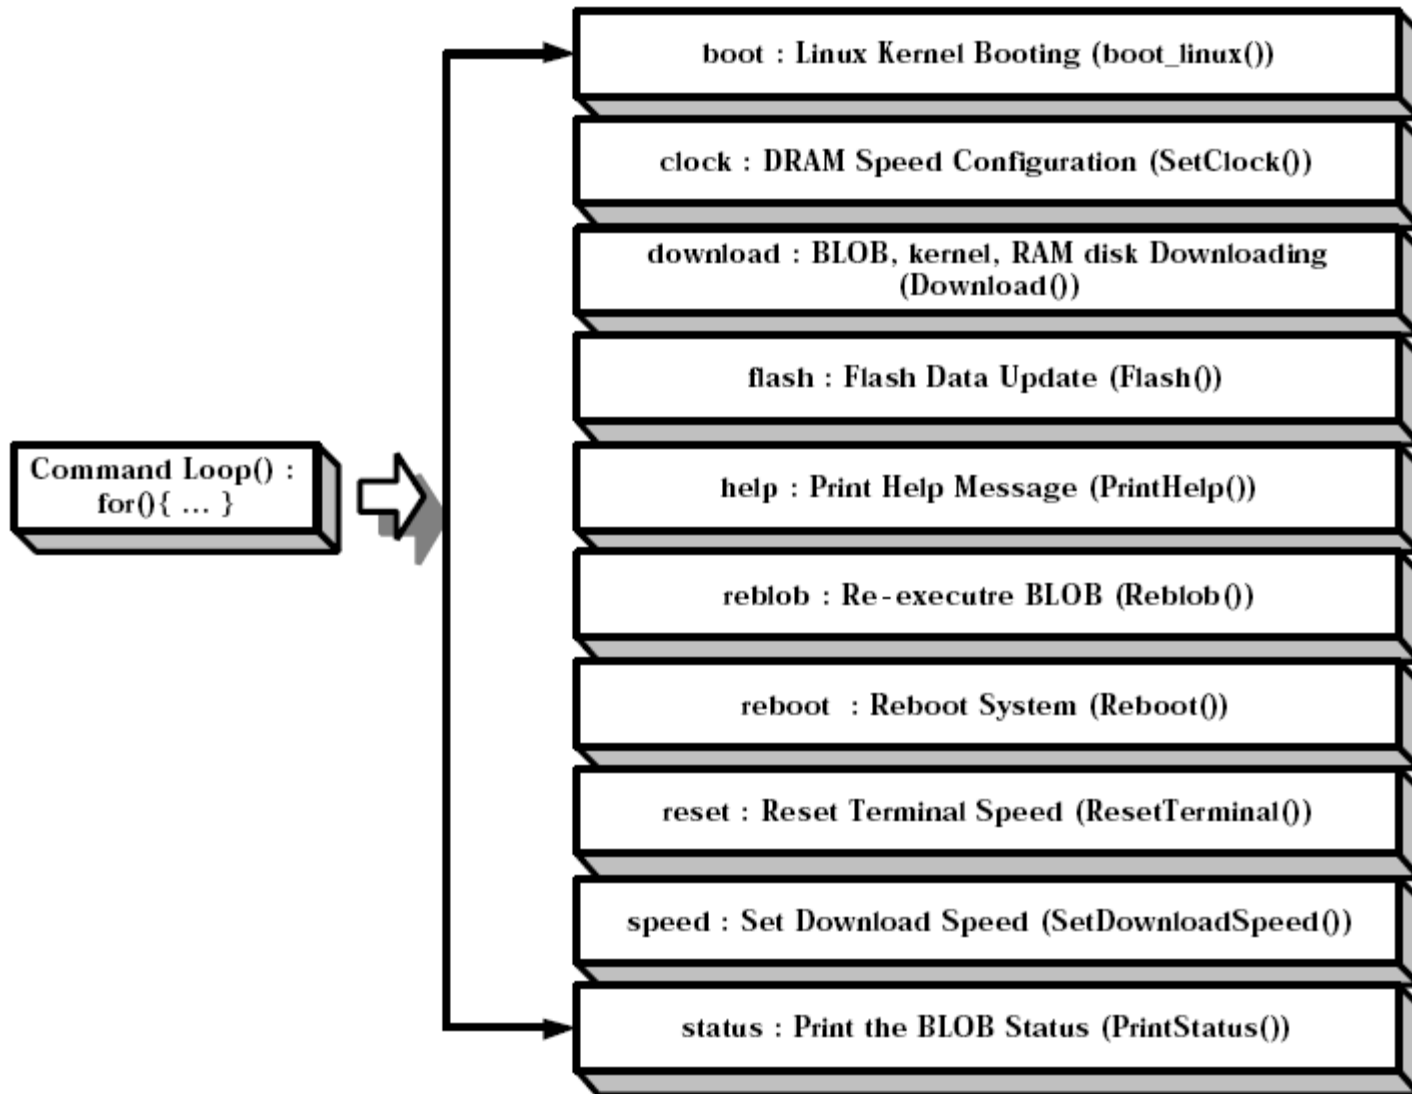
Boot-loader functions

- Initialize the hardware setting
- Basic monitor and debugger
- Pass the control to OS

Boot-loader Example



Boot-loader Example



ARM Examples

- Vector table

```
.text

/* Jump vector table as in table 3.1 in [1] */
.globl _start
_start:      b      reset
             b      undefined_instruction
             b      software_interrupt
             b      prefetch_abort
             b      data_abort
             b      not_used
             b      irq
             b      fiq
```

Reset_Handler

```
/* the actual reset code */
reset:

    /* First, mask **ALL** interrupts */
    ldr    r0, IC_BASE
    mov    r1, #0x00
    str    r1, [r0, #ICMR]

    /* switch CPU to correct speed */
    ldr    r0, PWR_BASE
    LDR    r1, cpuspeed
    str    r1, [r0, #PPCR]

    /* setup memory */
    bl     memsetup

    /* init LED */
    bl     ledinit
```


Reset_Handler (Cont.)

```
/* check if this is a wake-up from sleep */
ldr    r0, RST_BASE
ldr    r1, [r0, #RCSR]      Reset status register
and    r1, r1, #0x0f
teq    r1, #0x08
bne    normal_boot /* no, continue booting */

/* yes, a wake-up. clear RCSR by writing a 1 (see
9.6.2.1 from [1]) */
mov    r1, #0x08
str    r1, [r0, #RCSR]      ;

/* get the value from the PSPR and jump to it */
ldr    r0, PWR_BASE
ldr    r1, [r0, #PSPR]      Power manager scratch pad register
mov    pc, r1
```

Reset_Handler (Cont.)

normal_boot:

/* enable I-cache */

mrc p15, 0, r1, c1, c0, 0 @ read control reg

orr r1, r1, #0x1000 @ set Icache

mcr p15, 0, r1, c1, c0, 0 @ write it back

/* check the first 1MB in increments of 4k */

mov r7, #0x1000

mov r6, r7, lsl #8 /* 4k << 2^8 = 1MB */

ldr r5, MEM_START

mem_test_loop:

mov r0, r5

bl testram

teq r0, #1

beq badram

add r5, r5, r7

subs r6, r6, r7

bne mem_test_loop

```

        /* the first megabyte is OK, so let's clear it */
        mov     r0, #((1024 * 1024) / (8 * 4))    /* 1MB in
steps of 32 bytes */
        ldr     r1, MEM_START
...
clear_loop:
        stmia   r1!, {r2-r9}
        subs    r0, r0, #(8 * 4)
        bne     clear_loop

        /* relocate the second stage loader */
        add     r2, r0, #(128 * 1024)             /* blob is 128kB */
        add     r0, r0, #0x400                    /* skip first 1024
bytes */
        ldr     r1, MEM_START
        add     r1, r1, #0x400                    /* skip over here
as well */
        ..
copy_loop:
        ldmia   r0!, {r3-r10}
        stmia   r1!, {r3-r10}
        cmp     r0, r2
        ble     copy_loop

```

Reset_Handler (Cont.)

```
/* set up the stack pointer */
```

```
ldr    r0, MEM_START
```

```
add    r1, r0, #(1024 * 1024)
```

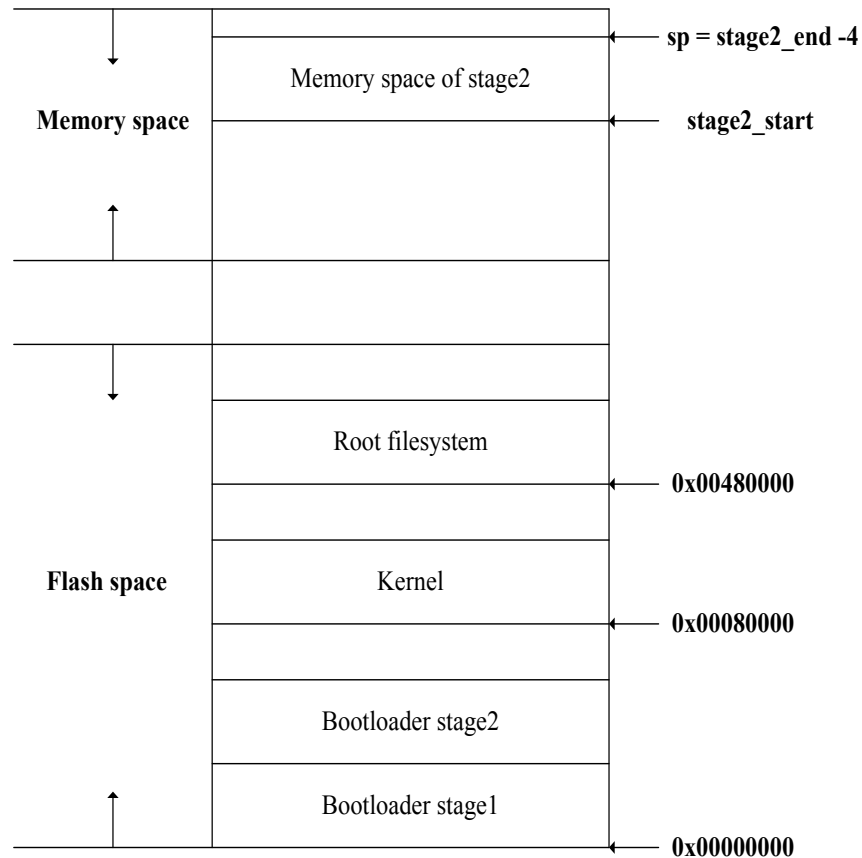
```
sub    sp, r1, #0x04
```

```
/* blob is copied to ram, so jump to it */
```

```
add    r0, r0, #0x400
```

```
mov    pc, r0
```

Bootloader memory map



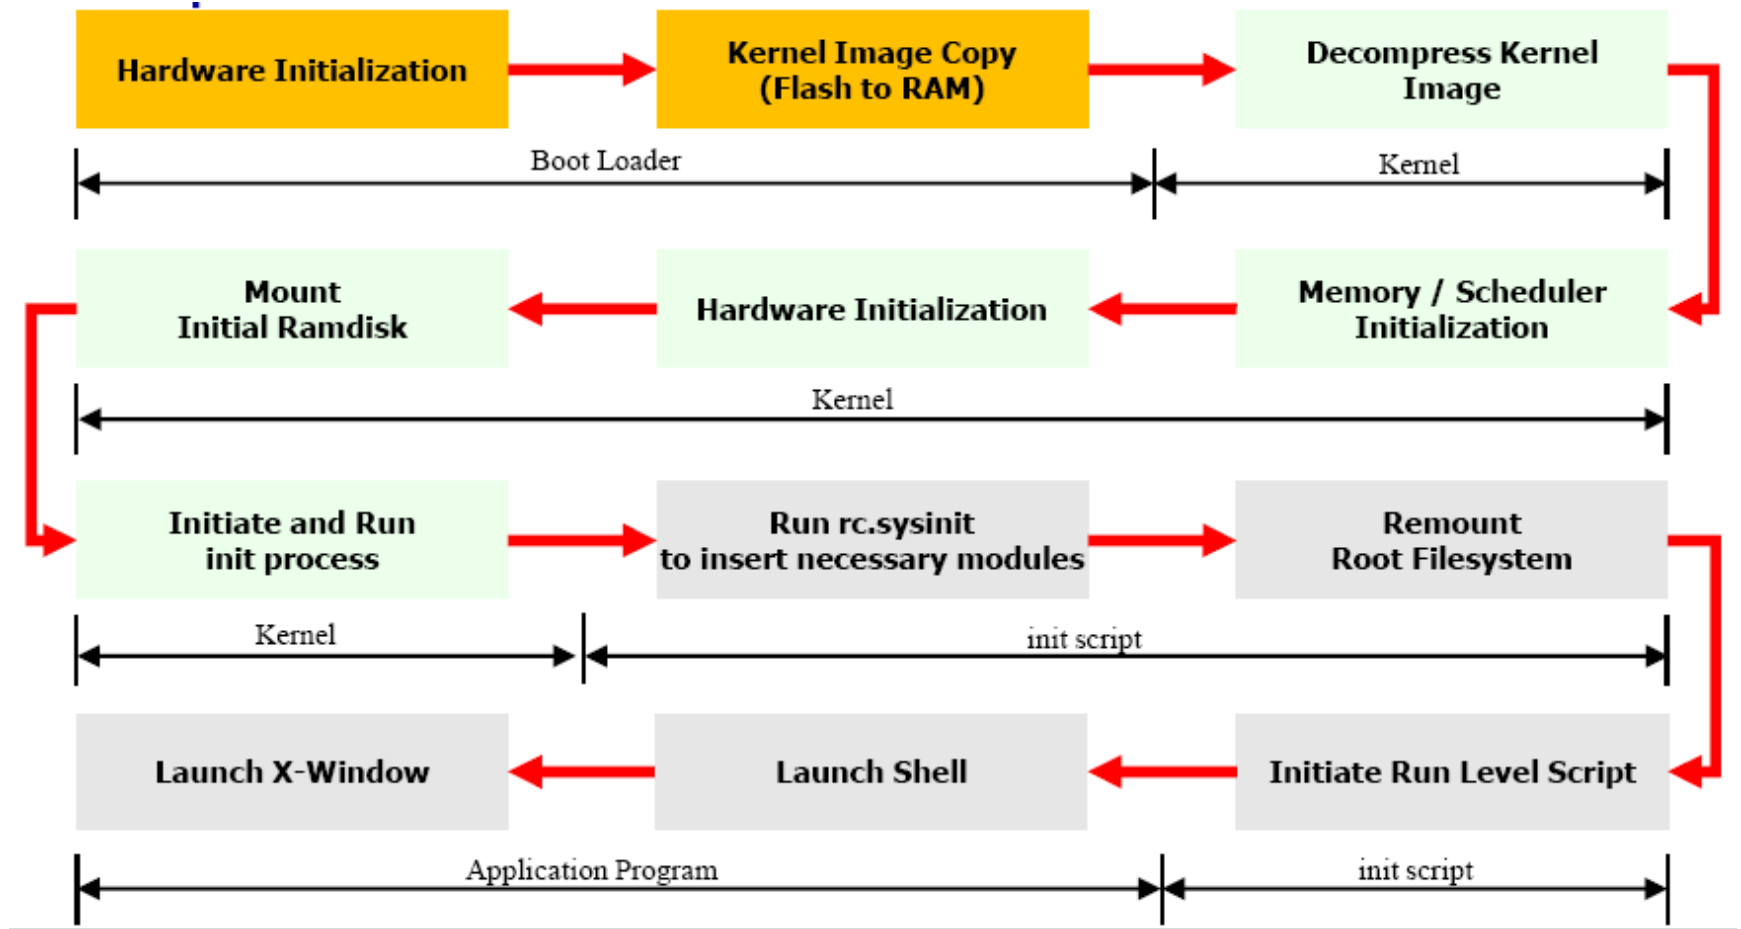
Boot-loader C program

```
int main(void)
{
~
    led_on();
~
    SerialInit(baud9k6);
    TimerInit();
~
    SerialOutputString(PACKAGE " version " VERSION "\n"
                      "Copyright (C) 1999 2000 2001 ")
~
    get_memory_map();
~
    SerialOutputString("Running from ");
    if(RunningFromInternal())
        SerialOutputString("internal");
    else
        SerialOutputString("external");
...
}
```

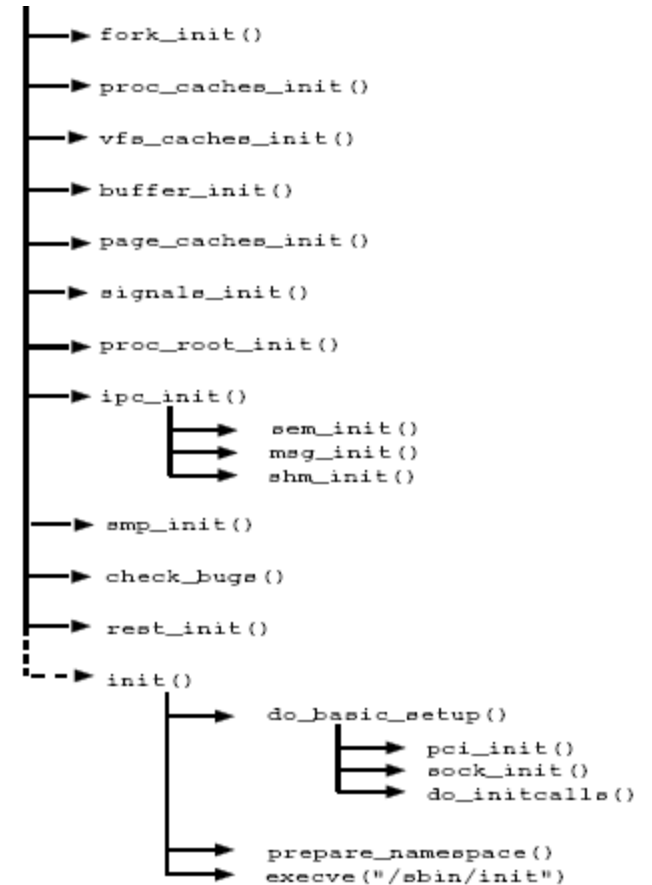
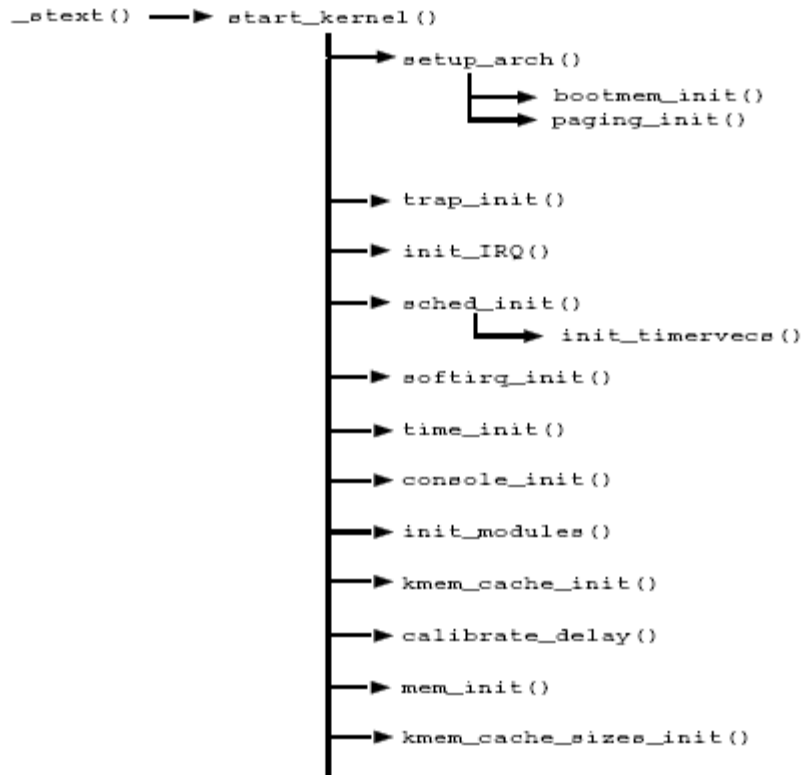
Boot-loader C program

```
/* wait 10 seconds before starting autoboot */
SerialOutputString("Autoboot in progress, press any
key...");
for(i = 0; i < 10; i++) {
    SerialOutputByte('.');
    retval = SerialInputBlock(commandline, 1, 1);
~
if(retval == 0) {
~
    boot_linux(commandline);    }
~
for(;;) {
~
    if(numRead > 0) {
        if(MyStrNCmp(commandline, "boot", 4) == 0) {
            boot_linux(commandline + 4);
        } else if(MyStrNCmp(commandline, "clock", 5) == 0) {
            SetClock(commandline + 5);
        } else if(MyStrNCmp(commandline, "download ", 9) == 0) {
~
        return 0;
} /* main */
```

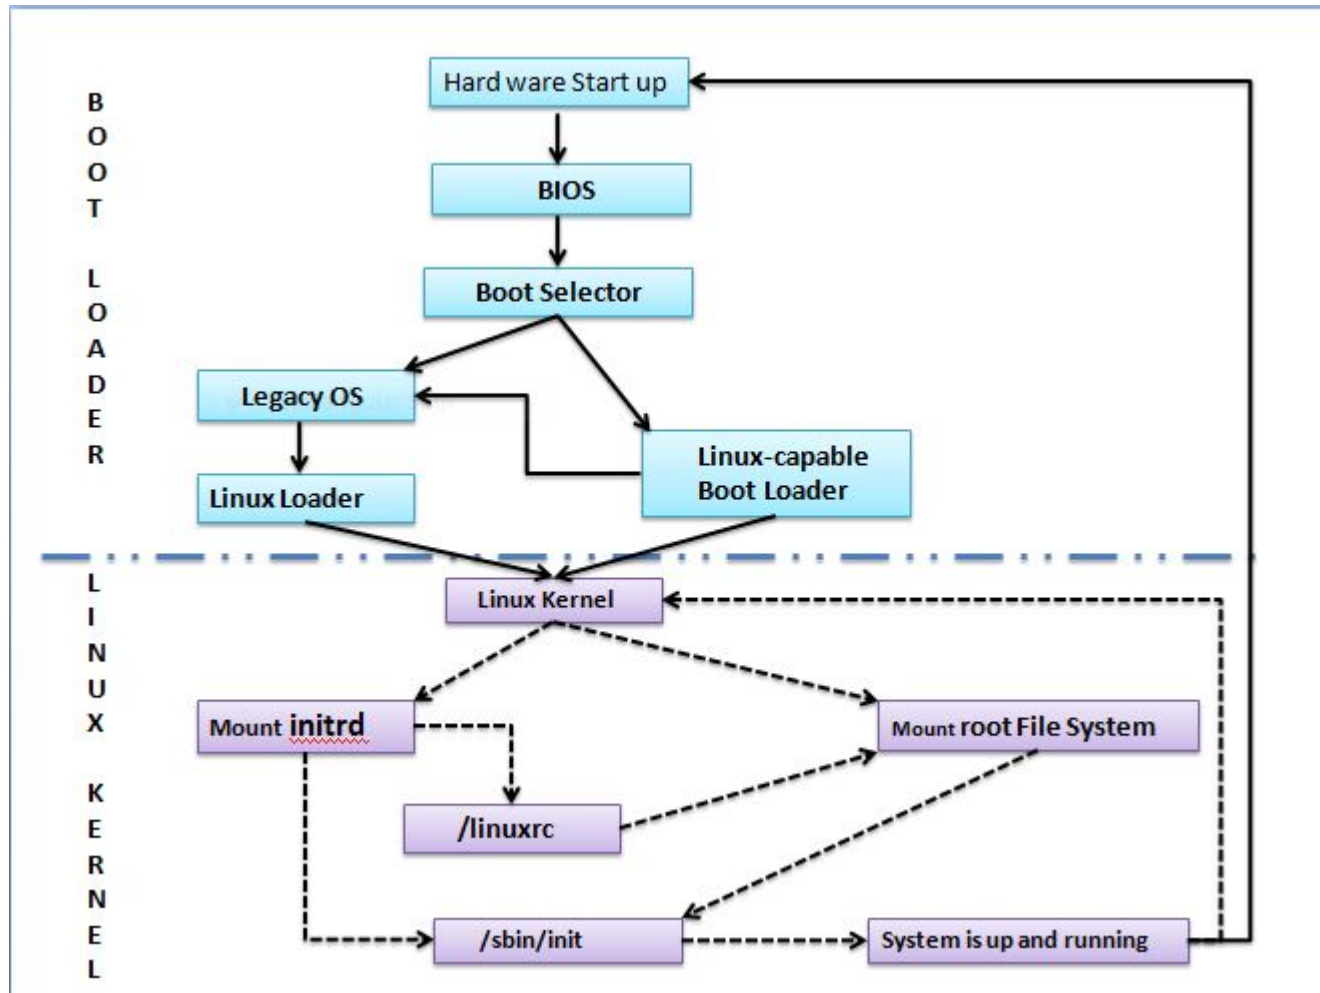
How does OS boot?



Linux booting



Mount root disk

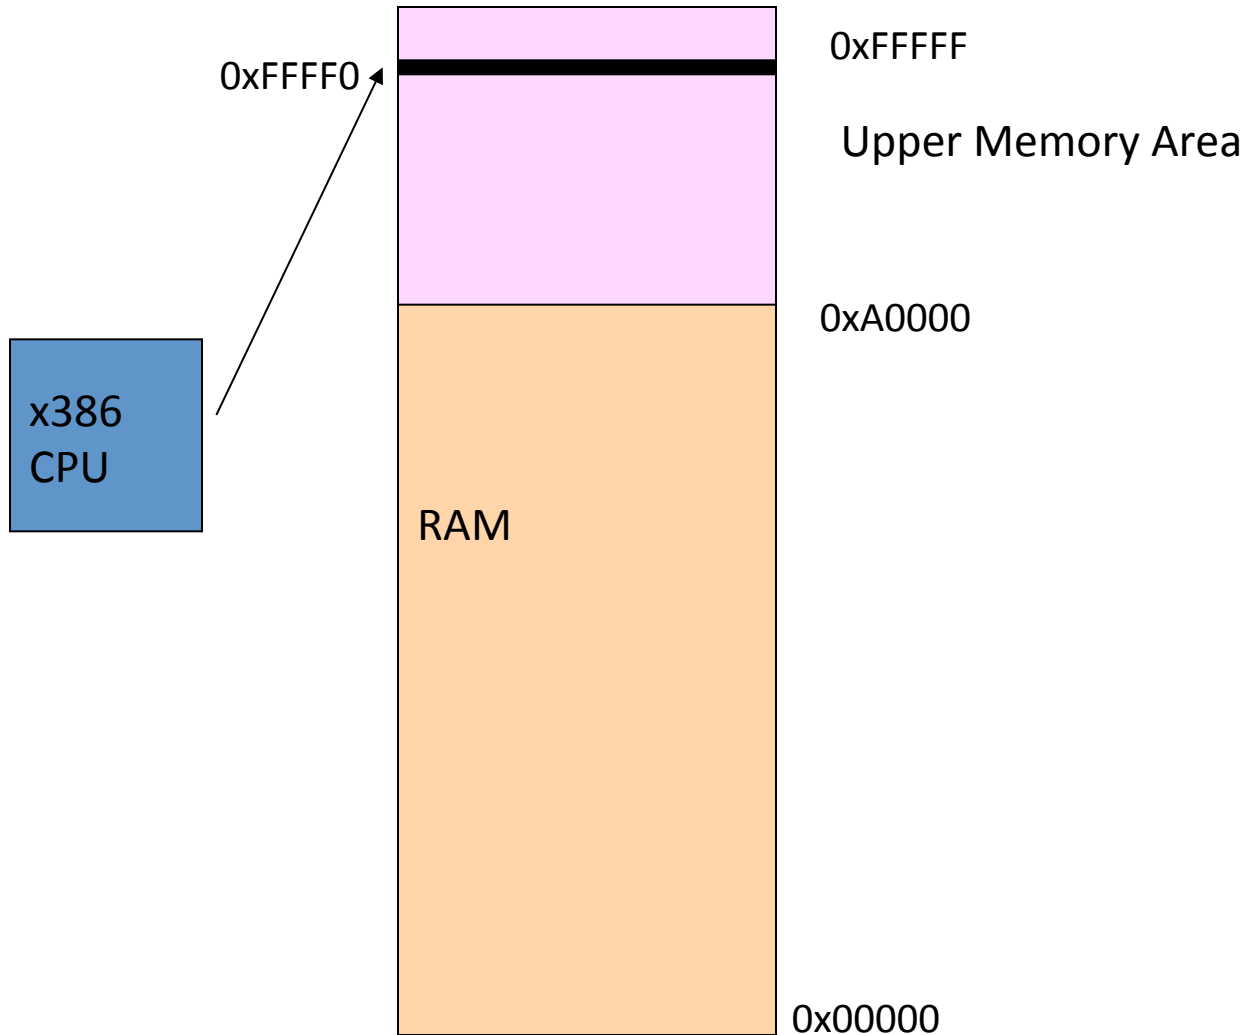


initrd (initramfs): boot loader initialized RAM disk

A case study of PC and Linux 0.11

- Basic Input and Output System
- Firmware
- For Intel boot architecture
 - Old but backward compatible
- Functions include
 - H/W initialization and parameter configurations
 - Device drivers
 - Debuggers/monitor
 - Load OS/another bootloader

PC Booting (Cont)

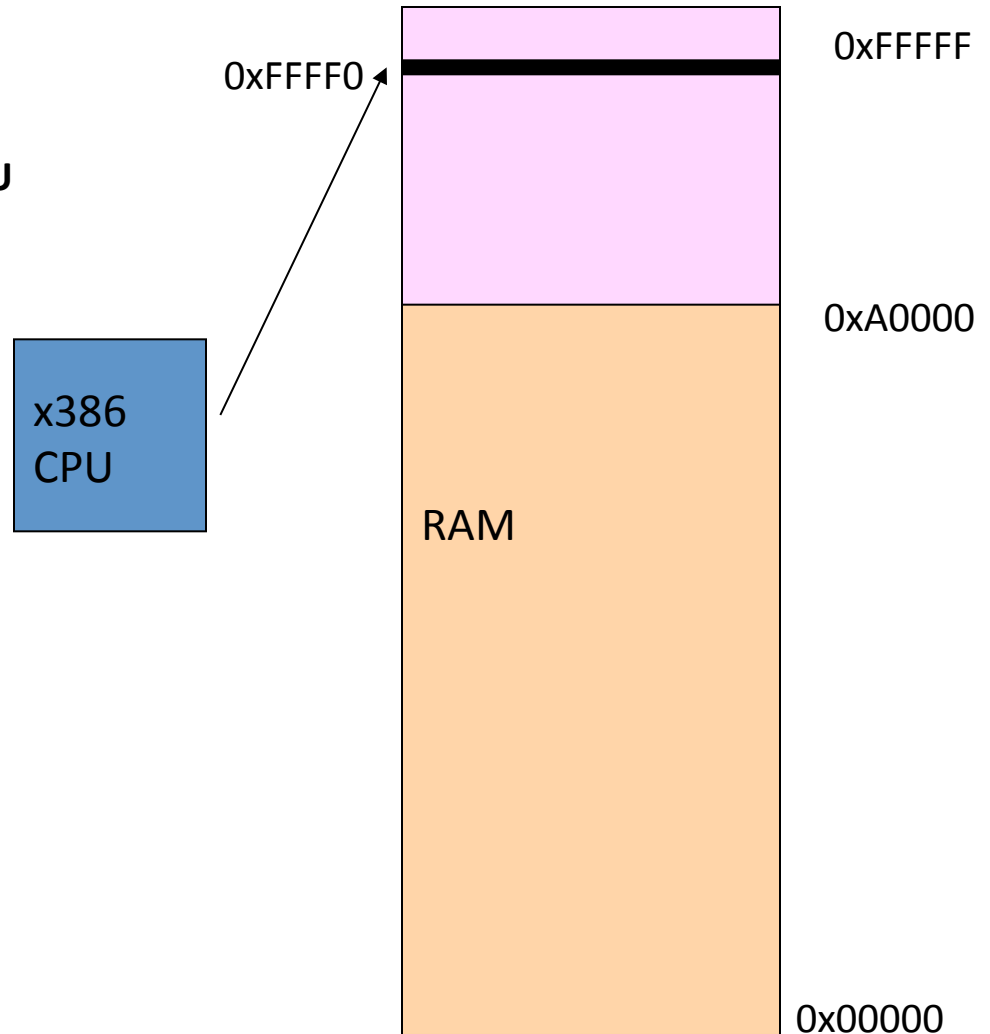


UMA

Address	First 16K (x0000h- x3FFFh)	Second 16K (x4000h- x7FFFh)	Third 16K (x8000h- xBFFFh)	Fourth 16K (xC000h- xFFFFh)
A0000-AFFFFh	VGA Graphics Mode Video			
B0000- BFFFFh	VGA Monochrome Text Mode Video RAM		VGA Color Text Mode Video RAM	
C0000- CFFFFh	VGA Video BIOS ROM		IDE Hard Disk BIOS ROM	Optional Adapter ROM BIOS or RAM UMBs
D0000- DFFFFh	Optional Adapter ROM BIOS or RAM UMBS			
F0000- FFFFFh	System BIOS ROM			

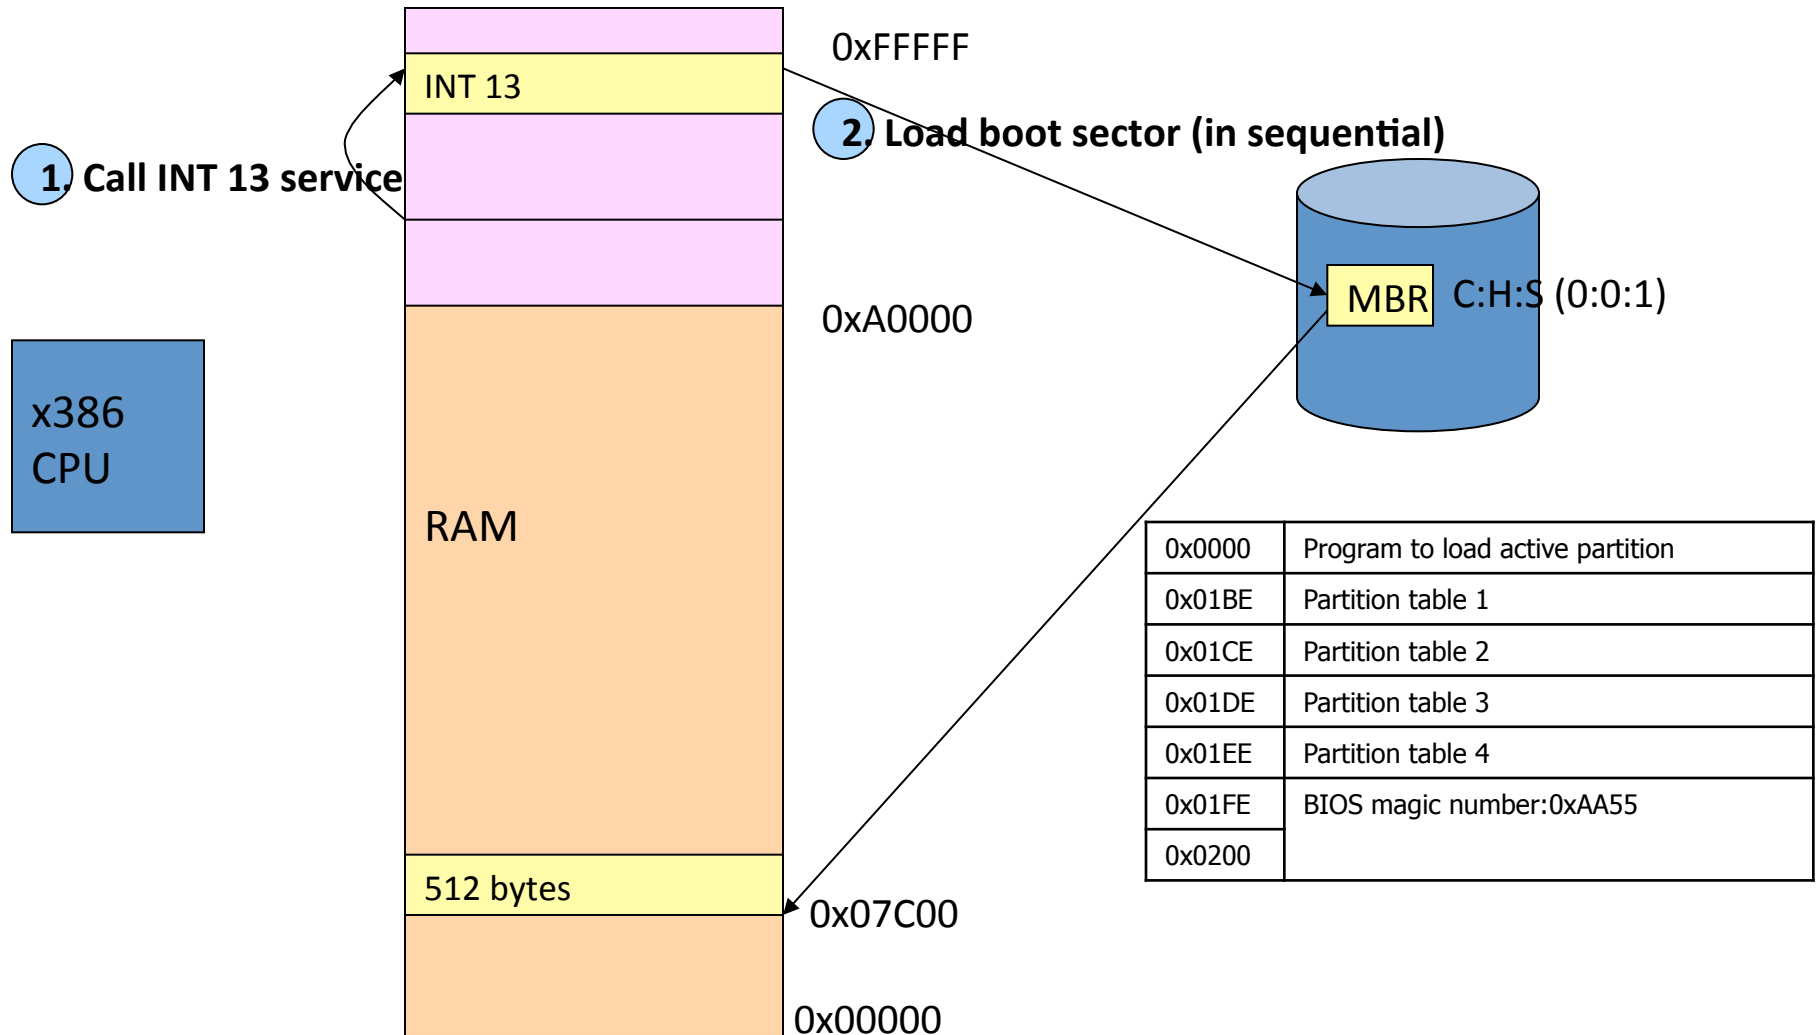
PC Booting (Cont)

1. Power supply sends POWER GOOD to CPU
2. CPU resets
3. Run FFFF:0000 @ BIOS ROM
4. Jump to a real BIOS start address
5. POST
6. Beep if there is an error
7. Read CMOS data/settings
8. Run 2nd-stage boot

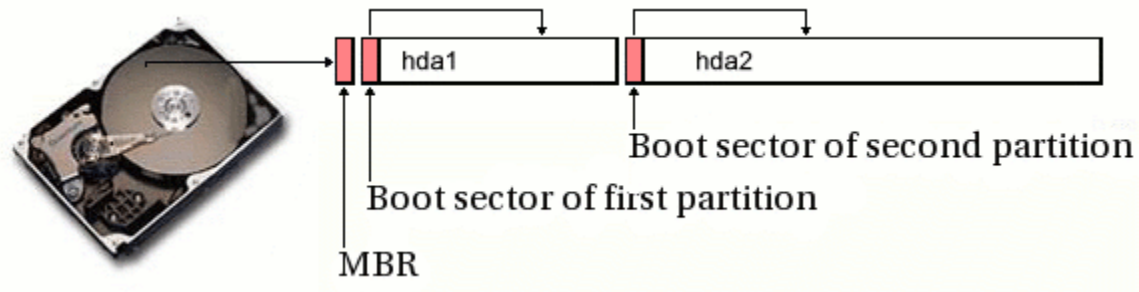


INT	Address	Type	Description
00h	0000:0000h	Processor	Divide by zero
01h	0000:0004h	Processor	Single step
02h	0000:0008h	Processor	Non maskable interrupt (NMI)
03h	0000:000Ch	Processor	Breakpoint
04h	0000:0010h	Processor	Arithmetic overflow
05h	0000:0014h	Software	Print screen
06h	0000:0018h	Processor	Invalid op code
07h	0000:001Ch	Processor	Coprocessor not available
08h	0000:0020h	Hardware	System timer service routine
09h	0000:0024h	Hardware	Keyboard device service routine
0Ah	0000:0028h	Hardware	Cascade from 2nd programmable interrupt controller
0Bh	0000:002Ch	Hardware	Serial port service - COM post 2
0Ch	0000:0030h	Hardware	Serial port service - COM port 1
0Dh	0000:0034h	Hardware	Parallel printer service - LPT 2
0Eh	0000:0038h	Hardware	Floppy disk service
0Fh	0000:003Ch	Hardware	Parallel printer service - LPT 1
10h	0000:0040h	Software	Video service routine
11h	0000:0044h	Software	Equipment list service routine
12h	0000:0048H	Software	Memory size service routine
13h	0000:004Ch	Software	Hard disk drive service
14h	0000:0050h	Software	Serial communications service routines
15h	0000:0054h	Software	System services support routines
16h	0000:0058h	Software	Keyboard support service routines
17h	0000:005Ch	Software	Parallel printer support services

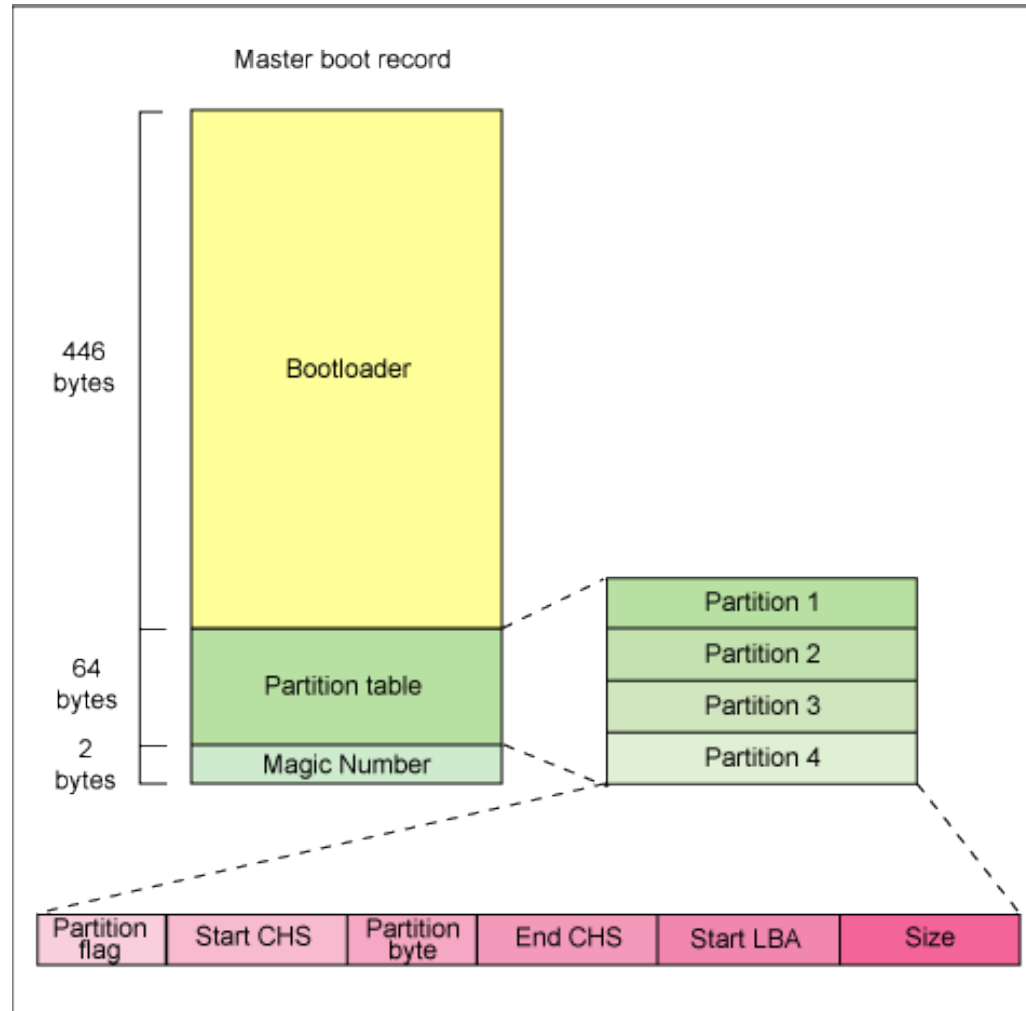
PC Booting (Cont)



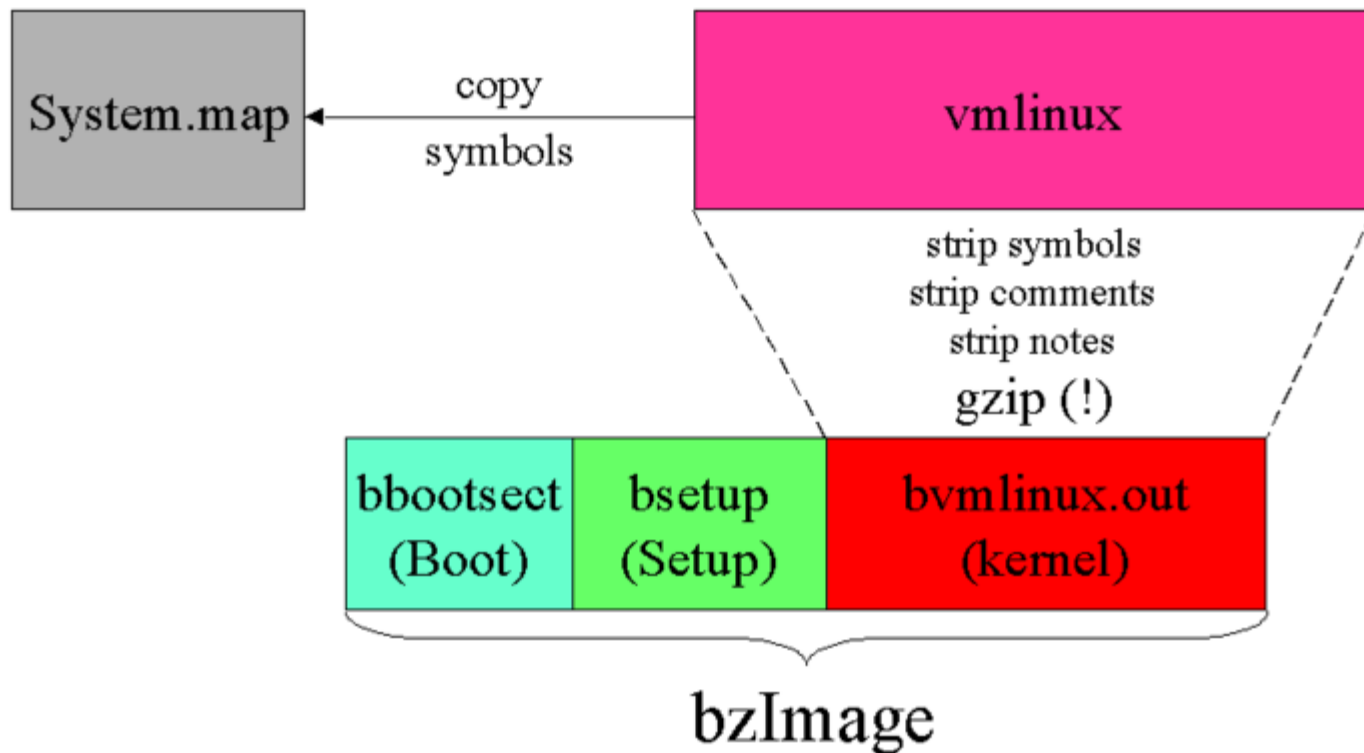
MBR



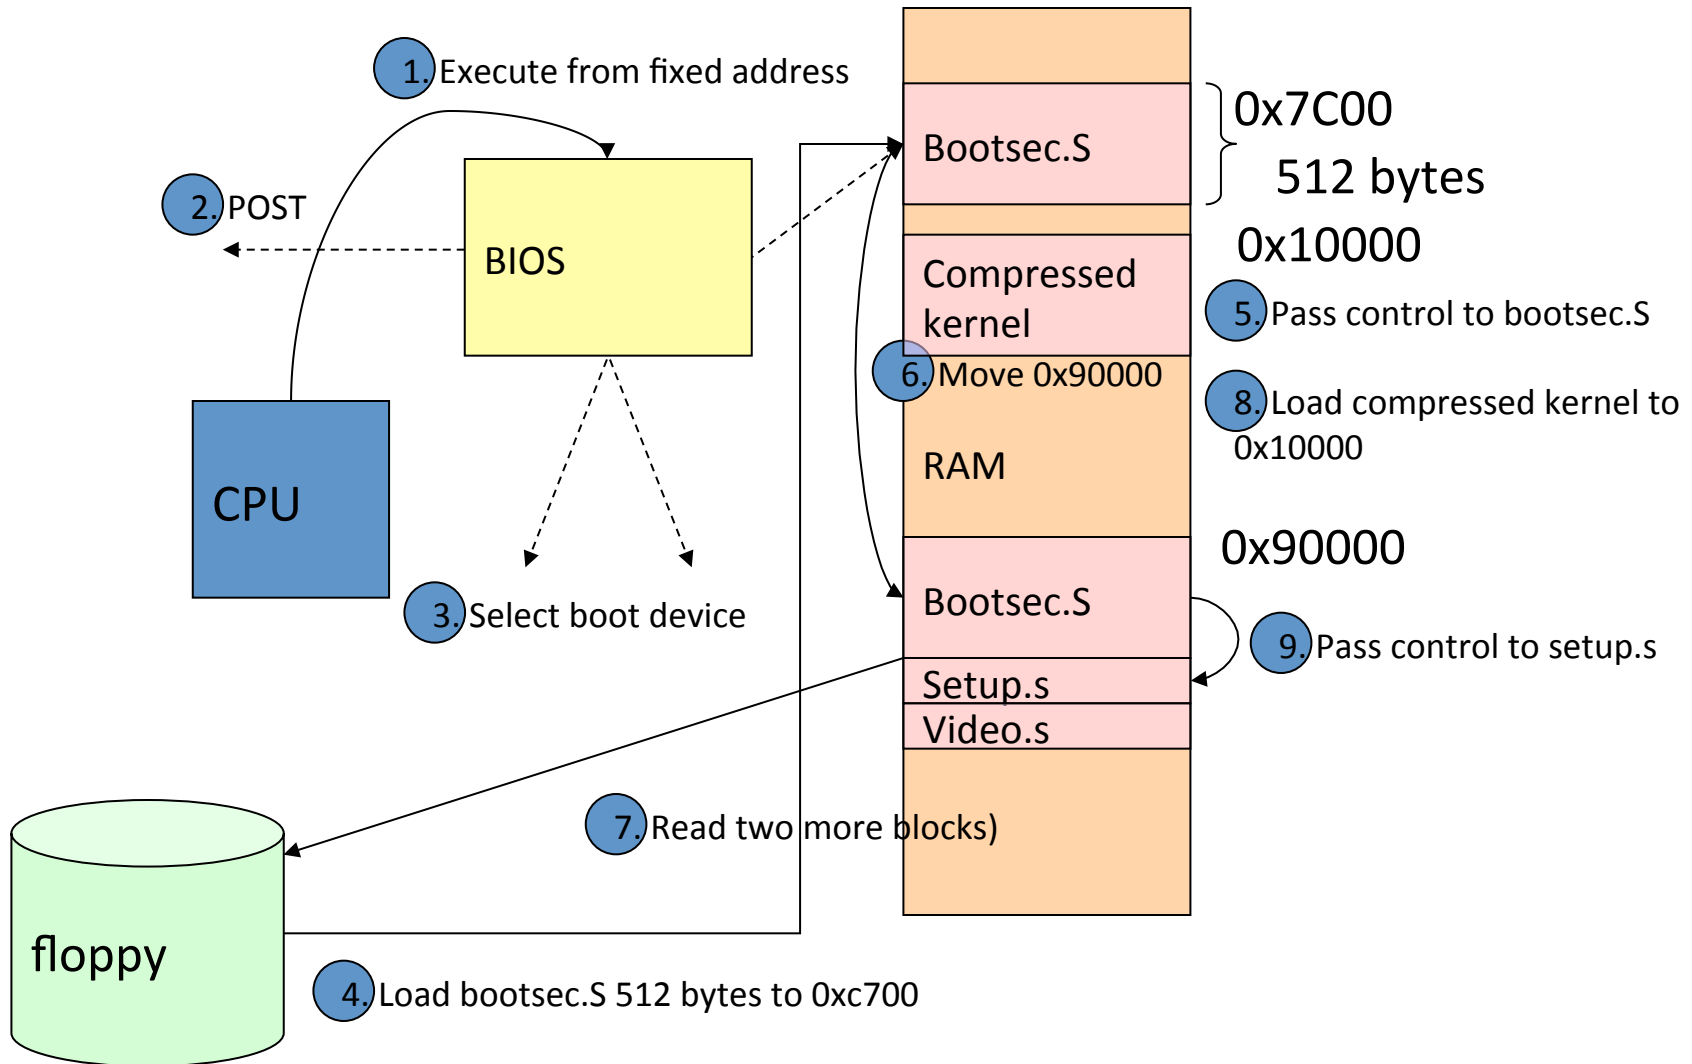
MBR (Master Boot Record)



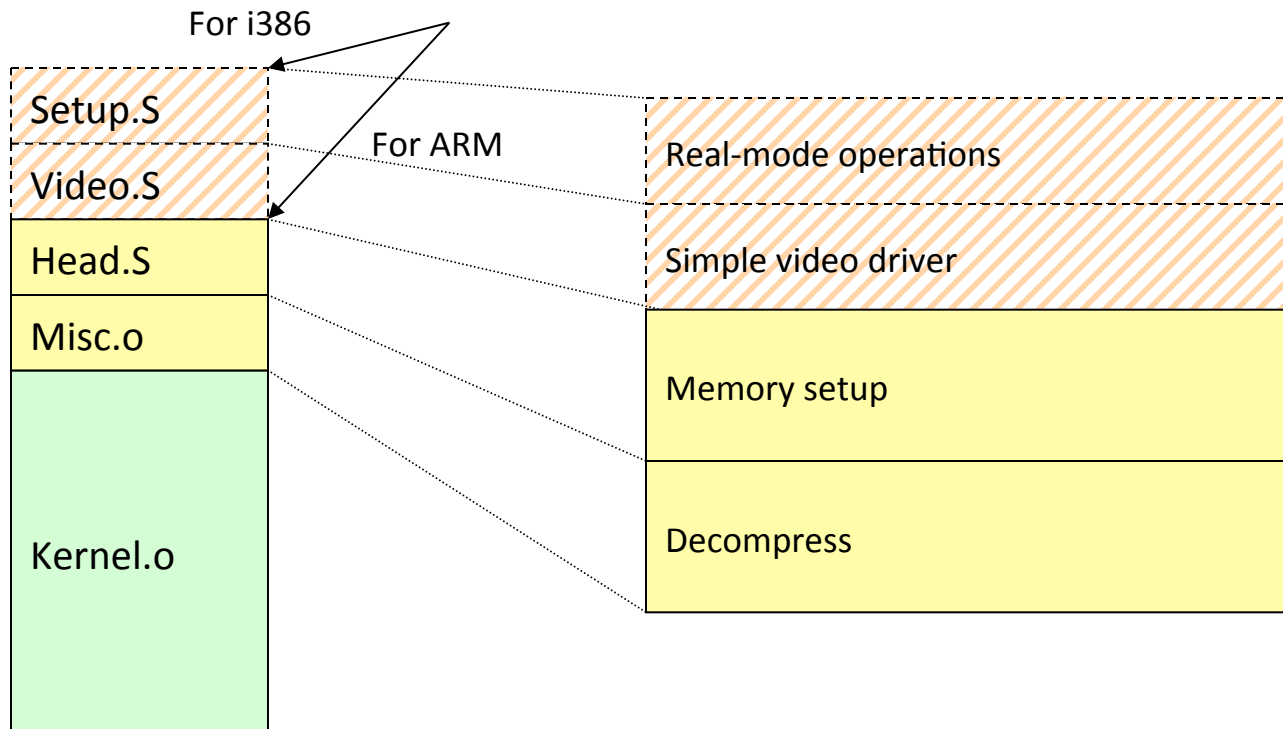
Anatomy of bzImage



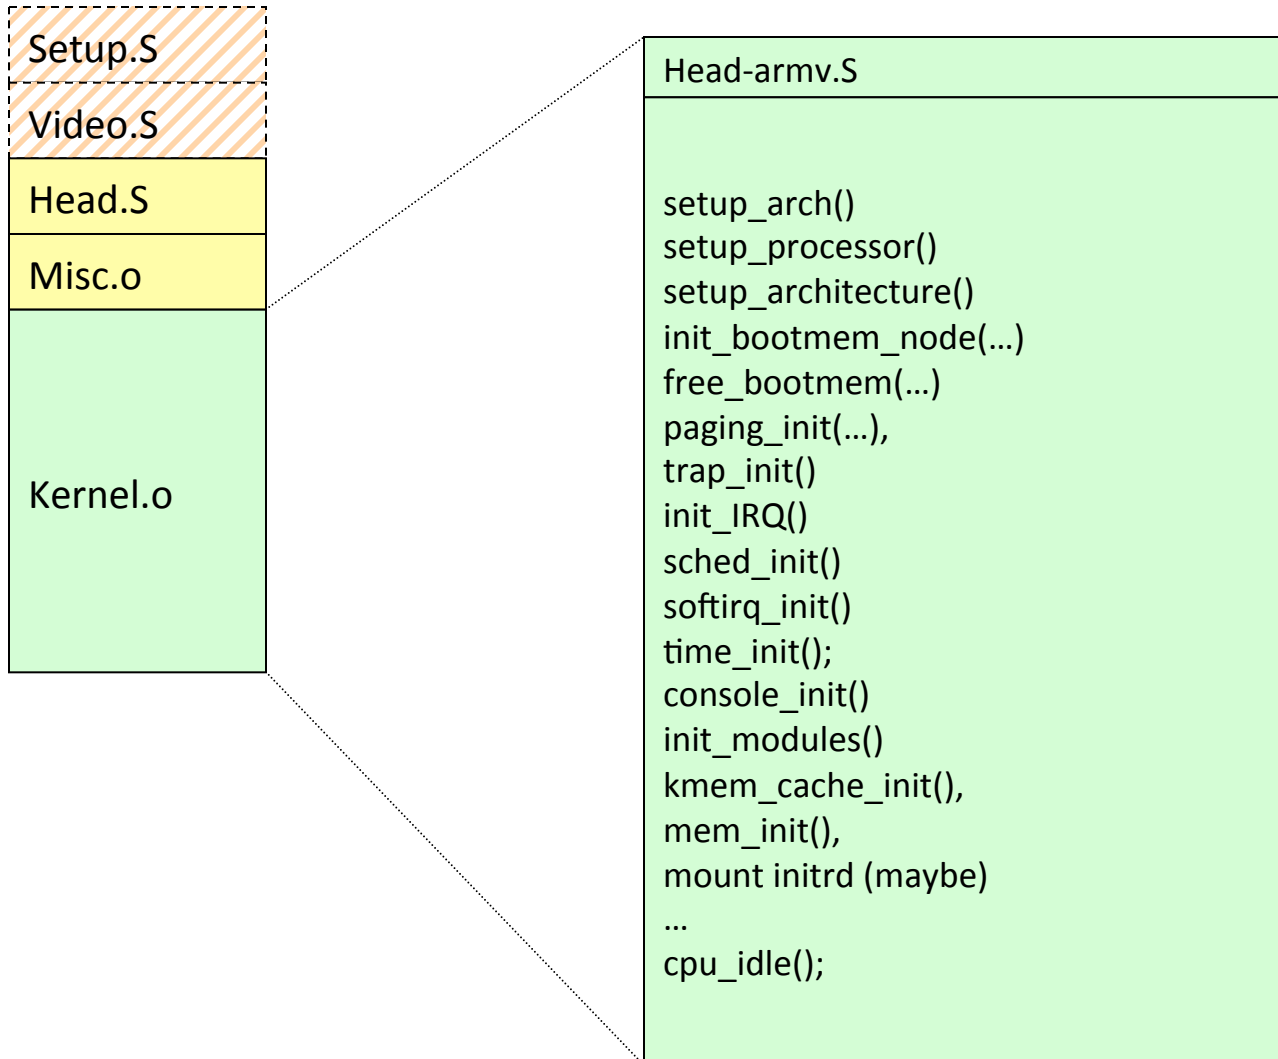
Linux Boot Example



Kernel Image Structure



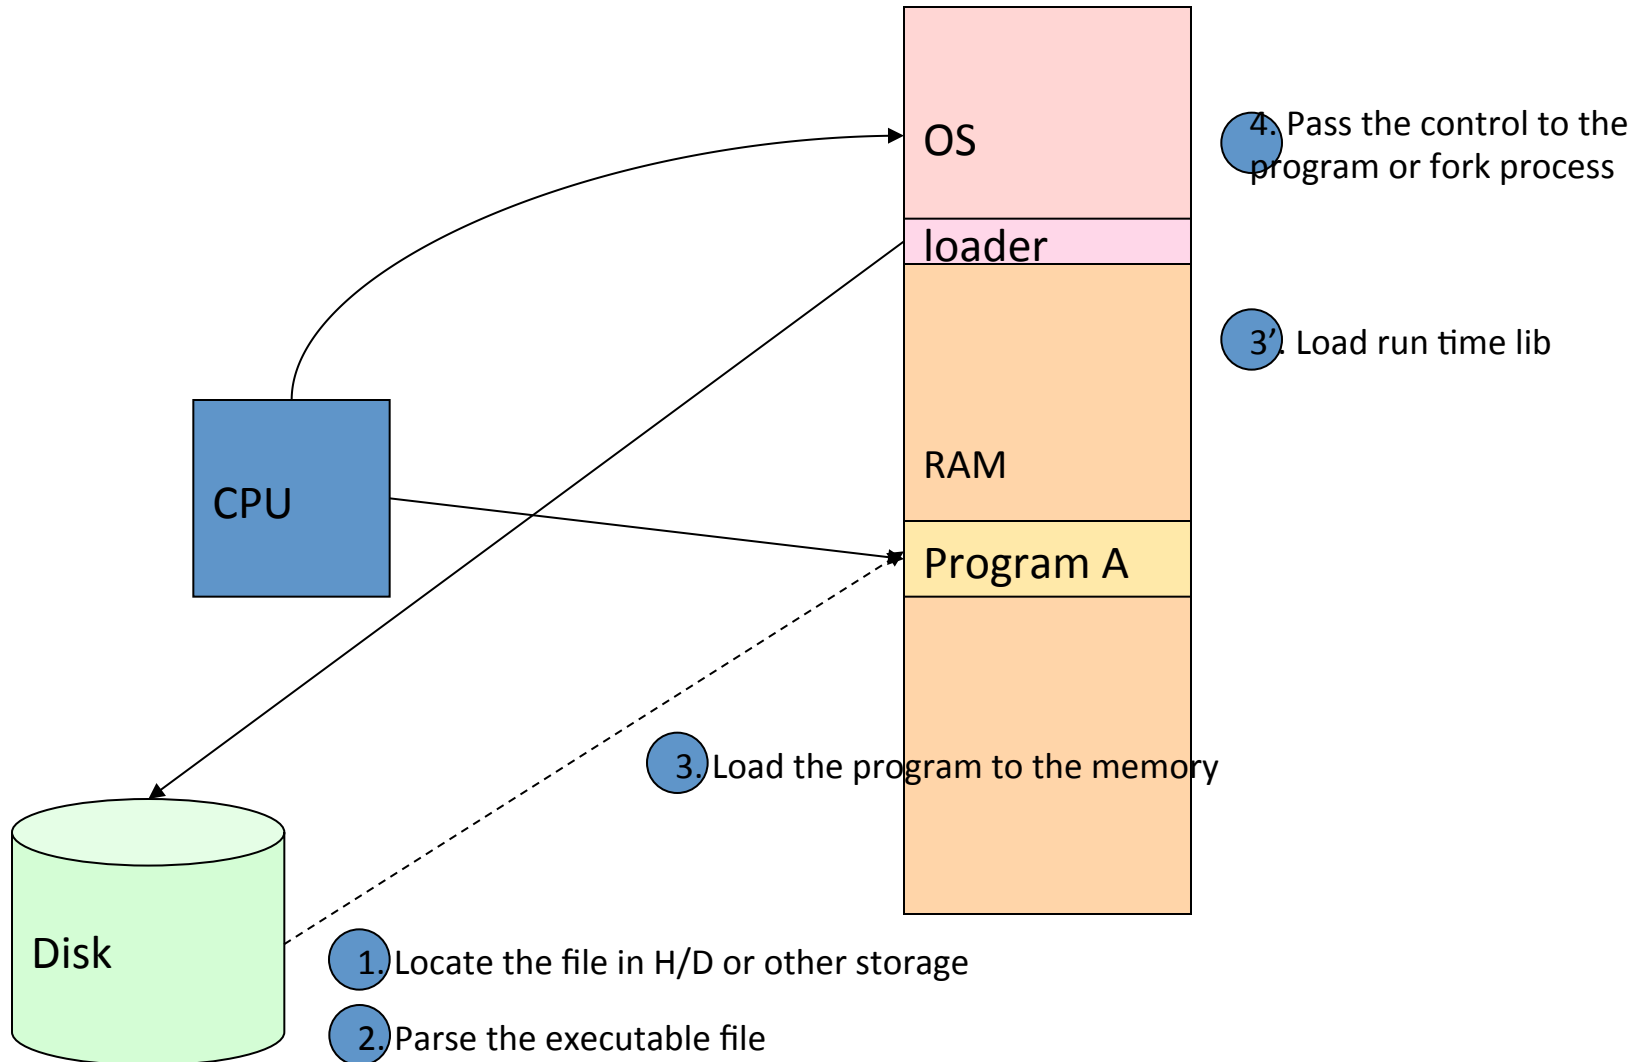
Kernel Image Structure



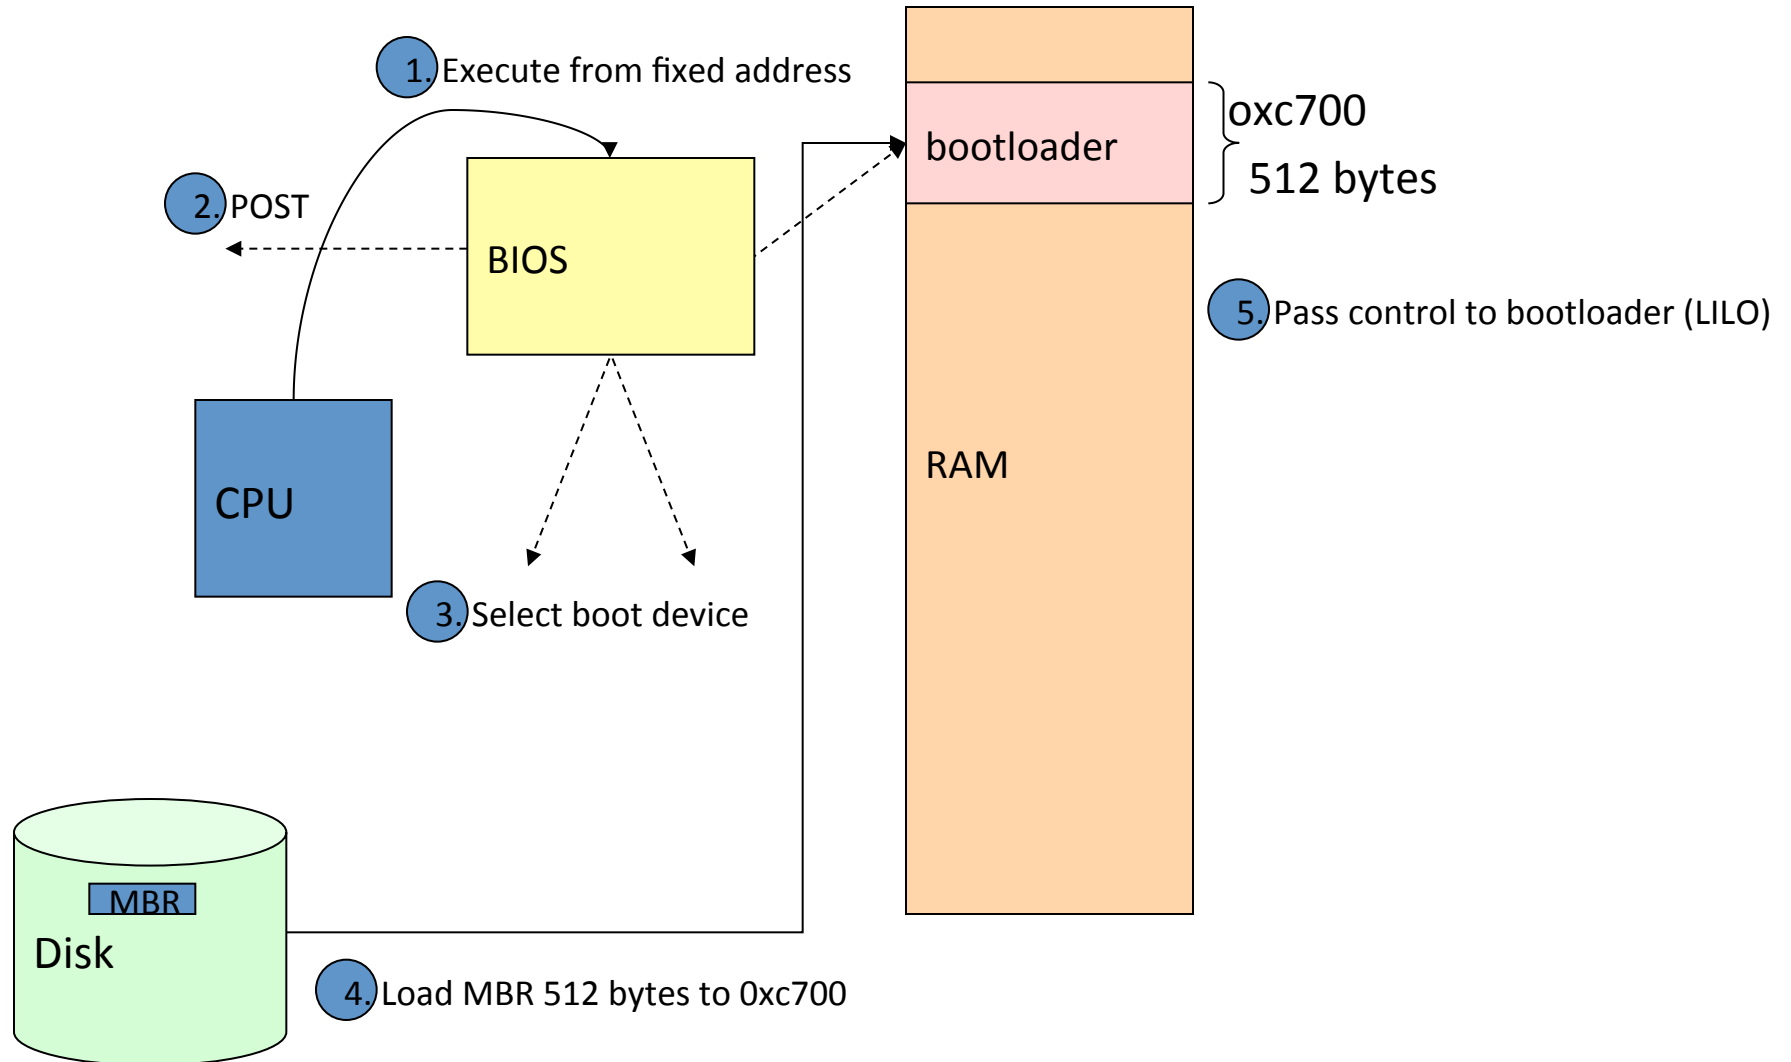
Loader

- Normally a user level program
- Might use OS system services to access the storage and control the memory allocation
- Load and interpret executable files to memory and ask OS to run it

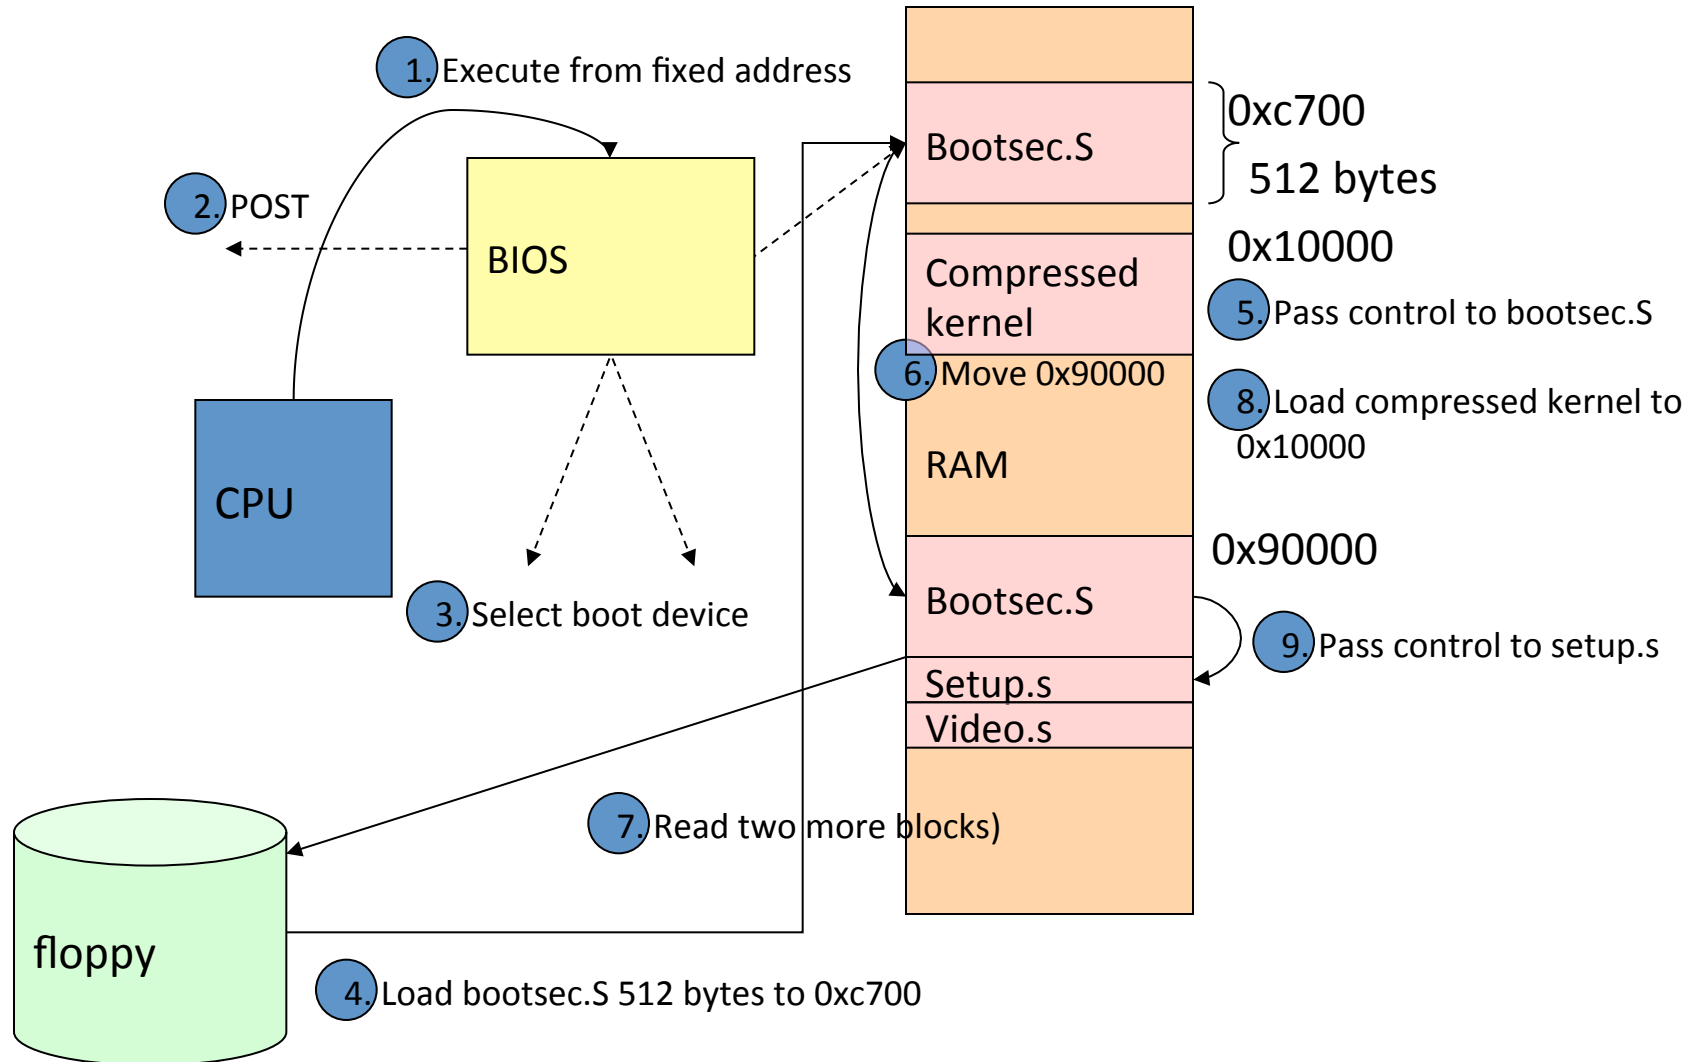
Loader



Linux Boot Example



Linux Boot Example



Bootloader Speedups

- Remove waiting time
- Removing unnecessary initialization routines
- Uncompressed kernel
- DMA Copy Of Kernel On Startup
- Fast Kernel Decompression
- Kernel XIP

Remove waiting time

- Boot loader await a couple of seconds to connect debug port
- Since boot loader cannot detect connection of debug port, boot loader awaits a couple of seconds on normal boot sequence.
- U-boot : `setenv bootdelay 0`

Removing unnecessary initialization routines

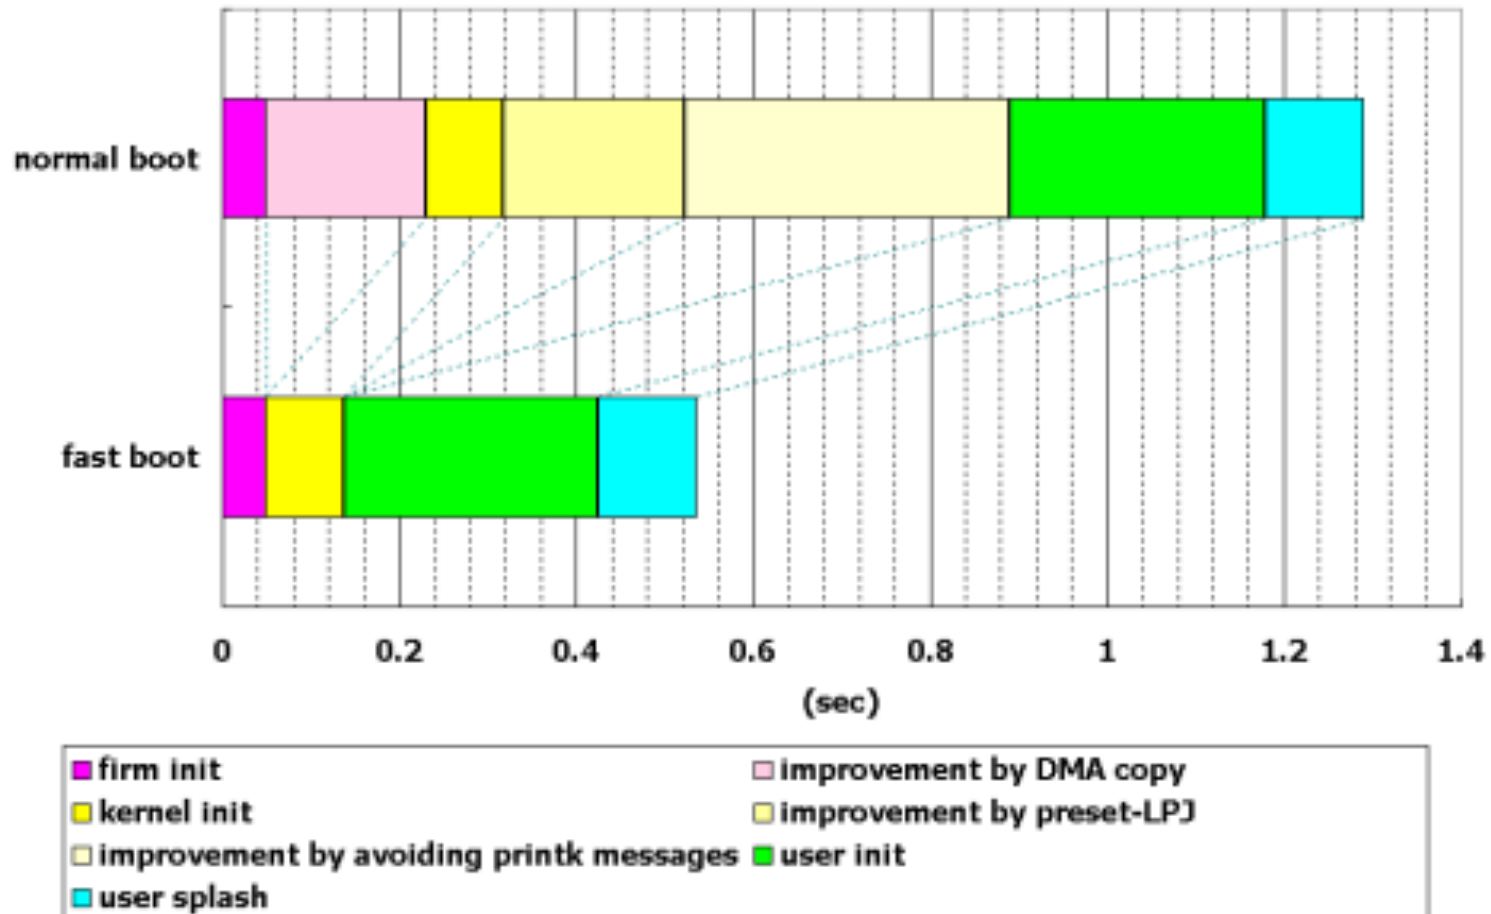
- Should not perform initialize other devices except for necessities to load kernel
- Examples
 - Removing initialization of LCD
 - Removing initialization of timer

Uncompressed kernel

- Fast boot, but image size has been larger
 - 2MB – 2.5MB non-compressed (ARM)
 - 1MB – 1.5MB compressed (ARM)
- It should be different results, performance of CPUs, speed of flash memory
 - Profiling is required
- Make Image vs. make zImage

DMA Copy Of Kernel On Startup

- Turn on DMA in bootloader



Fast Kernel Decompression

- Use other compression/decompression algorithm
 - Slow compression, good compression ratio, fast decompression
- GZIP vs. Sony UCL
- Small kernel size, fast kernel loading time

Fast Kernel Decompression

Image file:	initrd-2.6.5-1.358		Power PC
method	UCL	GZIP	improved %
parameter	-b4194304	-8	.
source file size	819200	819200	.
compressed size	187853	189447	.
compression rate	77.1%	76.9%	0.3%
compression time: user (sec)	5.13	2.03	-152.5%
sys (sec)	0.09	0.06	-36.5%
total (sec)	5.22	2.09	-149.0%
decompression time: user (sec)	0.12	0.3	59.7%
sys (sec)	0.1	0.08	-16.9%
total (sec)	0.22	0.39	43.0%

Fast Kernel Decompression

Image file:	vmlinux-2.4.20 for ibm-440gp		PowerPC
method	UCL	GZIP	improved %
parameter	-b4194304	-8	.
source file size	1810351	1810351	.
compressed size	790250	776807	.
compression rate	56.3%	57.1%	-1.3%
compression time: user (sec)	17.29	6.07	-185.0%
sys (sec)	0.04	0.02	-92.4%
total (sec)	17.33	6.09	-184.6%
decompression time: user (sec)	0.12	0.16	26.1%
sys (sec)	0.03	0.04	35.8%
total (sec)	0.15	0.2	28.2%

Kernel XIP

- Direct addressing on NOR flash memory
- Cannot compress kernel
- PowerPC 405LP/266 MHZ

Boot Stage	Non-XIP Time	XIP Time
Copy kernel to RAM	85 ms	12 ms *
Decompress kernel	453 ms	0 ms
Kernel time to initialize (time to first user space program)	819 ms	882 ms
Total kernel boot time	1357 ms	894 ms
Reduction:	*	463 ms

Kernel XIP (Cont.)

- TI OMAP 5912/196 MHZ

Boot Stage	Non-XIP Time Kernel compressed	Non-XIP Time Kernel not compressed	XIP Time
Copy kernel to RAM	56 ms	120 ms	0 ms
Decompress kernel	545 ms	0 ms	0 ms
Kernel time to initialize (time to first user space program)	88 ms	208 ms	110 ms
Total kernel boot time	689 ms	208 ms	110 ms
Reduction:	*	481 ms	579 ms

Embedded Linux Speedups

- Removing unnecessary message printout
- Remove unnecessary functions and device drivers
- Modularization of device driver
- Asynchronous function calls
- Avoid performance measurement routine
- RTC no sync
- Using read-only file system
- Using lazy mount technique on R/W file systems
- Deferred Initcalls

Removing unnecessary message printout

- Add “quiet” to kernel command line
- Example 1

Hardware	KMC SH board, using VGA console
Kernel Version	CELF-1 (040126)
Configuration	relatively small configuration (details not available)
Time without "quiet" option	637878 usec
Time with "quiet" option	461893 usec
Time savings	176 milliseconds

Example 2

Hardware	TI OMAP board, using serial console
Kernel Version	CELF-1 (040126)
Configuration	Kernel booted with XIP, CRAMFS root file system, with preset-LPJ
Time without "quiet" option	551735 usec
Time with "quiet" option	280676 usec
Time savings	271 milliseconds

Embedded Linux Speedups

- Remove unnecessary functions and device drivers
 - Reduce kernel size
- Modularization of device driver
 - Upload module after boot sequence finished
 - e.g. CAM device

Asynchronous function calls

- For some initialization functions, it could be done asynchronously
- Asynchronous function calls has been merged in mainline starting from 2.6.29

Avoid performance measurement routine

- kernel updates loops_per_jiffy value at every booting to using timer
 - This value should be fixed until hardware changes or modify setting
 - The problem is this routine uses some loops (Delaying)
- Method examples
 - Find out loops_per_jiffy after basic boot sequence
 - Modify init/calibrate.c in kernel source
- TI OMAP 1510/158MHz
 - 212 milliseconds

RTC no sync

- `get_cmos_time()` is used to read the value of the external real-time clock (RTC) when the kernel boots
- This routine delays until the edge of the next second rollover, in order to ensure that the time value in the kernel is accurate with respect to the RTC
- Introduces up to 1 second of variability in the total bootup time
- Reconfigure kernel with "CONFIG_RTC_NO_SYNC_ON_READ" is enabled
- “Fast boot options” labeled as “No SYNC on read of Real Time Clock”

Using read-only file system

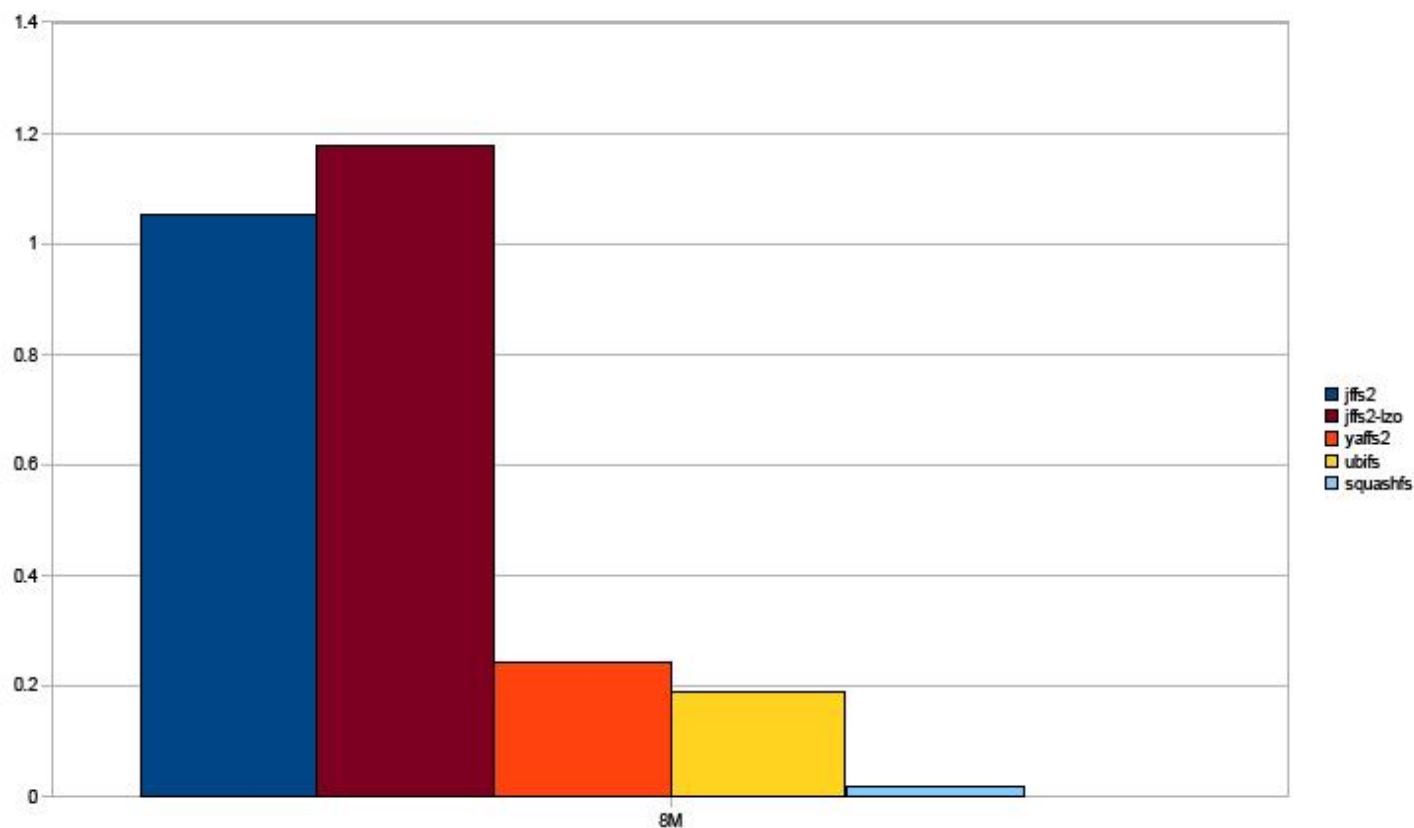
- Without using write functionality, simplify using file system
- There are no delay time on mount read-only file systems
- One of simplest, RomFS
- The cramfs has compression functionality
- Method examples
 - Classify read-only, R/W, temporary files that will include file system
 - Position cramfs read-only files
 - Using lazy mount technique, to include R/W files
 - Temporary file uses tmpfs

Using lazy mount technique on R/W file systems

- slow mount speed file systems like jffs2 mount after finishing boot sequence on user request
- On booting, using read-only file instead
- JFFS2 mount of 8M filesystem partition is over 1 second
- UBIFS mount of 8M partition is under .2 seconds
- Squashfs mount of 8M partition is under 50 milliseconds



Zoom - Mount time (seconds) - 8M



Deferred Initcalls

- Find modules that are not required for core functionality of product
 - Ex: USB on a camera – `uhci_hcd_usb`, `ehci_hcd_init`
- Change module init routine declaration
 - `module_init(foo_init)` to
 - `deferred_module_init(foo_init)`
- Modules marked like this are not initialized during kernel boot
- After main init, do:
 - `echo 1 >/proc/deferred_initcalls`
- Deferred initcalls are run

Application Speedups

- Using binary script, not shell script
- Using init process with simplified and optimized
- Parallel RC scripts
- Application XIP
- Using pre-link shared lib
- Optimize of application programs
- Move from glibc to uClibc

Using binary script, not shell script

- Disadvantages performance of shell script
 - Interpreter
 - Interpret every line and perform
 - Interpret meaningless line and phrases
 - Use fork & exec technique when perform commands
- Make binary script without fork & exec technique

Using init process with simplified and optimized

- BusyBox optimizations
 - The modified init scripts must be run with "bash" as well as the BusyBox "ash" shell
- Follow the rules listed below to reduce execution time for the init scripts:
 - Do not use unnecessary codes in the scripts.
 - Replace external commands and utilities with the BusyBox built-ins as far as possible.
 - Do not use the piped commands as far as possible.
 - Reduce the number of commands within a pipe.
 - Do not use the back-quoted commands as far as possible.
 - The main goal of such optimization is to reduce the number of the "fork/exec" calls during a script execution.

Parallel RC scripts

- run RC scripts in parallel
 - to take advantage of the multi-processing capabilities of the OS (such as overlapping execution with I/O, etc.)
- Need shell/shell script support

Pre-linking

- CPU: ARM9
- Base OS : MontaVistaLinux CEE 3.1 (2.4.20 kernel, glibc-2.3)
- GUI:X Window, GTK+(1.2.10) / Motif like toolkit



User Response Time/Boot up time Application of prelink

- Normal dynamic linked ELF: Over 2 sec. to start up a multimedia application process (fork ~ main)

Stages of Processing in Process Start up	elapsed time
1. Layout and map shared libs to virtual address space	96ms
2. Resolve symbol references	2354ms
3. Init of each ELF file	29ms

- Prelink eliminates symbol reference resolution processes.

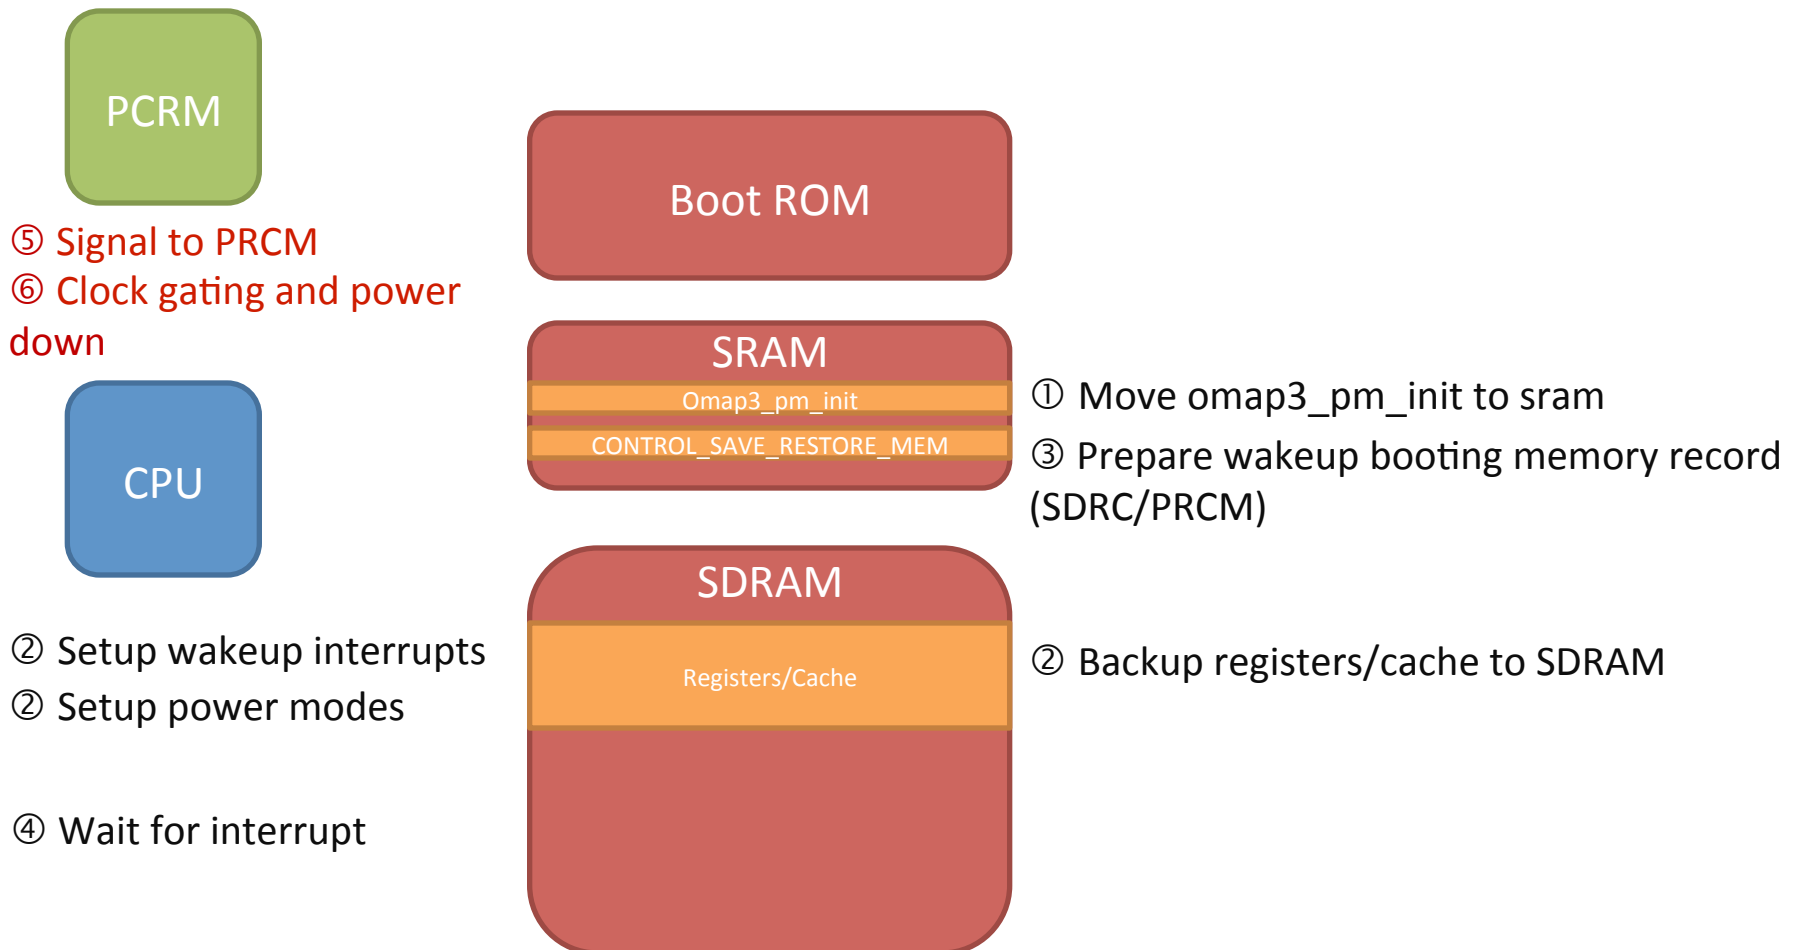
2,479msec → 125msec

- Boot up time also reduced: Over 1min. → 20sec.

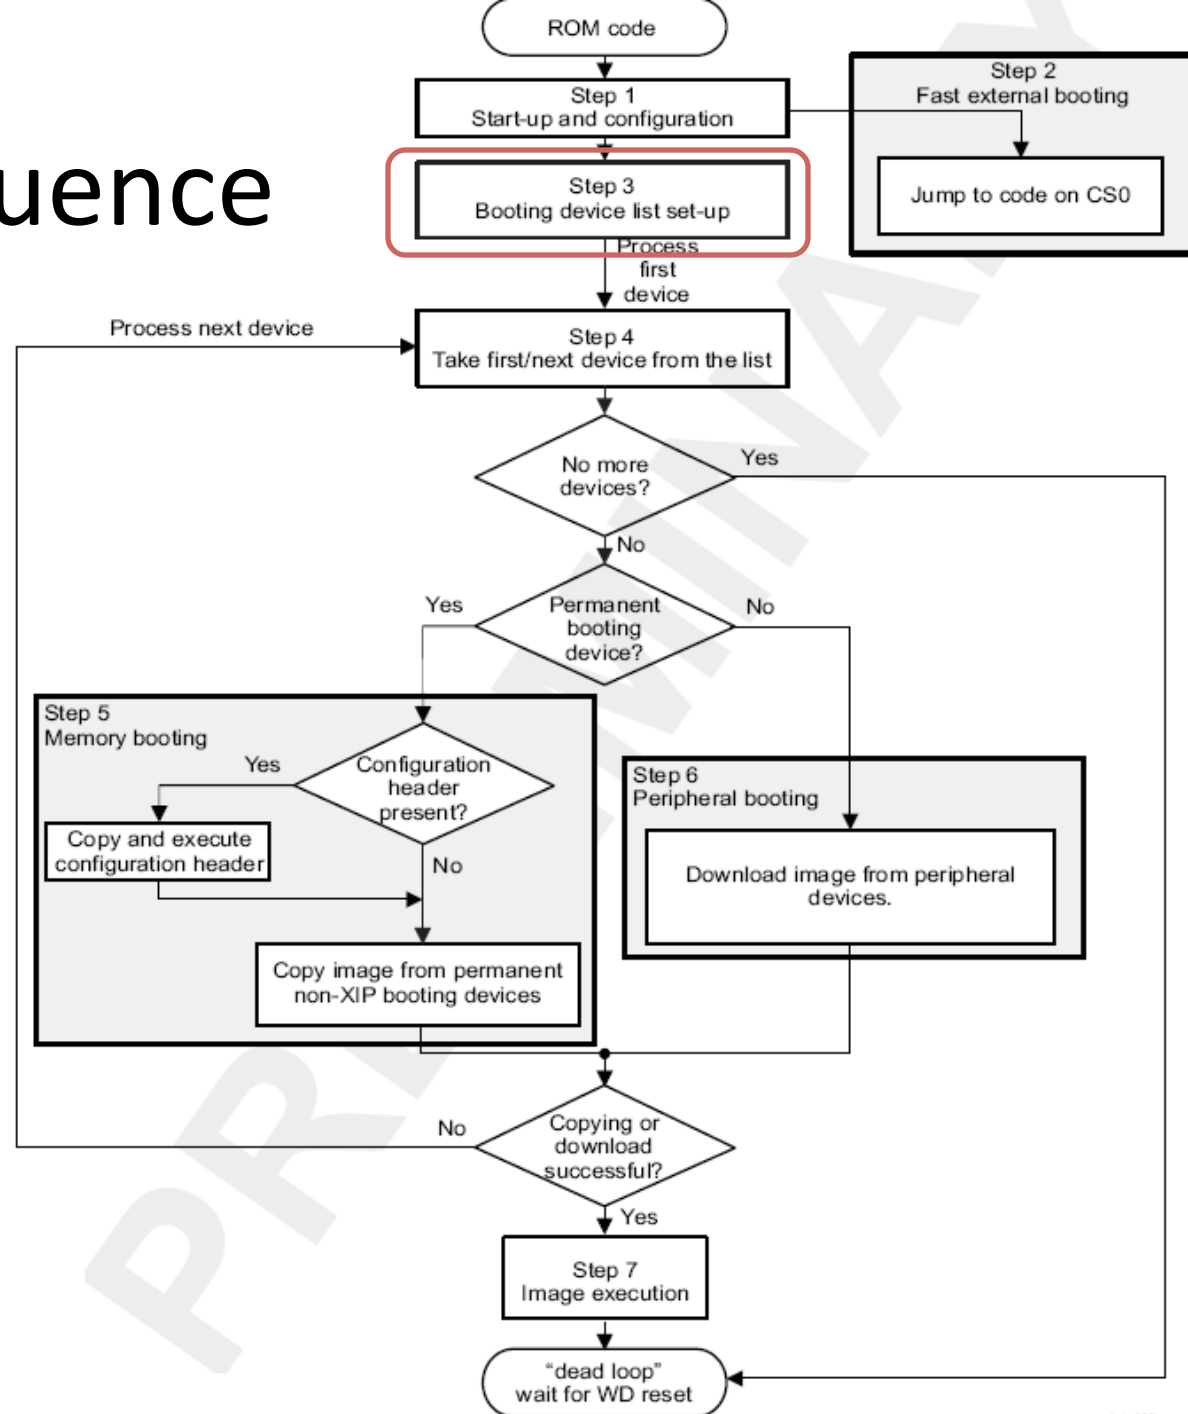
Suspend to RAM

OMAP3430

- Enter dormant mode

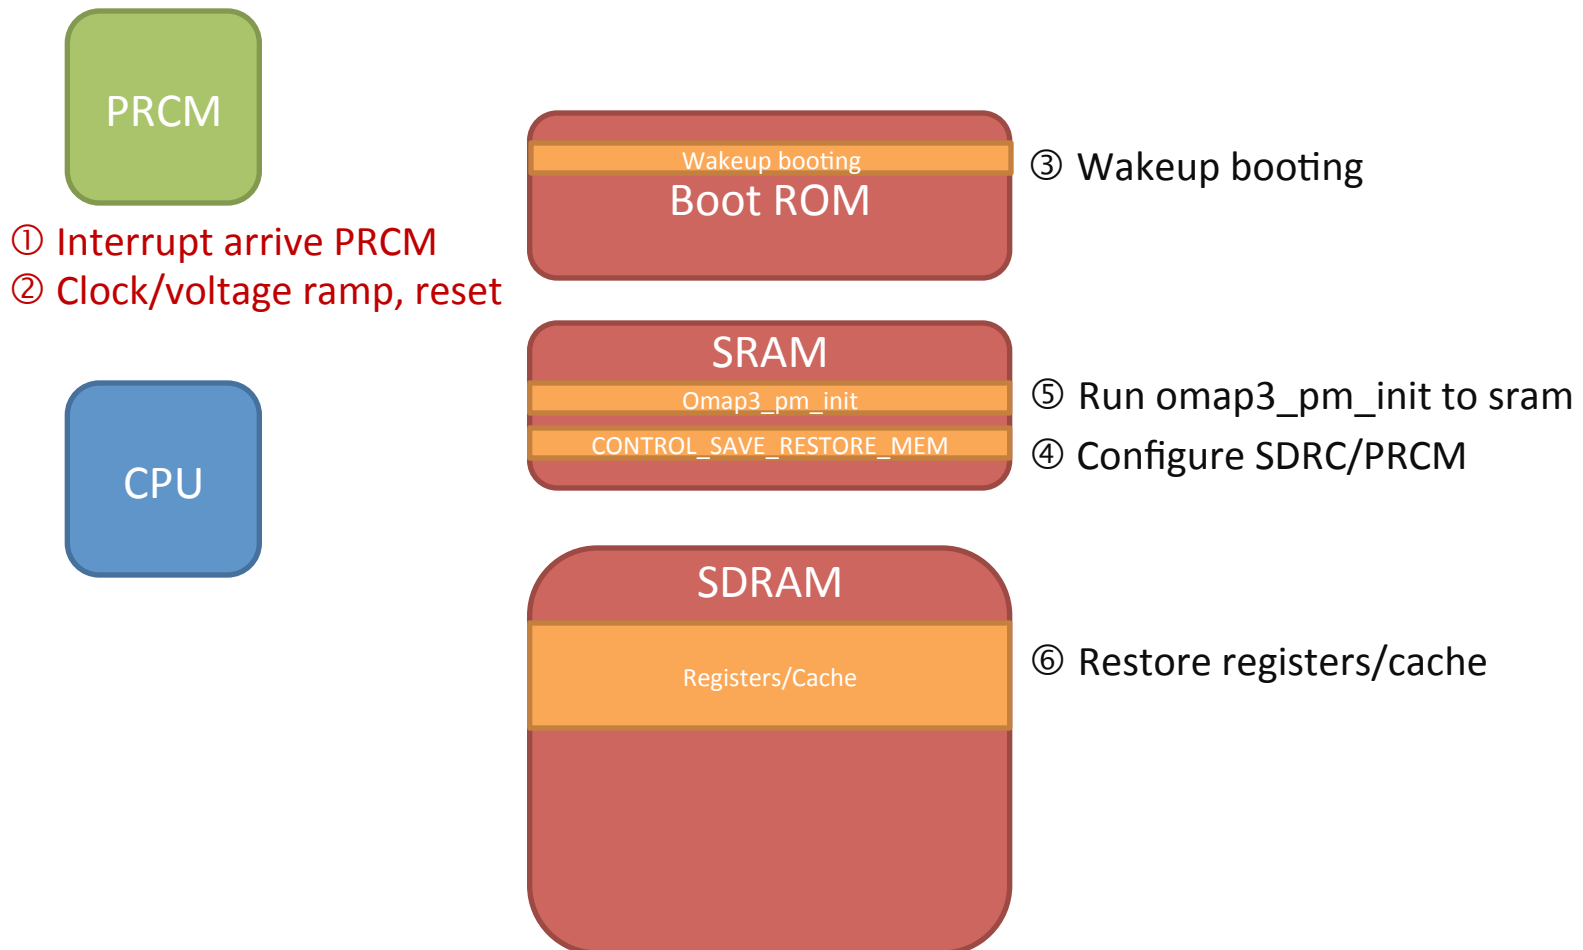


Booting sequence



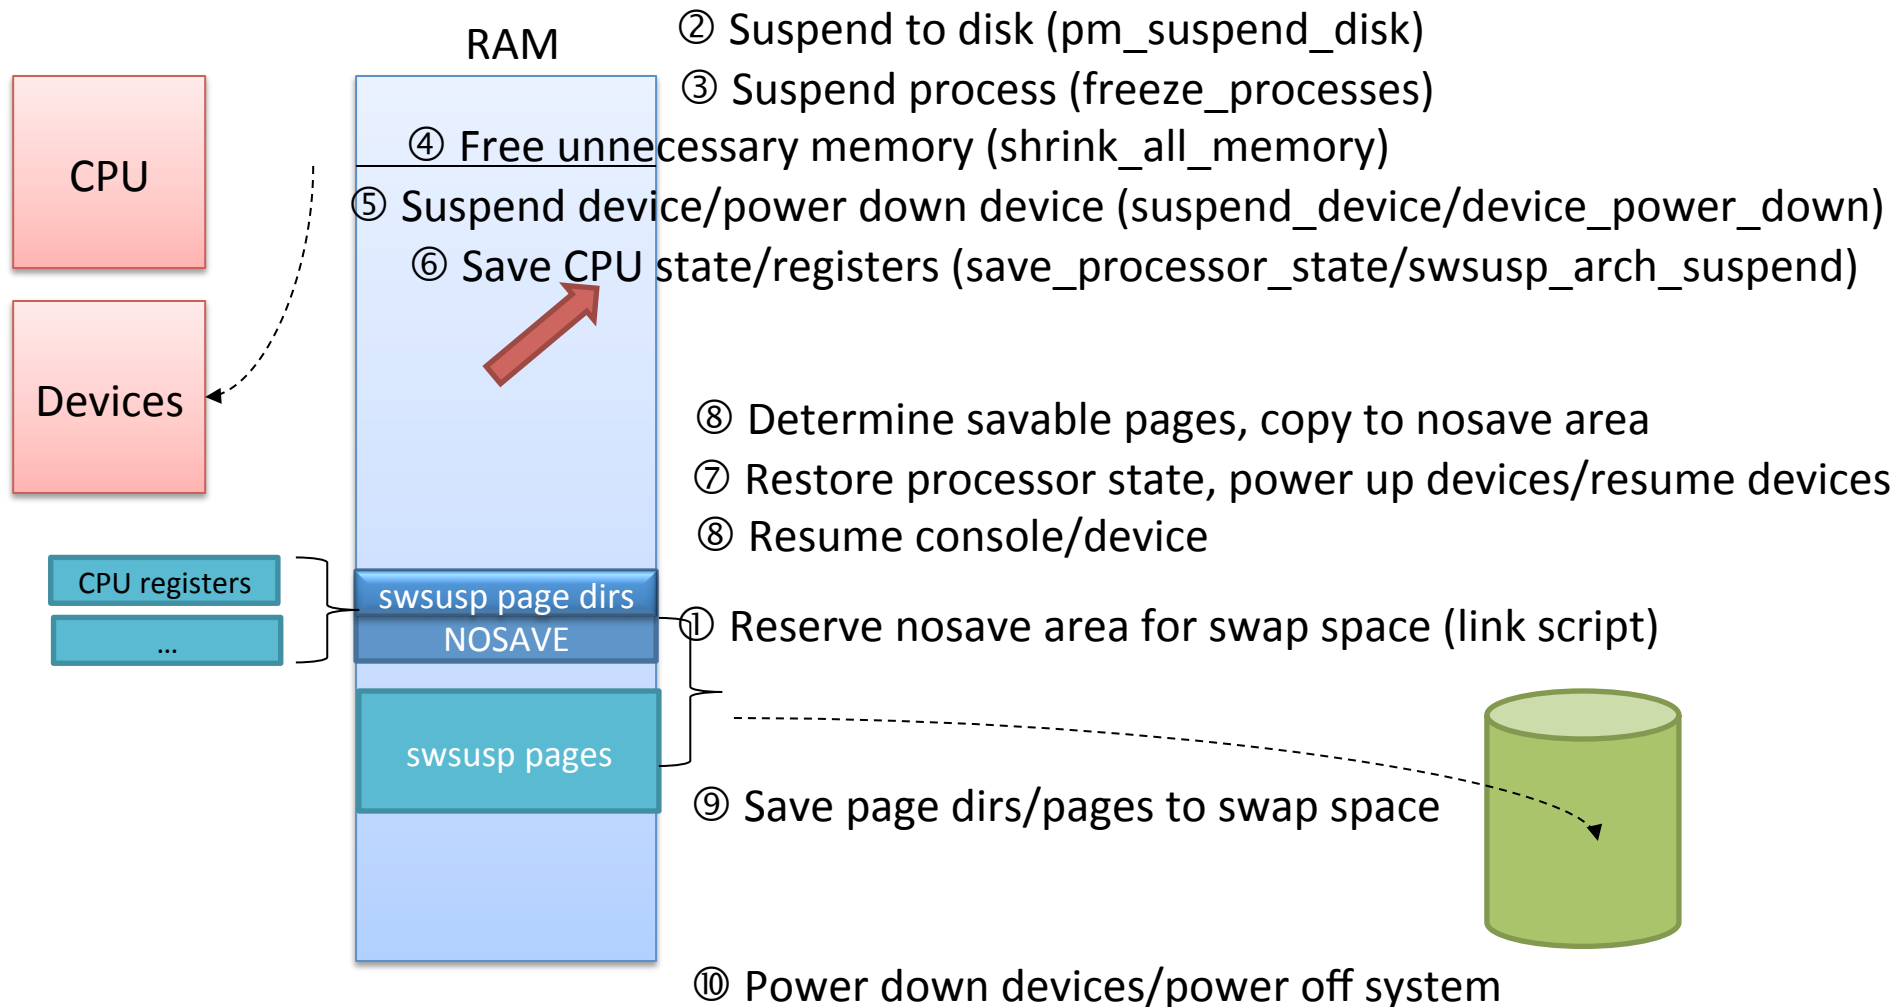
OMAP3430

- Exit dormant mode



Suspend to Disk

Suspend Flow



Resume Flow

① CPU booting/bootloader/kernel loading/decompress kernel/start_kernel

② late_initcall(software_resume)

③ Check signature and header of swsusp image

④ Allocate memory for swsusp image

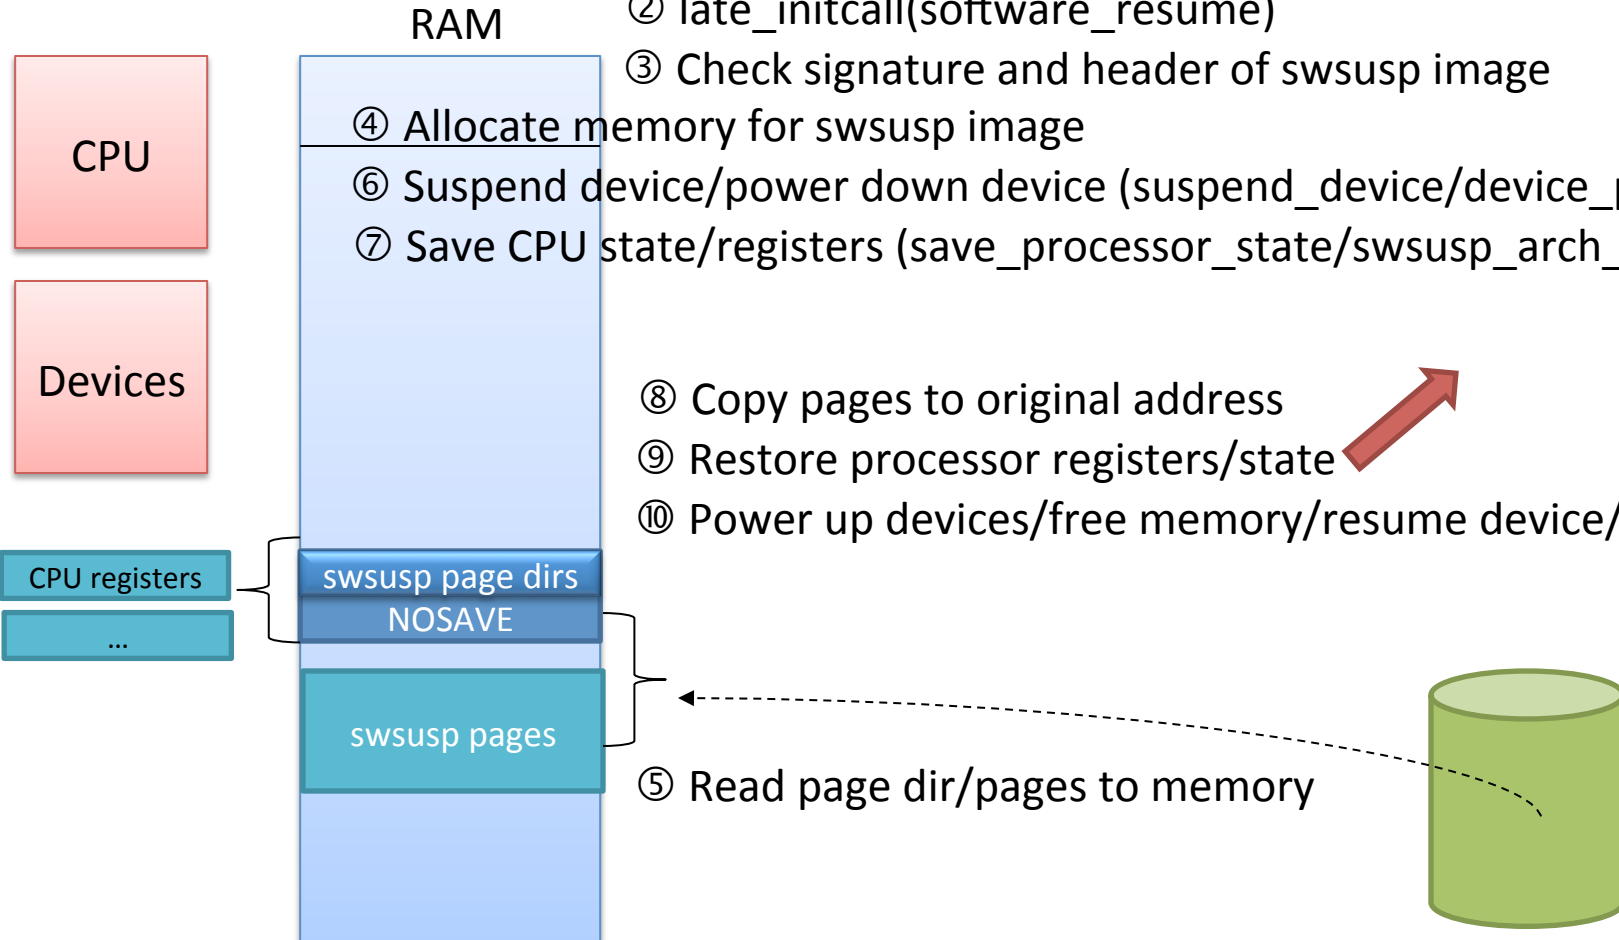
⑥ Suspend device/power down device (suspend_device/device_power_down)

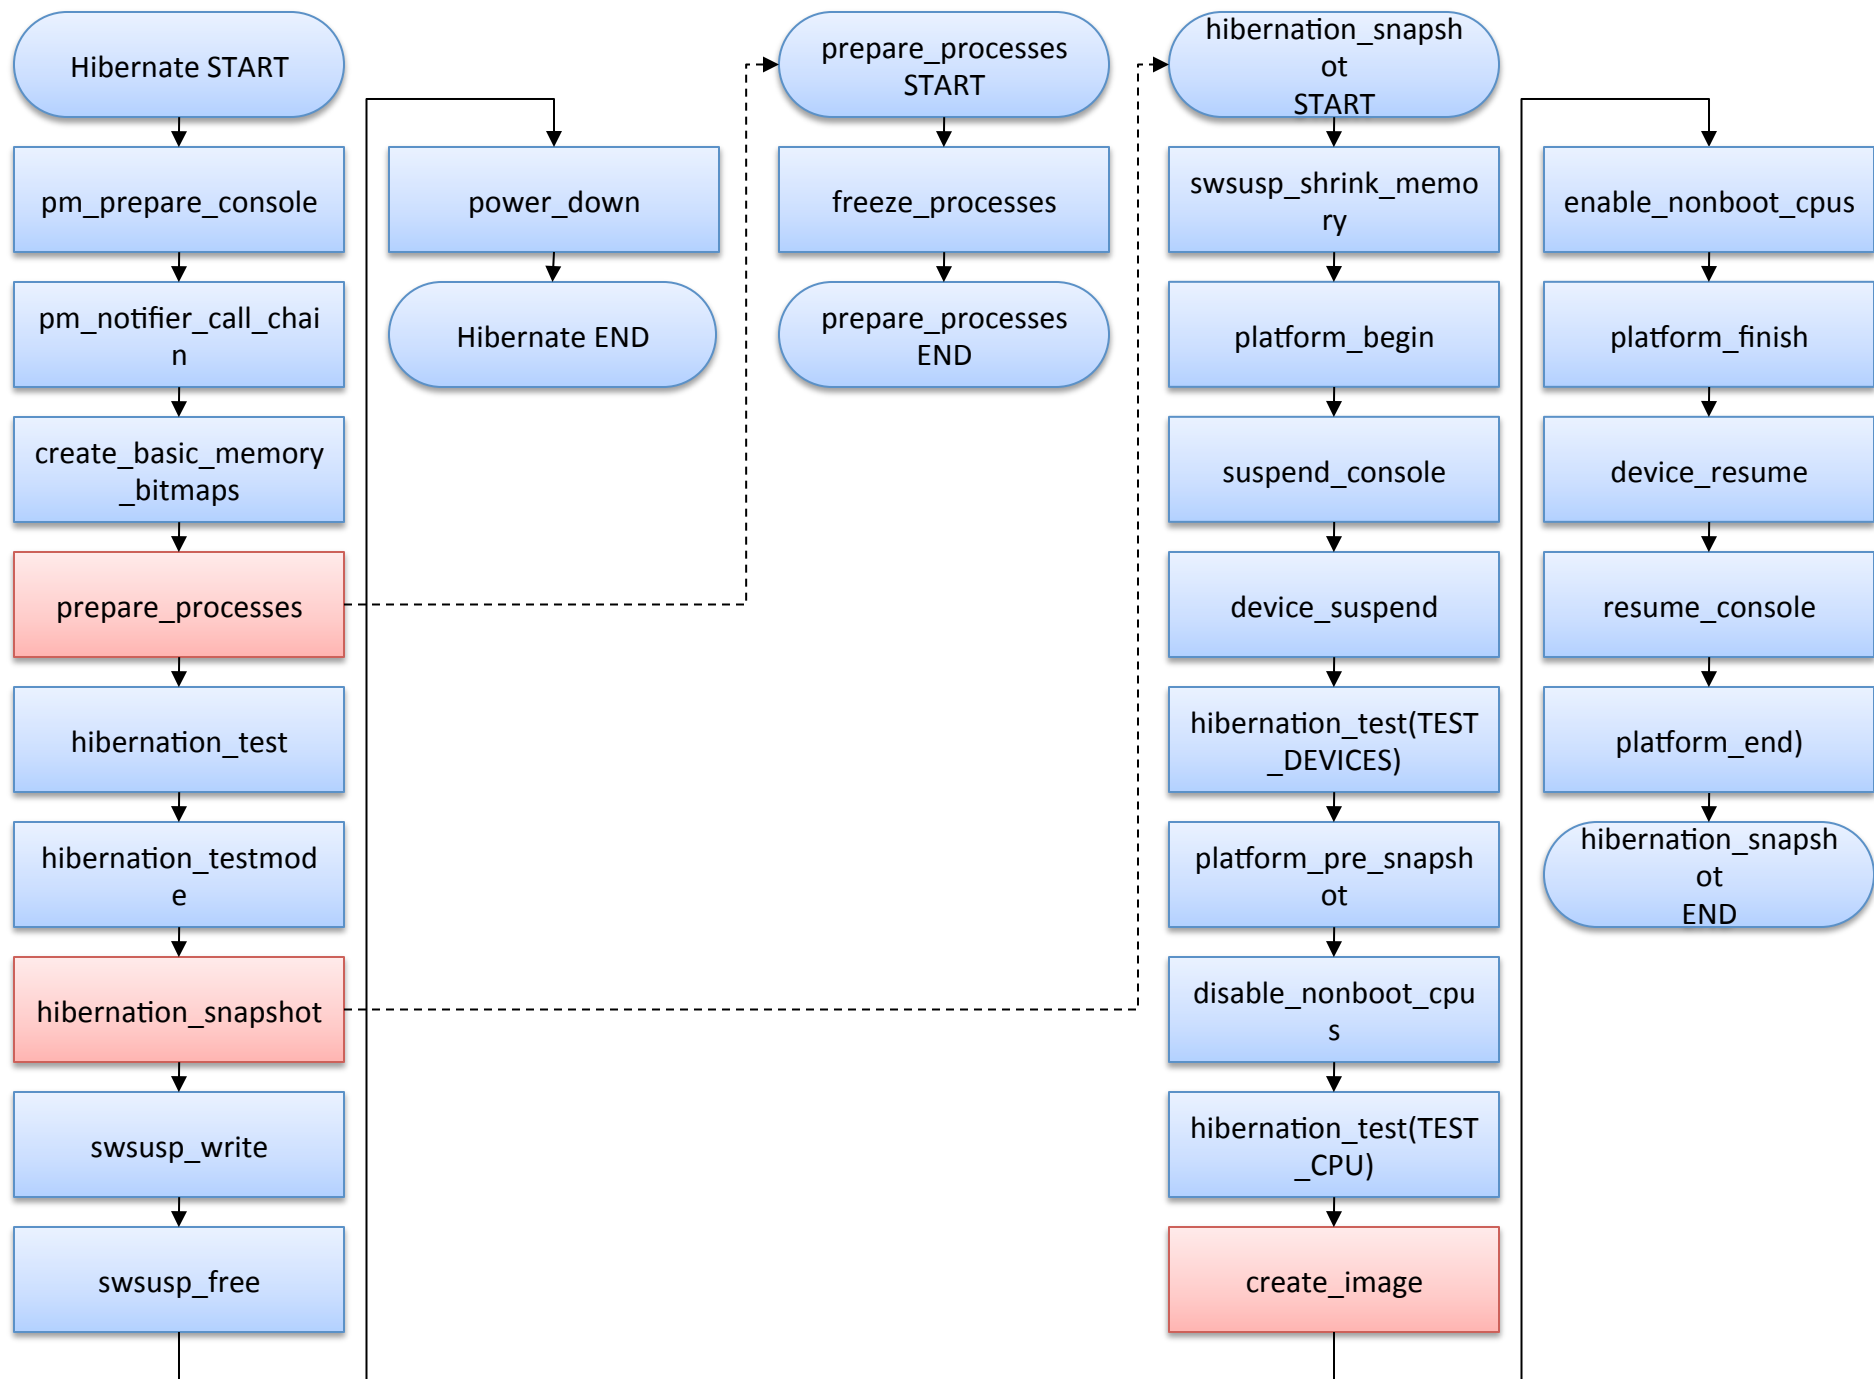
⑦ Save CPU state/registers (save_processor_state/swsusp_arch_suspend)

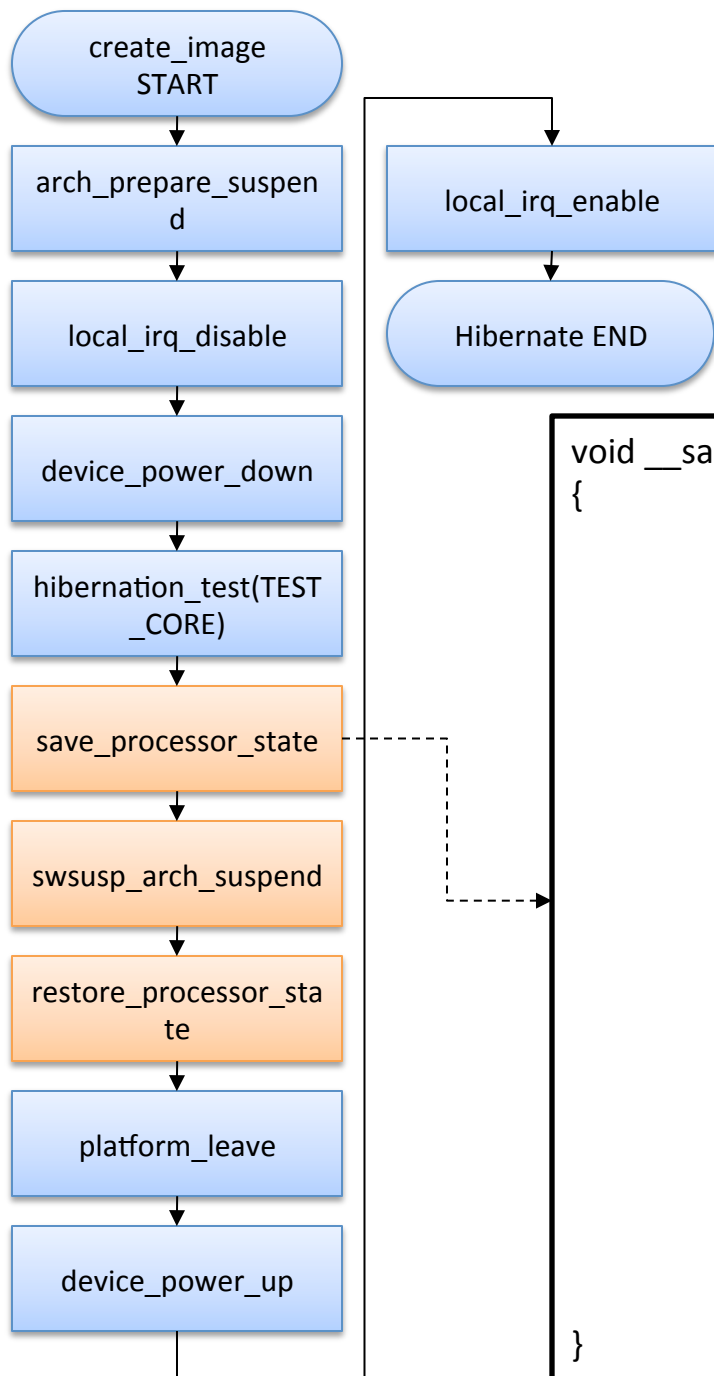
⑧ Copy pages to original address

⑨ Restore processor registers/state

⑩ Power up devices/free memory/resume device/thaw processes







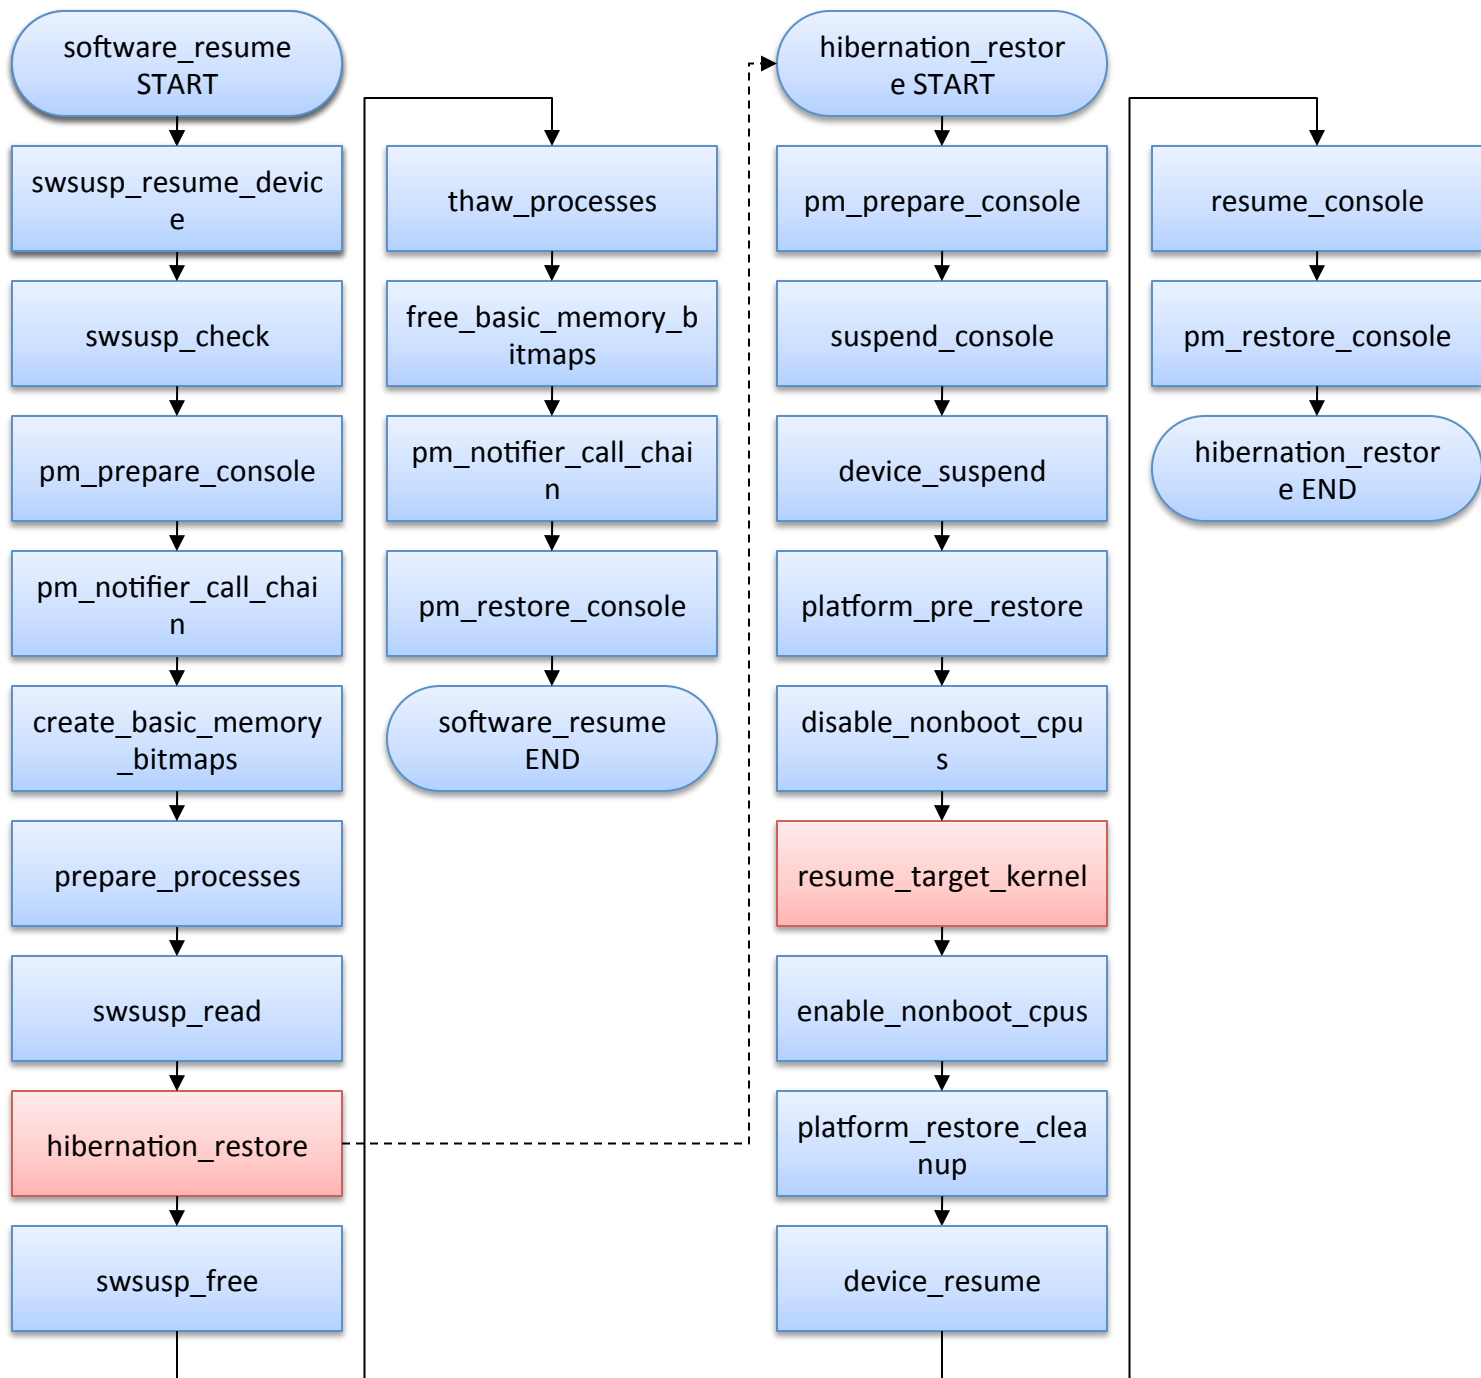
```
void __save_processor_state(struct saved_context *ctxt)
{
```

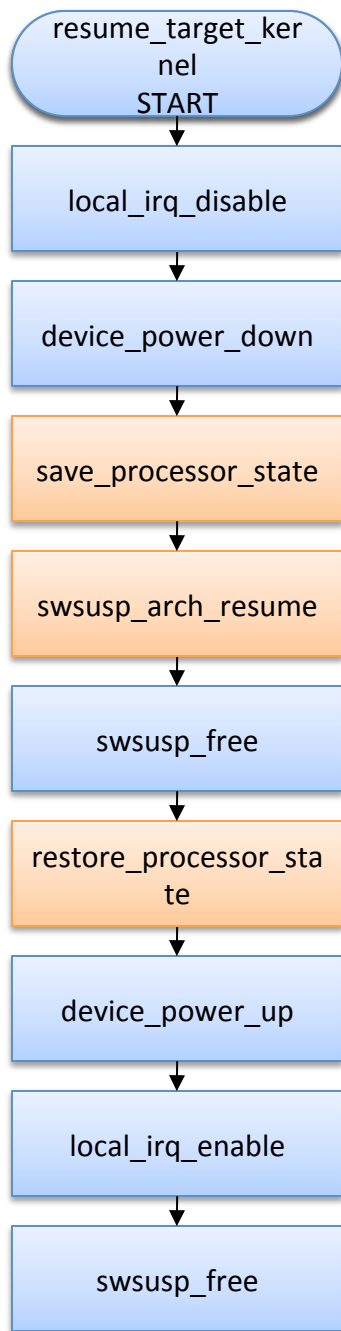
```
    /* save preempt state and disable it */
    preempt_disable();
```

```
    /* save coprocessor 15 registers */
```

```
    asm volatile ("mrc p15, 0, %0, c1, c0, 0" : "=r" (ctxt->CR));
    asm volatile ("mrc p15, 0, %0, c3, c0, 0" : "=r" (ctxt->DACR));
    asm volatile ("mrc p15, 0, %0, c5, c0, 0" : "=r" (ctxt->D_FSR));
    asm volatile ("mrc p15, 0, %0, c5, c0, 1" : "=r" (ctxt->I_FSR));
    asm volatile ("mrc p15, 0, %0, c6, c0, 0" : "=r" (ctxt->FAR));
    asm volatile ("mrc p15, 0, %0, c9, c0, 0" : "=r" (ctxt->D_CLR));
    asm volatile ("mrc p15, 0, %0, c9, c0, 1" : "=r" (ctxt->I_CLR));
    asm volatile ("mrc p15, 0, %0, c9, c1, 0" : "=r" (ctxt->D_TCMRR));
    asm volatile ("mrc p15, 0, %0, c9, c1, 1" : "=r" (ctxt->I_TCMRR));
    asm volatile ("mrc p15, 0, %0, c10, c0, 0" : "=r" (ctxt->TLBLR));
    asm volatile ("mrc p15, 0, %0, c13, c0, 0" : "=r" (ctxt->FCSE));
    asm volatile ("mrc p15, 0, %0, c13, c0, 1" : "=r" (ctxt->CID));
    asm volatile ("mrc p15, 0, %0, c2, c0, 0" : "=r" (ctxt->TTBR));
```

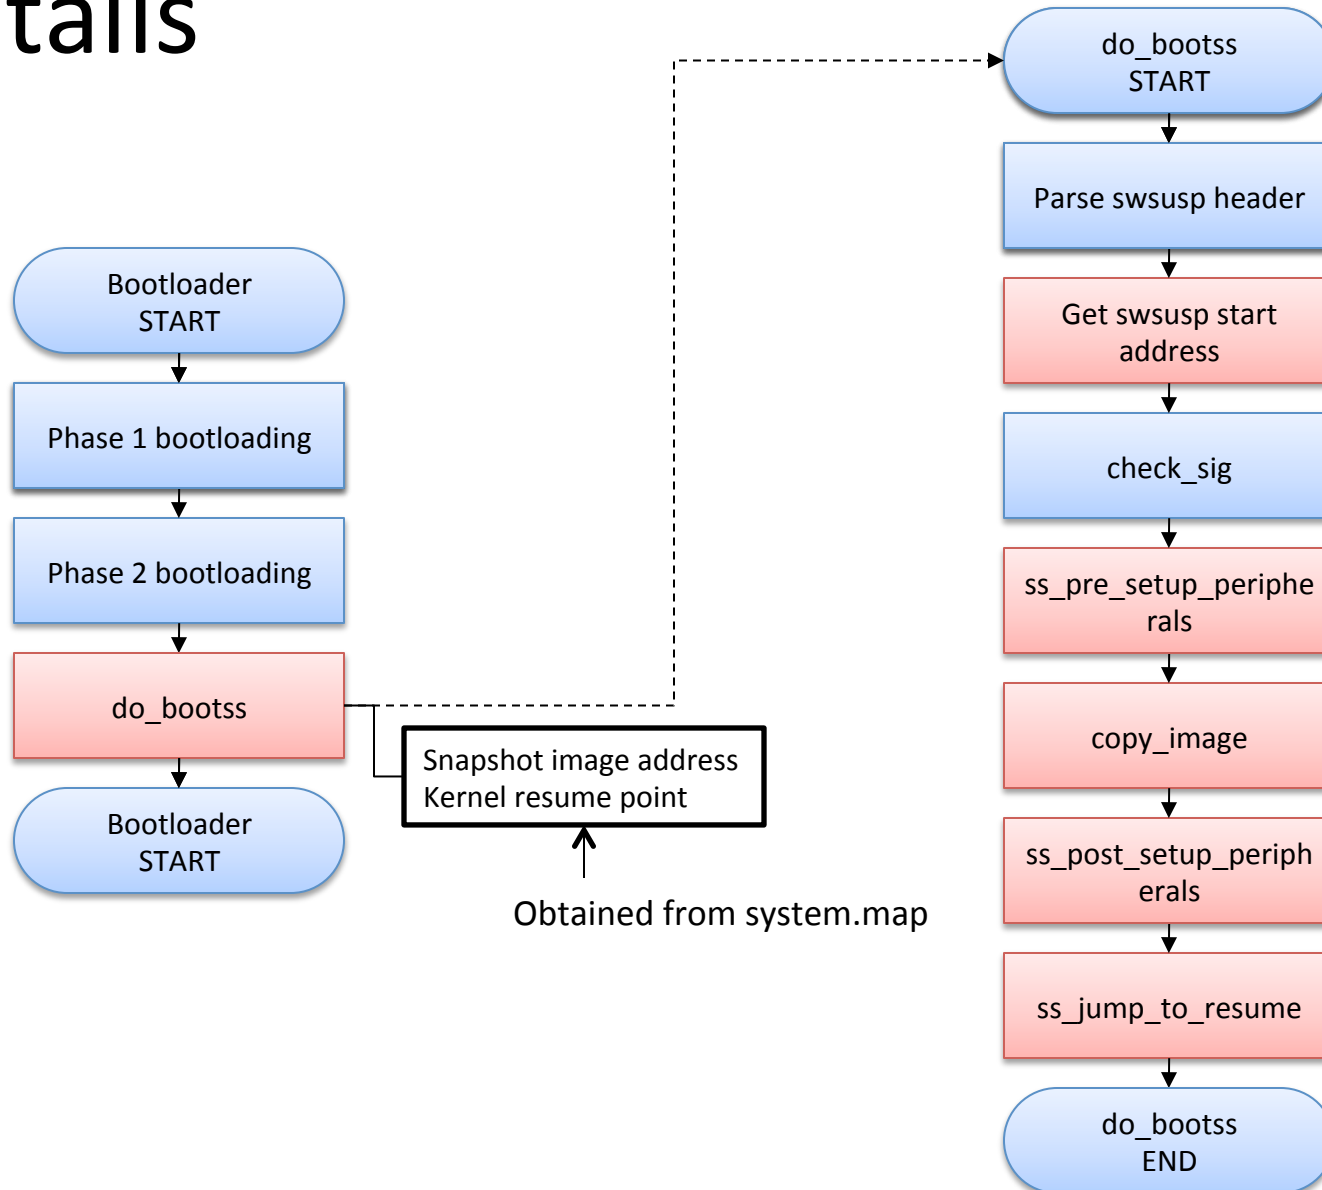
```
}
```





Snapshot Booting Implementation

Details



Preserving hibernation system image

