

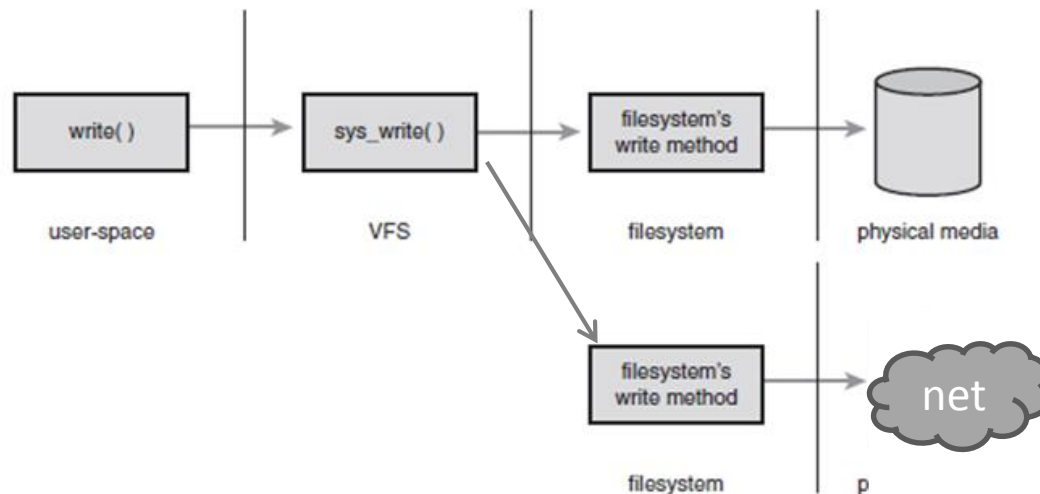
# Unit 7: File System Driver

Prof. Li-Pin Chang

ESSLab@NCTU

# Linux VFS

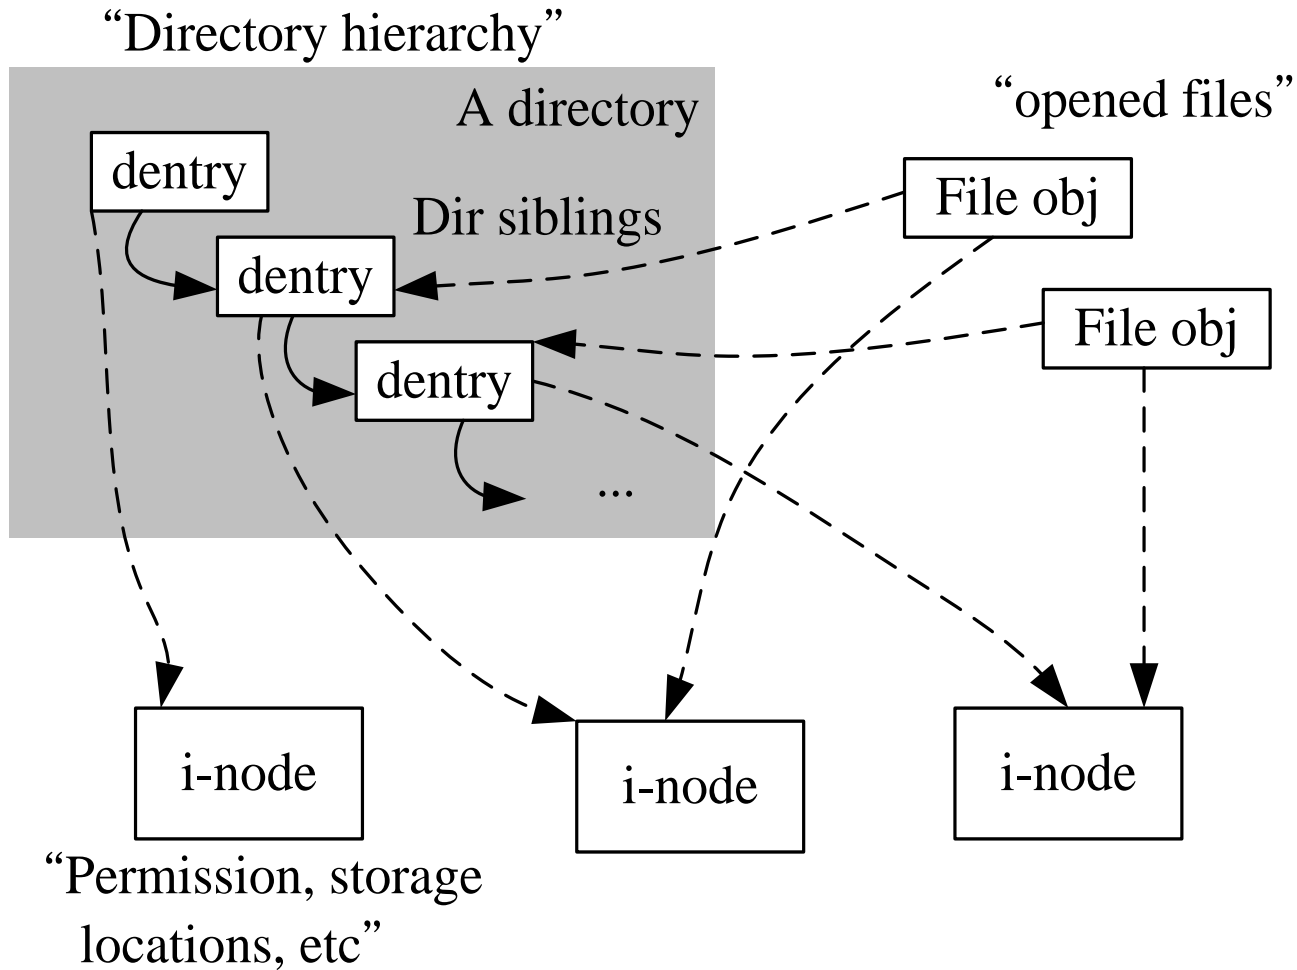
- Virtual File-system Switch
  - Provides an abstraction layer between user space and file system drivers
  - A set of common operations for various file system implementations
    - ext234 (disk)
    - nfs (networ)
    - yaffs2 (flash)



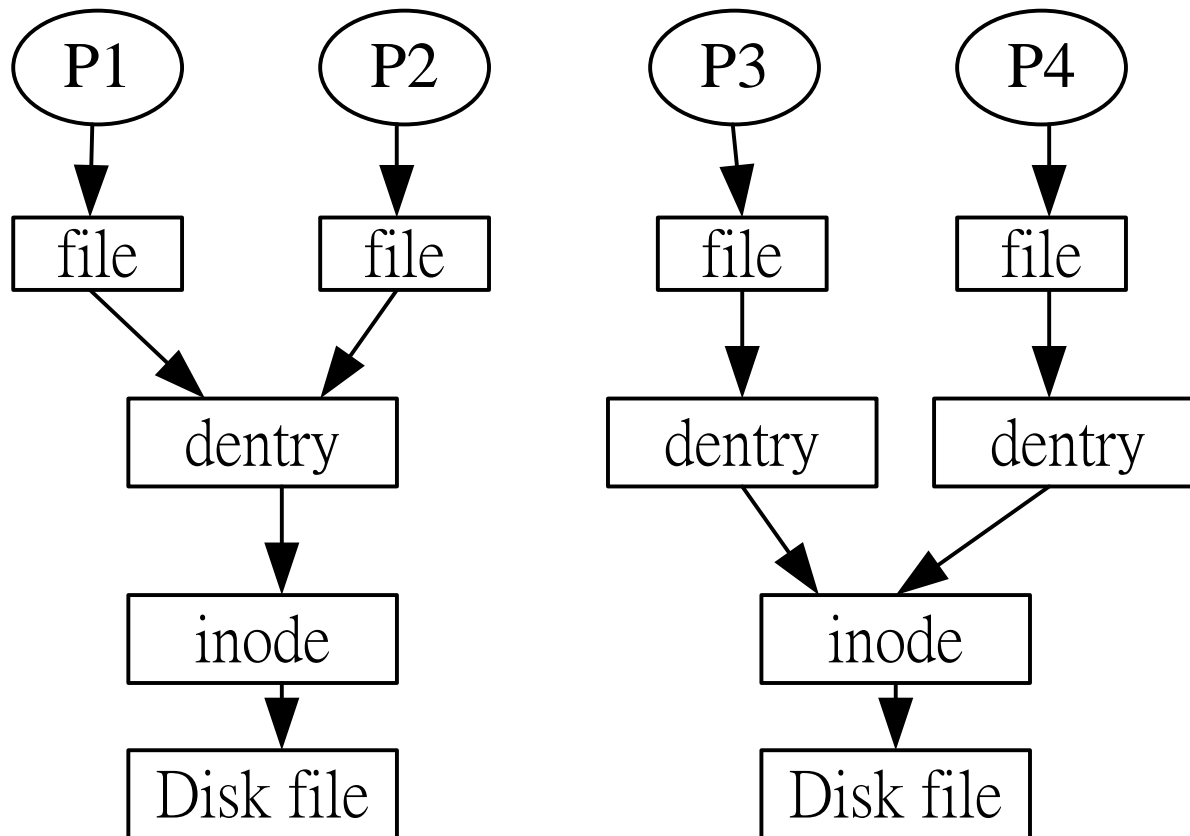
# VFS Memory Objects

- *Superblock* object
  - represent the entire file system
- *Inode* object
  - represent an individual file
- *file* object
  - represent an opened file
- *dentry* object
  - represent an individual directory entry
- Each type of object is associated with a set of operations

# VFS Memory Objects



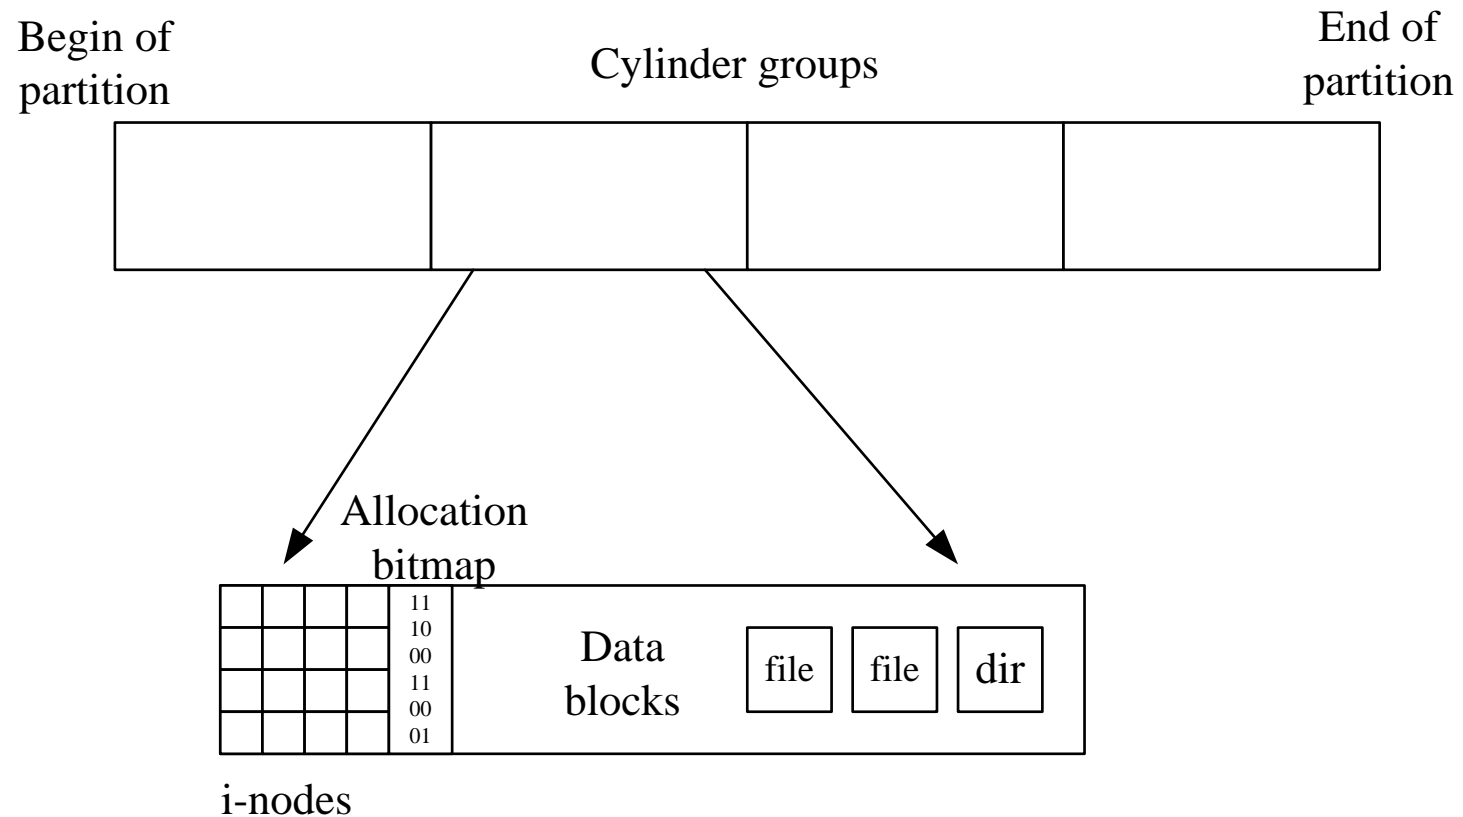
# VFS Object Relationship



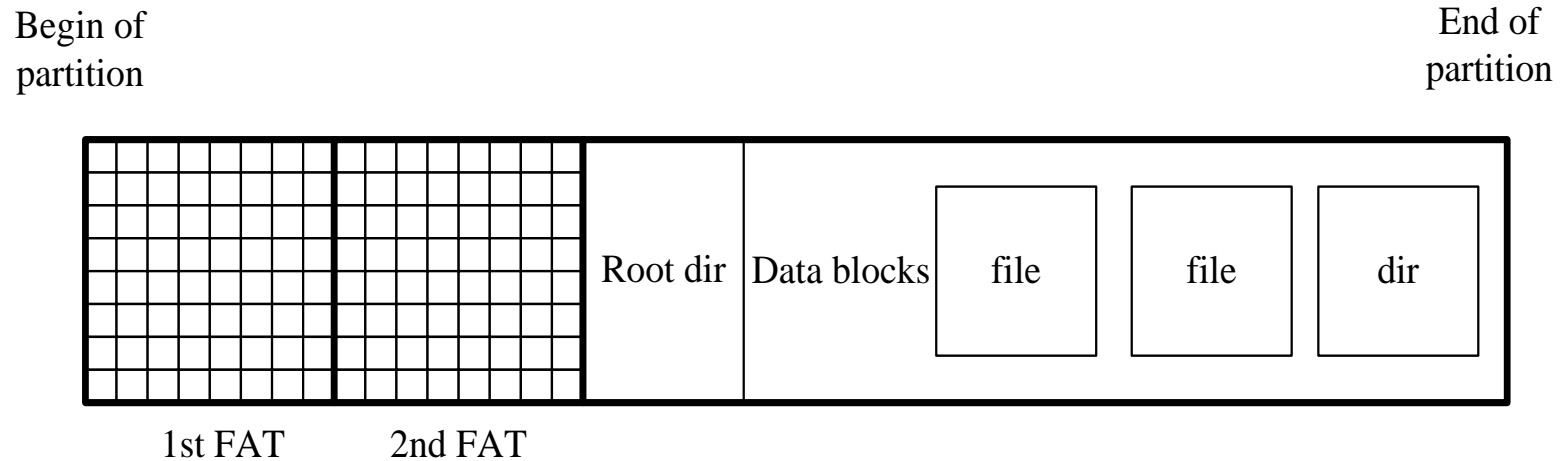
# File-System Metadata

- **In disk**, specific to file system implementation
  - Linux ext234 file system
    - Super block
    - Inode
    - Allocation bitmaps
  - Microsoft FAT file system
    - File allocation tables
    - Directory
- File system drivers must create VFS memory objects based on their metadata
  - Ext234 copies inodes from disk to memory
  - FAT file system does not have disk inodes; it creates inodes on the fly

# Layout of the Linux ext 2/3/4 file system



# The layout of FAT 12/16/32 file system





# ramfs

- In this lecture, we study ramfs
- ramfs is a file system that allocates memory pages for storing file data
  - Losing all pages after unmounting!
  - Pages cannot be swapped to disk
    - Fast, but uses more memory
    - Use *tmpfs* instead if you need to create a very large ram file system
- ramfs does not write to disk
  - `mount -t ramfs none /mnt`

# ramfs Initialization

```
static struct file_system_type ramfs_fs_type = {
    .name      = "ramfs",
    .get_sb     = ramfs_get_sb,
    .kill_sb    = ramfs_kill_sb,
};

int ramfs_get_sb(struct file_system_type *fs_type,
    int flags, const char *dev_name, void *data, struct vfsmount *mnt)
{
    return get_sb_nodev(fs_type, flags, data, ramfs_fill_super, mnt);
}

static void ramfs_kill_sb(struct super_block *sb)
{
    kfree(sb->s_fs_info);
    kill_litter_super(sb);
}

static int __init init_ramfs_fs(void)
{
    return register_filesystem(&ramfs_fs_type);
}
```

```

static int ramfs_fill_super(struct super_block * sb, void * data, int silent)
{
    struct ramfs_fs_info *fsi;
    struct inode *inode = NULL;
    struct dentry *root;
    int err;

    save_mount_options(sb, data);

    fsi = kzalloc(sizeof(struct ramfs_fs_info), GFP_KERNEL);

    ...

    sb->s_maxbytes      = MAX_LFS_FILESIZE;
    sb->s_blocksize      = PAGE_CACHE_SIZE;
    sb->s_blocksize_bits = PAGE_CACHE_SHIFT;
    sb->s_magic          = RAMFS_MAGIC;
    sb->s_op             = &ramfs_ops;
    sb->s_time_gran      = 1;

    inode = ramfs_get_inode(sb, S_IFDIR | fsi->mount_opts.mode, 0); ←
    if (!inode) {
        err = -ENOMEM;
        goto ↓fail;
    }

    root = d_alloc_root(inode); ←
    sb->s_root = root;
    if (!root) {
        err = -ENOMEM;
        goto ↓fail;
    }

    return 0;
fail:
    kfree(fsi);
    sb->s_fs_info = NULL;
    iput(inode);
    return err;
} ? end ramfs_fill_super ?

```

Allocate an *inode* for  
the root directory

Allocate an *dentry* object  
for the root inode

# Superblock Object

- An object that contains information regarding the whole file system
  - The device that the file system is mounted on
  - Block size (multiple of disk sectors)
  - A dirty flag (needing fsck or not)
- A superblock is created when file system is mounted on a device
  - `ramfs_getsb()`

```

struct super_block {
    struct list_head    s_list;           /* list of all superblocks */
    dev_t               s_dev;           /* identifier */
    unsigned long       s_blocksize;     /* block size in bytes */
    unsigned char       s_blocksize_bits; /* block size in bits */
    unsigned char       s_dirt;          /* dirty flag */
    unsigned long long  s_maxbytes;      /* max file size */
    struct file_system_type s_type;      /* filesystem type */
    struct super_operations s_op;        /* superblock methods */
    struct dqquot_operations *dq_op;     /* quota methods */
    struct quotactl_ops  *s_qcop;        /* quota control methods */
    struct export_operations *s_export_op; /* export methods */
    unsigned long       s_flags;         /* mount flags */
    unsigned long       s_magic;         /* filesystem's magic number */
    struct dentry        *s_root;        /* directory mount point */
    struct rw_semaphore  s_umount;       /* unmount semaphore */
    struct semaphore     s_lock;         /* superblock semaphore */
    int                  s_count;        /* superblock ref count */
    int                  s_need_sync;    /* not-yet-synced flag */
    atomic_t             s_active;       /* active reference count */
    void                 *s_security;    /* security module */
    struct xattr_handler **s_xattr;      /* extended attribute handlers */

    struct list_head     s_inodes;       /* list of inodes */
    struct list_head     s_dirty;        /* list of dirty inodes */
    struct list_head     s_io;           /* list of writebacks */
    struct list_head     s_more_io;      /* list of more writeback */
    struct hlist_head     s_anon;        /* anonymous dentries */
    struct list_head     s_files;        /* list of assigned files */
    struct list_head     s_dentry_lru;   /* list of unused dentries */
    int                  s_nr_dentry_unused; /* number of dentries on list */
    struct block_device  *s_bdev;        /* associated block device */
    struct mtd_info       *s_mtd;        /* memory disk information */
    struct list_head     s_instances;    /* instances of this fs */
    struct quota_info     s_dquot;       /* quota-specific options */
    int                  s_frozen;       /* frozen status */
    wait_queue_head_t    s_wait_unfrozen; /* wait queue on freeze */
    char                 s_id[32];       /* text name */
    void                 *s_fs_info;     /* filesystem-specific info */
    fmode_t              s_mode;         /* mount permissions */
    struct semaphore     s_vfs_rename_sem; /* rename semaphore */
    u32                  s_time_gran;    /* granularity of timestamps */
    char                 *s_subtype;     /* subtype name */
    char                 *s_options;     /* saved mount options */
};

```

## Essential members

- s\_dirt
- s\_op
- \*s\_root
- s\_inodes
- s\_dirty
- s\_files
- \*s\_bdev

# Superblock Operations

- A structure of ~20 callbacks
  - Operations applied to file system
    - Getting file system statistics
    - Syncing all dirty data
  - inode allocation/de-allocation (in ram and in disk)
- A file system may or may not register a callback for each superblock operation
  - If a callback is NULL, the VFS simply does nothing or calls the generic handler
- ramfs superblock operations

```
static const struct super_operations ramfs_ops = {  
    .statfs      = simple_statfs,  
    .drop_inode  = generic_delete_inode,  
    .show_options = generic_show_options,  
};
```

```
struct super_operations {
    struct inode *(*alloc_inode)(struct super_block *sb);
    void (*destroy_inode)(struct inode *);
    void (*dirty_inode) (struct inode *);
    int (*write_inode) (struct inode *, int);
    void (*drop_inode) (struct inode *);
    void (*delete_inode) (struct inode *);
    void (*put_super) (struct super_block *);
    void (*write_super) (struct super_block *);
    int (*sync_fs)(struct super_block *sb, int wait);
    int (*freeze_fs) (struct super_block *);
    int (*unfreeze_fs) (struct super_block *);

    int (*statfs) (struct dentry *, struct kstatfs *);
    int (*remount_fs) (struct super_block *, int *, char *);
    void (*clear_inode) (struct inode *);
    void (*umount_begin) (struct super_block *);
    int (*show_options)(struct seq_file *, struct vfsmount *);
    int (*show_stats)(struct seq_file *, struct vfsmount *);
    ssize_t (*quota_read)(struct super_block *, int, char *, size_t, loff_t);
    ssize_t (*quota_write)(struct super_block *, int, const char *, size_t, loff_t);
    int (*bdev_try_to_free_page)(struct super_block*, struct page*, gfp_t);
};
```

# inode

- inode is an object associated with a disk file
  - Contains all the information needed to manipulate a file
    - File size, timestamps, disk locations, permissions, etc
  - Unique to each file, independent of file name
- An inode is constructed in memory when a file is accessed
  - Some file systems have inode in disk (ext234), so just copy it from disk
  - Many file systems have no disk inodes (FAT, NTFS, ...), construct on the fly



# inode lists

- An inode resides in one of the following lists (via *i\_list*)
  - A global list named *inode\_unused*
    - Clean inodes but currently no processes use them
  - A global list named *inode\_inuse*
    - Clean inodes that some processes refer to them
  - A list *s\_dirty* in the corresponding superblock
    - Modified inodes awaiting write back
- A superblock also chains all its inodes in a list *s\_inode*
  - Via *i\_sb\_list*

```

struct inode {
    struct hlist_node    i_hash;           /* hash list */
    struct list_head     i_list;           /* list of inodes */
    struct list_head     i_sb_list;        /* list of superblocks */
    struct list_head     i_dentry;        /* list of dentries */
    unsigned long        i_ino;            /* inode number */
    atomic_t             i_count;          /* reference counter */
    unsigned int         i_nlink;          /* number of hard links */
    uid_t                i_uid;            /* user id of owner */
    gid_t                i_gid;            /* group id of owner */
    kdev_t               i_rdev;           /* real device node */
    u64                  i_version;        /* versioning number */
    loff_t               i_size;           /* file size in bytes */
    seqcount_t           i_size_seqcount;  /* serializer for i_size */
    struct timespec      i_atime;          /* last access time */
    struct timespec      i_mtime;          /* last modify time */
    struct timespec      i_ctime;          /* last change time */
    unsigned int         i_blkbits;        /* block size in bits */
    blkcnt_t             i_blocks;         /* file size in blocks */
    unsigned short       i_bytes;          /* bytes consumed */
    umode_t              i_mode;           /* access permissions */
    spinlock_t           i_lock;           /* spinlock */
    struct rw_semaphore  i_alloc_sem;      /* nests inside of i_sem */
    struct semaphore     i_sem;            /* inode semaphore */
    struct inode_operations *i_op;         /* inode ops table */
    struct file_operations *i_fop;         /* default inode ops */
    struct super_block    *i_sb;           /* associated superblock */
    struct file_lock      *i_flock;        /* file lock list */
    struct address_space  *i_mapping;       /* associated mapping */
    struct address_space  i_data;          /* mapping for device */
    struct dquot          *i_dquot[MAXQUOTAS]; /* disk quotas for inode */
    struct list_head     i_devices;        /* list of block devices */
    union {
        struct pipe_inode_info *i_pipe;    /* pipe information */
        struct block_device     *i_bdev;    /* block device driver */
        struct cdev              *i_cdev;   /* character device driver */
    };
    unsigned long         i_dnotify_mask;  /* directory notify mask */
    struct dnotify_struct *i_dnotify;       /* dnotify */
    struct list_head      inotify_watches; /* inotify watches */
    struct mutex          inotify_mutex;    /* protects inotify_watches */
    unsigned long         i_state;          /* state flags */
    unsigned long         dirtied_when;     /* first dirtying time */
    unsigned int          i_flags;          /* filesystem flags */
    atomic_t             i_writecount;      /* count of writers */
    void                  *i_security;      /* security module */
    void                  *i_private;       /* fs private pointer */
};

```

## Essential members

- `i_list` (the current list)
- `i_sb_list` (all inodes of the same sb)
- `i_dentry`
- `i_ino`
- `i_nlink` (hard link count)
- `i_uid`
- `i_gid`
- `i_count` (ref count)
- `i_size`
- `i_blocks`
- `i_?time`
- `i_mode`
- `*i_op`
- `*i_sb`
- `*i_bdev`

```
};
```

# inode Structure

- Lists
  - i\_list : pointers for the list that describes the inode's current state
  - i\_sb\_list: pointers for the list of inodes of the superblock
- Permissions
  - i\_uid & i\_gid: file user ID and group ID
  - i\_mode : file access mode rwxrwxrwx
- Hard-links
  - i\_nlink : number of hard links
    - When down to 0, the inode (and the file) can be **deleted from disk**
  - i\_dentry : the head of the list of dentry objects referencing this inode
- Time
  - i\_atime, i\_mtime, i\_ctime: access, modify, create times

# inode Structure

- Reference
  - `i_count`: how many processes refer to this inode
  - When down to 0, the inode can be freed from memory
- inode operations
  - `*i_op`: a pointer to a structure of callbacks
- Device
  - `*i_bdev`: a pointer to the underlying block device
- File-system private data
  - `*i_private`: a pointer to FS-specific extension, e.g., direct/indirect/double indirect/triple indirect pointers of ext2 file system

```

struct inode_operations {
    int (*create) (struct inode *,struct dentry *,int, struct nameidata *);
    struct dentry * (*lookup) (struct inode *,struct dentry *, struct nameidata *);
    int (*link) (struct dentry *,struct inode *,struct dentry *);
    int (*unlink) (struct inode *,struct dentry *);
    int (*symlink) (struct inode *,struct dentry *,const char *);
    int (*mkdir) (struct inode *,struct dentry *,int);
    int (*rmdir) (struct inode *,struct dentry *);
    int (*mknod) (struct inode *,struct dentry *,int,dev_t);
    int (*rename) (struct inode *, struct dentry *,
                   struct inode *, struct dentry *);
    int (*readlink) (struct dentry *, char __user *,int);
    void * (*follow_link) (struct dentry *, struct nameidata *);
    void (*put_link) (struct dentry *, struct nameidata *, void *);
    void (*truncate) (struct inode *);
    int (*permission) (struct inode *, int);
    int (*setattr) (struct dentry *, struct iattr *);
    int (*getattr) (struct vfsmount *mnt, struct dentry *, struct kstat *);
    int (*setxattr) (struct dentry *, const char *,const void *,size_t,int);
    ssize_t (*getxattr) (struct dentry *, const char *, void *, size_t);
    ssize_t (*listxattr) (struct dentry *, char *, size_t);
    int (*removexattr) (struct dentry *, const char *);
    void (*truncate_range)(struct inode *, loff_t, loff_t);
    long (*fallocate)(struct inode *inode, int mode, loff_t offset,
                      loff_t len);
    int (*fiemap)(struct inode *, struct fiemap_extents_info *, u64 start,
                  u64 len);
};

```

# inode Operations (1/3)

- inode operations are structural operations, i.e., they may modify disk metadata
- *create*(dir, dentry, mode, nameidata)
  - Creates a new inode for a **regular file** associated with a dentry object in some directory. (new file)
- *lookup*(dir, dentry, nameidata)
  - Searches a directory for an inode corresponding to the filename included in a dentry object. (existing file)
- *link*(old\_dentry, dir, new\_dentry)
  - Creates a new **hard link** that refers to the file specified by old\_dentry in the directory dir; the new hard link has the name specified by new\_dentry.
- *unlink*(dir, dentry)
  - Removes the hard link of the file specified by a dentry object from a directory. (delete the file when hard link=0)

# inode Operations (2/3)

- *mkdir*(dir, dentry, mode)
  - Creates a new inode for **a directory** associated with a dentry object in some directory.
- *rmdir*(dir, dentry)
  - Removes from a directory the subdirectory whose name is included in a dentry object.

# inode Operations (3/3)

- *mknod*(dir, dentry, mode, rdev)
  - Creates a new disk inode for a **special file** associated with a dentry object in some directory.
  - The mode and rdev parameters specify, respectively, the file type and the device's major and minor numbers.
- *rename*(old\_dir, old\_dentry, new\_dir, new\_dentry)
  - Moves the file identified by old\_entry from the old\_dir directory to the new\_dir one.
- *truncate*(inode)
  - Shrink the size of the file associated with an inode.



# ramfs inode Operations

```
static const struct inode_operations ramfs_dir_inode_operations = {  
    .create      = ramfs_create,  
    .lookup      = simple_lookup,  
    .link        = simple_link,  
    .unlink      = simple_unlink,  
    .symlink     = ramfs_symlink,  
    .mkdir       = ramfs_mkdir,  
    .rmdir       = simple_rmdir,  
    .mknod       = ramfs_mknod,  
    .rename      = simple_rename,  
};  
  
const struct inode_operations ramfs_file_inode_operations = {  
    .getattr     = simple_getattr,  
};
```

- On the creation of directory, device node, regular file, ramfs just allocates an inode in memory

```

static int
ramfs_mknod(struct inode *dir, struct dentry *dentry, int mode, dev_t dev)
{
    struct inode *inode = ramfs_get_inode(dir->i_sb, mode, dev);      ← (1)
    int error = -ENOSPC;

    if (inode) {
        if (dir->i_mode & S_ISGID) {                                     ← (2)
            inode->i_gid = dir->i_gid;
            if (S_ISDIR(mode))
                inode->i_mode |= S_ISGID;
        }
        d_instantiate(dentry, inode);
        dget(dentry); /* Extra count - pin the dentry in core */    ← (3)
        error = 0;
        dir->i_mtime = dir->i_ctime = CURRENT_TIME;
    }
    return error;
}

```

*ramfs\_mknod()* creates an inode for the dentry specified in the 2<sup>nd</sup> parameter.

- (1) Calls *ramfs\_get\_inode()* to allocate a new inode
- (2) If the parent directory has a group id, then the inode inherits the group id from the parent
- (3) Increase the reference count of the dentry to prevent it from being deallocated from memory

```

struct inode *ramfs_get_inode(struct super_block *sb, int mode, dev_t dev)
{
    struct inode * inode = new_inode(sb);

    if (inode) {
        inode->i_mode = mode;
        inode->i_uid = current_fsuid();
        inode->i_gid = current_fsgid();
        inode->i_mapping->a_ops = &ramfs_aops; ← (1: address space operations)
        inode->i_mapping->backing_dev_info = &ramfs_backing_dev_info;
        mapping_set_gfp_mask(inode->i_mapping, GFP_HIGHUSER);
        mapping_set_unevictable(inode->i_mapping);
        inode->i_atime = inode->i_mtime = inode->i_ctime = CURRENT_TIME;
        switch (mode & S_IFMT) { ← (2: check inode type)
        default:
            init_special_inode(inode, mode, dev); ← (3a: device node)
            break;
        case S_IFREG:
            inode->i_op = &ramfs_file_inode_operations; ← (3b: regular file)
            inode->i_fop = &ramfs_file_operations;
            break;
        case S_IFDIR:
            inode->i_op = &ramfs_dir_inode_operations; ← (3c: directory file)
            inode->i_fop = &simple_dir_operations;

            /* directory inodes start off with i_nlink == 2 (for "." entry) */
            inc_nlink(inode);
            break;
        case S_IFLNK:
            inode->i_op = &page_symlink_inode_operations; ← (3: symbolic link)
            break;
        }
    } ? end if inode ?
    return inode;
} ? end ramfs_get_inode ?

```

(1) Set address-space operations

(2) S\_IFMT is a bitmak to extract the file type code from the mode value

S\_IFREG: regular file, S\_IFDIR: directory, S\_IFLNK: link

S\_ISCHR, S\_ISBLK, S\_ISFIFO, S\_ISSOCK: device node

(3) Set inode operations and file operations

ramfs\_mknod() is called to create an inode for

- A regular file : ramfs\_create() → ramfs\_mknod(S\_IFREG)
- A directory : ramfs\_mkdir() → ramfs\_mknod(S\_IFDIR)
- A device node : ramfs\_mknod() → ramfs\_mknod(S\_ISCHR, S\_ISBLK, S\_ISFIFO, or S\_ISSOCK)

```
static int ramfs_create(struct inode *dir, struct dentry *dentry, int mode, struct nameidata *nd)
{
    return ramfs_mknod(dir, dentry, mode | S_IFREG, 0); ← (1)
}
```

```
static int ramfs_mkdir(struct inode *dir, struct dentry *dentry, int mode)
{
    int retval = ramfs_mknod(dir, dentry, mode | S_IFDIR, 0); ← (2)
    if (!retval)
        inc_nlink(dir);
    return retval;
}
```

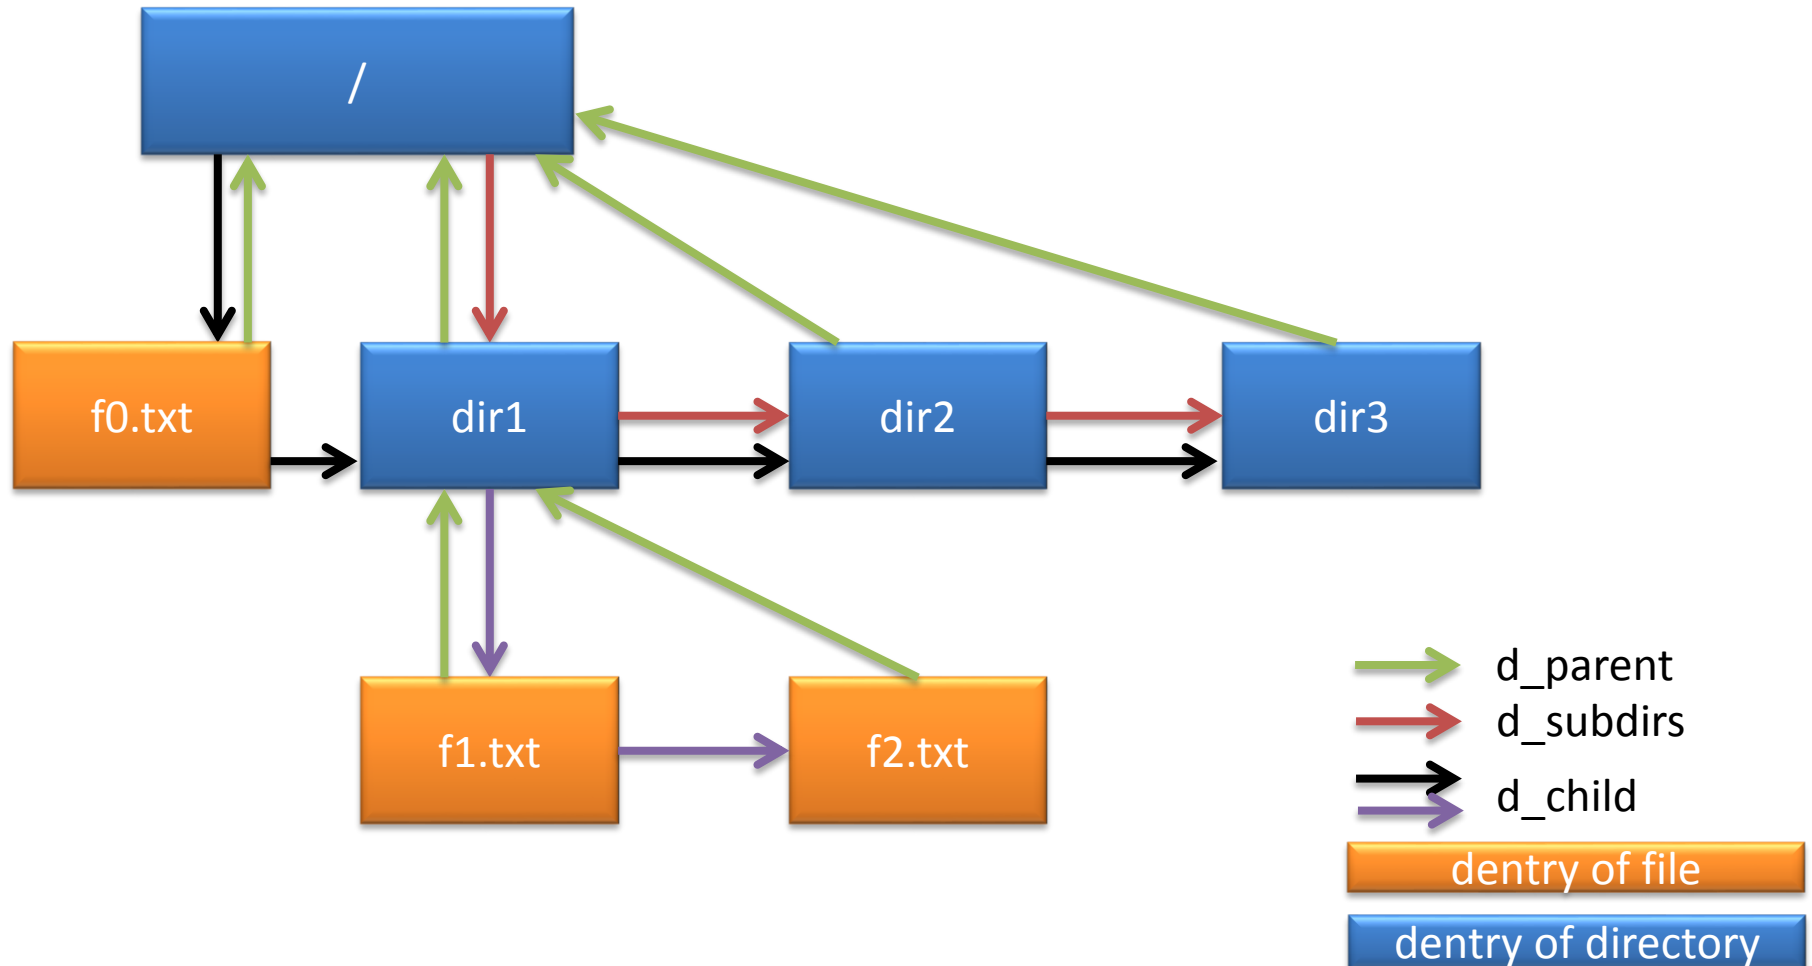
(1) Use S\_IFREG to create an inode for a regular file

(2) Use S\_IFDIR to create an inode for a directory

# Dentry Objects

- Describing the relation among pathname components
  - /mnt/dir1/file1.txt
  - mnt, dir1, file1.txt all have an dentry object
- VFS manages dentry creation/deletion/lookup
  - A FS driver is responsible to filling correct values in dentry objects
- An dentry object has one of the following statuses
  - Free: contains no valid information
  - In use:  $d\_count > 0$ , currently used by FS (ramfs set it  $\geq 1$ )
  - Unused:  $d\_count = 0$ , no processes refer to the corresponding file
  - Negative:  $d\_count < 0$ , the corresponding file has been deleted or the specified file does not exist
- Unused dentry objects are cached so that re-opening a previously opened file can be fast
  - Use an LRU list to recycle unused dentry objects
  - Use a hash function to speed up cache searching

# Relationship Among dentry Objects



```

struct dentry {
    atomic_t          d_count;      /* usage count */
    unsigned int      d_flags;      /* dentry flags */
    spinlock_t        d_lock;      /* per-dentry lock */
    int               d_mounted;    /* is this a mount point? */
    struct inode       *d_inode;    /* associated inode */
    struct hlist_node  d_hash;      /* list of hash table entries */
    struct dentry      *d_parent;   /* dentry object of parent */
    struct qstr        d_name;      /* dentry name */
    struct list_head   d_lru;       /* unused list */
    union {
        struct list_head d_child;  /* list of dentries within */
        struct rcu_head  d_rcu;    /* RCU locking */
    } d_u;
    struct list_head   d_subdirs;   /* subdirectories */
    struct list_head   d_alias;     /* list of alias inodes */
    unsigned long      d_time;      /* revalidate time */
    struct dentry_operations *d_op; /* dentry operations table */
    struct super_block  *d_sb;      /* superblock of file */
    void               *d_fsdata;   /* filesystem-specific data */
    unsigned char       d_iname[DNAME_INLINE_LEN_MIN]; /* short name */
};

```

## Essential members

- d\_count
- \*d\_inode
- \*d\_parent
- d\_name (filename)
- d\_lru
- d\_child
- d\_subdirs
- \*d\_sb

```

struct dentry_operations {
    int (*d_revalidate) (struct dentry *, struct nameidata *);
    int (*d_hash) (struct dentry *, struct qstr *);
    int (*d_compare) (struct dentry *, struct qstr *, struct qstr *);
    int (*d_delete) (struct dentry *);
    void (*d_release) (struct dentry *);
    void (*d_iput) (struct dentry *, struct inode *);
    char *(*d_dname) (struct dentry *, char *, int);
};

```

```
};
```

# File Objects

- A file object is the representation of an opened file
  - A process's view of an opened file
    - Such as the current file position
    - If many processes open the same file, then there will be many file objects per process and only one dentry and one inode
  - Created upon `open()` and destroyed upon `close()`
  - Does not correspond to any disk structures
- Operations on file objects are related to *user data access*, such as read and write
  - Notice that operations on inodes are related *disk data structure (metadata) management*



```

struct file {
    union {
        struct list_head    fu_list;        /* list of file objects */
        struct rcu_head      fu_rcuhead;     /* RCU list after freeing */
    } f_u;
    struct path              f_path;         /* contains the dentry */
    struct file_operations *f_op;           /* file operations table */
    spinlock_t               f_lock;         /* per-file struct lock */
    atomic_t                 f_count;        /* file object's usage count */
    unsigned int             f_flags;        /* flags specified on open */
    mode_t                   f_mode;         /* file access mode */
    loff_t                   f_pos;          /* file offset (file pointer) */
    struct fown_struct        f_owner;       /* owner data for signals */
    const struct cred         *f_cred;       /* file credentials */
    struct file_ra_state     f_ra;          /* read-ahead state */
    u64                      f_version;     /* version number */
    void                     *f_security;    /* security module */
    void                     *private_data; /* tty driver hook */
    struct list_head         f_ep_links;     /* list of epoll links */
    spinlock_t               f_ep_lock;     /* epoll lock */
    struct address_space     *f_mapping;     /* page cache mapping */
    unsigned long            f_mnt_write_state; /* debugging state */
};

```

```

struct path {
    struct vfsmount *mnt;
    struct dentry *dentry;
};

```

Essential members:

- f\_path
  - points to dentry
- f\_count
- f\_pos
- f\_ops
- f\_mode
  - How the process opened this file for access, e.g., read, write, or exec (similar to i\_mode)
- f\_flags
  - Process open flags, such as O\_SYNC, O\_DIRECT...

```

struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t (*aio_read) (struct kiocb *, const struct iovec *,
                        unsigned long, loff_t);
    ssize_t (*aio_write) (struct kiocb *, const struct iovec *,
                        unsigned long, loff_t);
    int (*readdir) (struct file *, void *, filldir_t);
    unsigned int (*poll) (struct file *, struct poll_table_struct *);
    int (*ioctl) (struct inode *, struct file *, unsigned int,
                unsigned long);
    long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
    long (*compat_ioctl) (struct file *, unsigned int, unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *, fl_owner_t id);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, struct dentry *, int datasync);
    int (*aio_fsync) (struct kiocb *, int datasync);
    int (*fasync) (int, struct file *, int);
    int (*lock) (struct file *, int, struct file_lock *);
    ssize_t (*sendpage) (struct file *, struct page *,
                        int, size_t, loff_t *, int);
    unsigned long (*get_unmapped_area) (struct file *,
                                        unsigned long,
                                        unsigned long,
                                        unsigned long,
                                        unsigned long);

    int (*check_flags) (int);
    int (*flock) (struct file *, int, struct file_lock *);
    ssize_t (*splice_write) (struct pipe_inode_info *,
                            struct file *,
                            loff_t *,
                            size_t,
                            unsigned int);
    ssize_t (*splice_read) (struct file *,
                            loff_t *,
                            struct pipe_inode_info *,
                            size_t,
                            unsigned int);
    int (*setlease) (struct file *, long, struct file_lock **);
};

```

# Operations on File Objects

- *lseek*(file, offset, origin)
  - Updates the file pointer
- *read*(file, buf, count, offset)
  - Reads count bytes from a file starting at position\*offset; the value \*offset is then increased
- *aio\_read*(req, buf, len, pos)
  - Starts an asynchronous I/O operation to read *len* bytes into *buf* from file position *pos*
- *write*(file, buf, count, offset)
  - Writes count bytes into a file starting at position\*offset; the value \*offset is then increased
- *aio\_write*(req, buf, len, pos)
  - Starts an asynchronous I/O operation to write *len* bytes from *buf* to file position *pos*

# Operations on File Objects

- *ioctl*(inode, file, cmd, arg)
  - Sends a command to an underlying hardware device. This method applies only to device files
- *mmap*(file, vma)
  - Performs a memory mapping of the file into a process address space
- *open*(inode, file)
  - Opens a file by creating a new file object and linking it to the corresponding inode object
- *release*(inode, file)
  - Releases the file object. Called when the last reference to an open file is closed— that is, when the `f_count` field of the file object becomes 0
- *fsync*(file, dentry, flag)
  - Flushes the file by writing all cached data to disk.

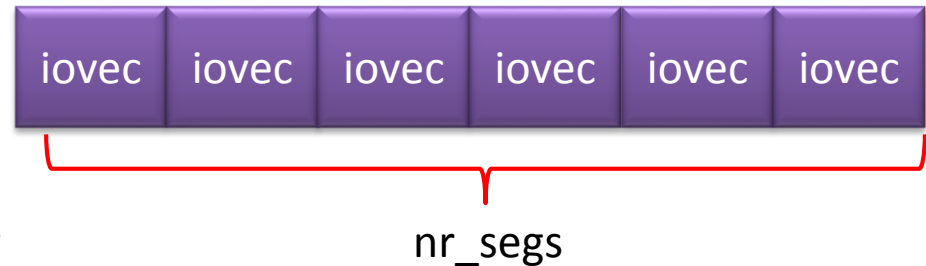
# ramfs File Operations

```
const struct file_operations ramfs_file_operations = {  
    .read           = do_sync_read,  
    .aio_read       = generic_file_aio_read,  
    .write          = do_sync_write,  
    .aio_write      = generic_file_aio_write,  
    .mmap           = generic_file_mmap,  
    .fsync          = simple_sync_file,  
    .splice_read    = generic_file_splice_read,  
    .splice_write   = generic_file_splice_write,  
    .llseek         = generic_file_llseek,  
};
```

- ramfs uses only the default handlers for file operations
  - These default handlers focus on the copy between user pages and kernel pages
  - How data are copied between kernel pages and disk blocks are implemented in *address-space operations*

# iovec

```
struct iovec
{
    void __user *iov_base;
    __kernel_size_t iov_len;
};
```



`iov_base`: a pointer to user-land buffer

`iov_len`: I/O length in bytes

- `*iovec`, an vector of file system I/O operations
  - A pointer to an array of `iovec`
  - Array has `nr_segs` elements (`iovec`)
  - The buffer pointed by `iov_base` is accessible in file system drivers
    - No further `kmap` is required

```

ssize_t
generic_file_aio_read(struct kiocb *iocb, const struct iovec *iov,
                     unsigned long nr_segs, loff_t pos)
{
    struct file *filp = iocb->ki_filp;
    ssize_t retval;
    unsigned long seg;
    size_t count;
    loff_t *ppos = &iocb->ki_pos;

    .....

    for (seg = 0; seg < nr_segs; seg++) {
        read_descriptor_t desc;

        desc.written = 0;
        desc.arg.buf = iov[seg].iov_base;
        desc.count = iov[seg].iov_len;
        if (desc.count == 0)
            continue;
        desc.error = 0;
        do_generic_file_read(filp, ppos, &desc, file_read_actor);
        retval += desc.written;
        if (desc.error) {
            retval = retval ? desc.error;
            break;
        }
        if (desc.count > 0)
            break;
    }
out:
    return retval;
} ? end generic_file_aio_read ?

```

```

/**
 * generic_file_aio_read - generic filesystem read routine
 * @iocb:      kernel I/O control block
 * @iov:      io vector request
 * @nr_segs:  number of segments in the iovec
 * @pos:      current file position
 *
 * This is the "read()" routine for all filesystems
 * that can use the page cache directly.
 */

```

For each iovec in the iov[], convert the iovec into a “read\_descriptor\_t” and calls do\_generic\_file\_read()

\*\*notice that do\_sync\_read() is only a wrapper of generic\_file\_aio\_read()

- Wait on each aio to complete

```

static void do_generic_file_read(struct file *filp, loff_t *ppos,
                                read_descriptor_t *desc, read_actor_t actor)
{
    // ...
    for (;;) {
        // ...
        cond_resched();
find_page:
        page = find_get_page(mapping, index);
        if (!page) {
            page_cache_sync_readahead(mapping,
                                     ra, filp,
                                     index, last_index - index);
            page = find_get_page(mapping, index);
            if (unlikely(page == NULL))
                goto ↓no_cached_page;
        }
page_ok:
        // ...
        ret = actor(desc, page, offset, nr);
        offset += ret;
        index += offset >> PAGE_CACHE_SHIFT;
        offset &= ~PAGE_CACHE_MASK;
        prev_offset = offset;

        page_cache_release(page);
        if (ret == nr && desc->count)
            continue;
        goto ↓out;
        // ...
readpage:
        // ...
        error = mapping->a_ops->readpage(filp, page);
        // ...
        goto ↑page_ok;
        // ...
no_cached_page:
        page = page_cache_alloc_cold(mapping);
        if (!page) {
            desc->error = -ENOMEM;
            goto ↓out;
        }
        error = add_to_page_cache_lru(page, mapping,
                                     index, GFP_KERNEL);
        // ...
        goto ↑readpage;
    } ? end for ;; ?

out:
    // ...
    file_accessed(filp);
} ? end do_generic_file_read ?

```

Check whether the desired page is already in page cache?

actor() is a pointer to file\_read\_actor(), which copies data from kernel pages to user buffer

Copy data to the new page. It will involve disk I/O in conventional file systems

Allocate a free page

Add the page to page cache



```

int file_read_actor(read_descriptor_t *desc, struct page *page,
                    unsigned long offset, unsigned long size)
{
    char *kaddr;
    unsigned long left, count = desc->count;

    if (size > count)
        size = count;

    /*
     * Faults on the destination of a read are common, so do it before
     * taking the kmap.
     */
    if (!fault_in_pages_writeable(desc->arg.buf, size)) {
        kaddr = kmap_atomic(page, KM_USER0);
        left = __copy_to_user_inatomic(desc->arg.buf,
                                       kaddr + offset, size);
        kunmap_atomic(kaddr, KM_USER0);
        if (left == 0)
            goto ↓success;
    }

    /* Do it the slow way */
    kaddr = kmap(page);
    left = __copy_to_user(desc->arg.buf, kaddr + offset, size);
    kunmap(page);

    if (left) {
        size -= left;
        desc->error = -EFAULT;
    }
    success:
    desc->count = count - size;
    desc->written += size;
    desc->arg.buf += size;
    return size;
} ? end file_read_actor ?

```

“fault in” the user buffer. If no page fault, do quick mapping for the kernel page and then copy data (hmm... everything is quick)

Already wasting a lot of time to fault in the user buffer. There is no need to use quick mapping. Do slow mapping for the kernel page and then copy data.

# ramfs Write

- `do_sync_write()` is a wrapper of `generic_file_aio_write()`
  - Wait on each aio write to complete
  - Descending into `generic_perform_write()`, which does the following for each page
    - Calls `a_op→begin_write()` to prepare the next page to write
    - Copy data from user buffers pointed by `iovec` to the page just prepared
    - Calls `a_op→end_write()` to mark the page dirty

# Address Space

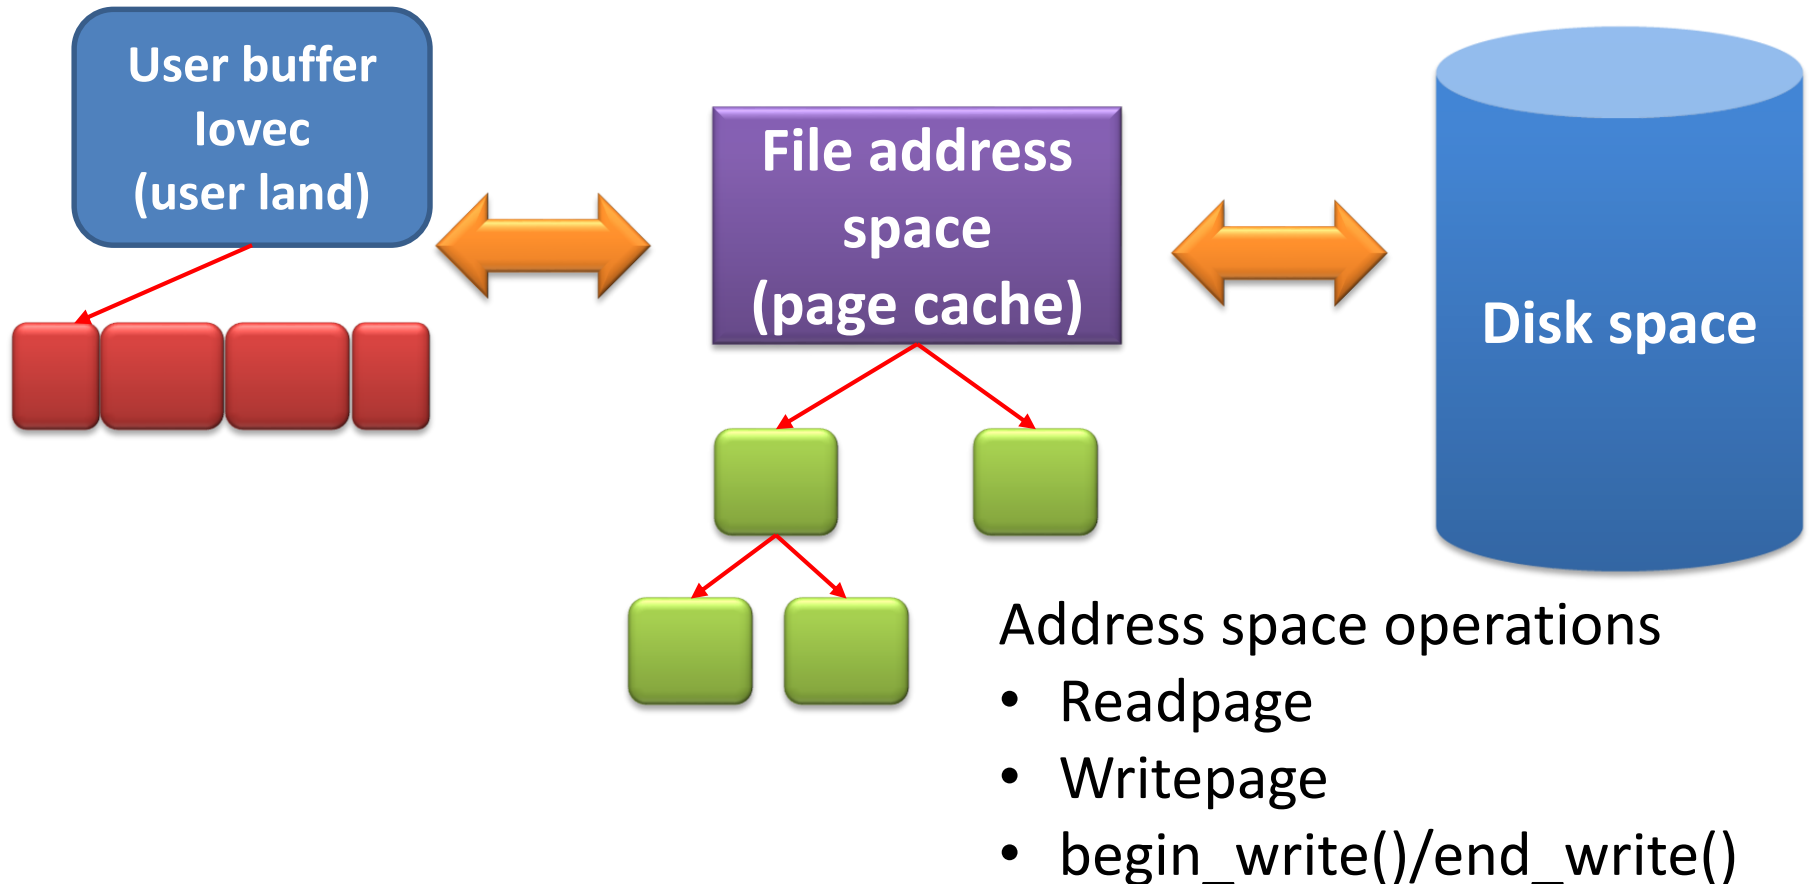
- “file address space” is more appropriate here
  - Not to be confused with process address space
  - A page owned by a process and is not mapped to any disk file is called an anonymous page
    - Pages of heap and stack
  - A page owned by a file is mapped to the file
    - Governed by the address space object “i\_data” embedded in the file’s inode
- One of the good reasons to adopt address space is to enable caching for file systems not using block devices
  - Cached pages are indexed by (inode#, page offset), not (dev #, block offset)
  - ramfs, nfs, yaffs2, ubifs, etc.

# Address Space

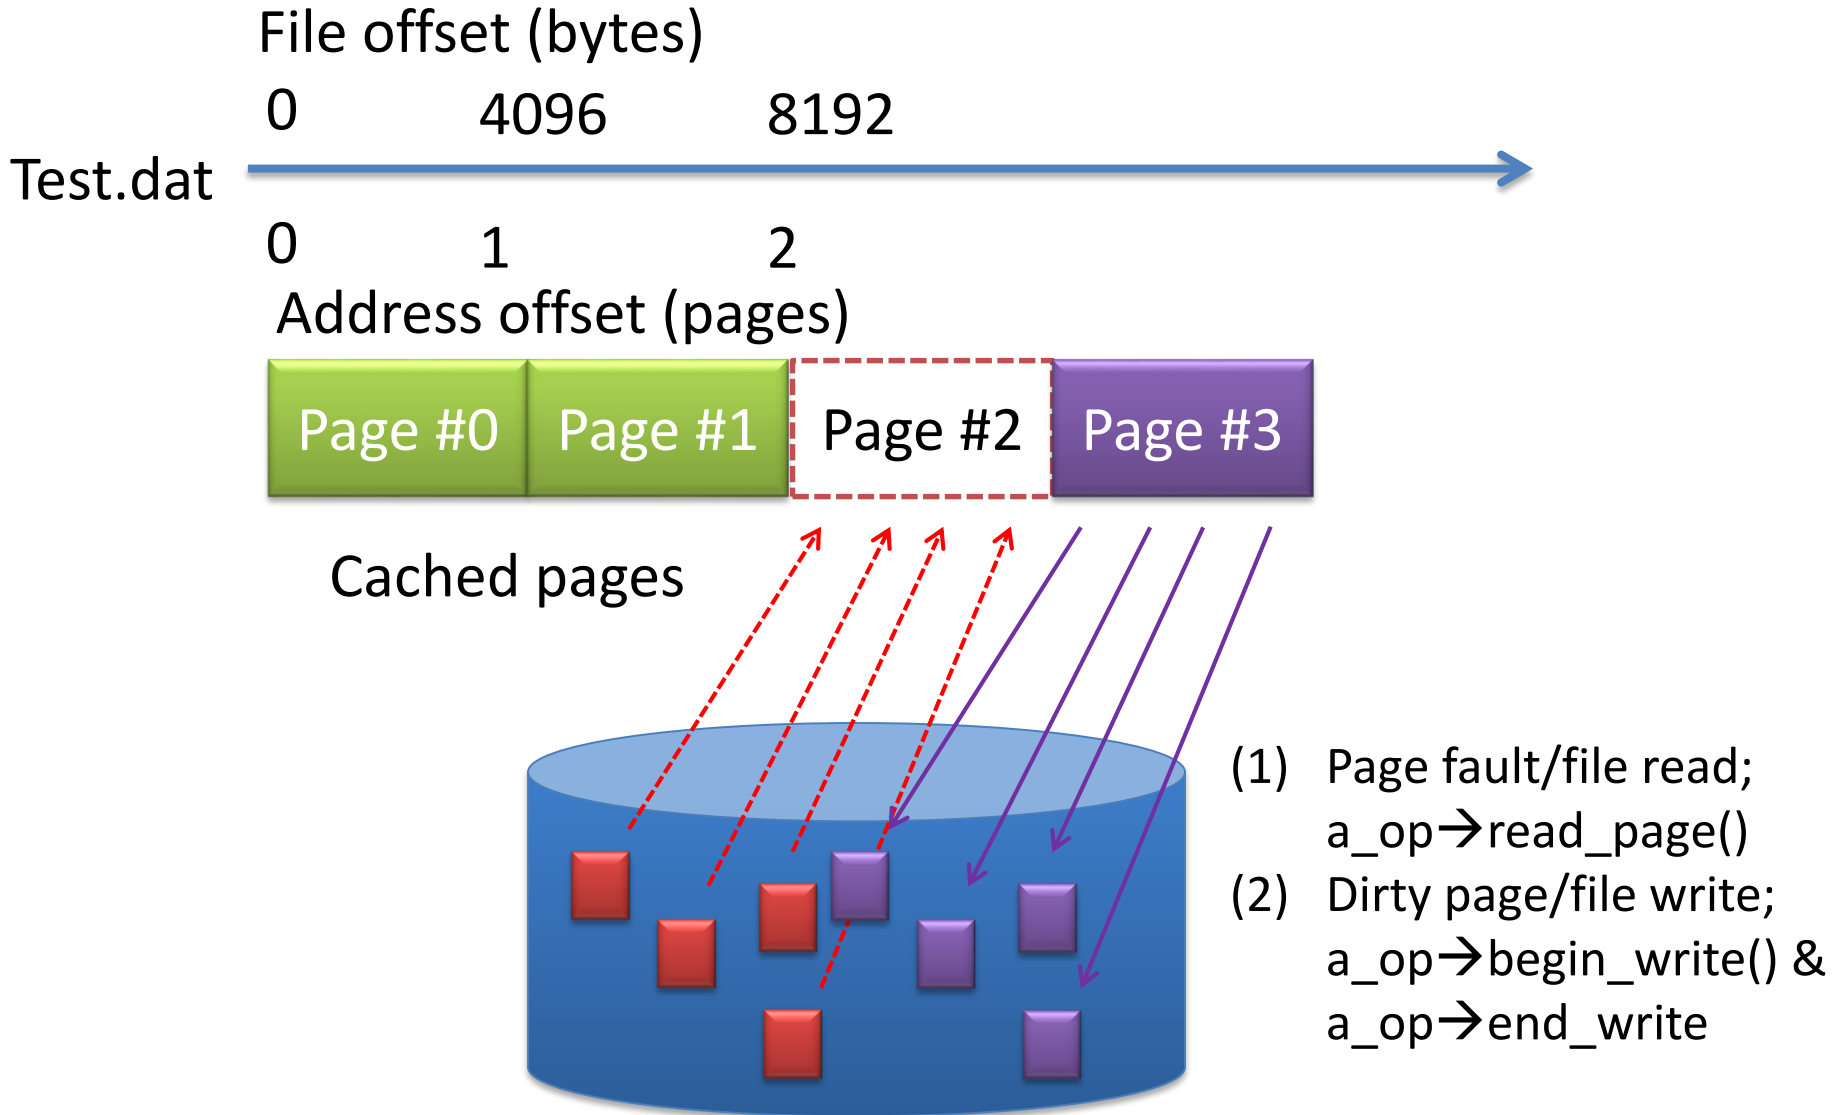
- A generic template of file operation
  - Modify VFS objects for the file operation
  - Calls address space operations
    - map pages to disk blocks, may involve metadata modification
    - Perform data copy
  - Mark modified pages dirty
- Thus, address space operations are to “map” and “copy” between pages and disk blocks

# User buffer, address space, and disk space

`generic_file_aio_read()`  
`generic_file_aio_write()`



# Address Space



# begin\_write() and end\_write()

- Part of file system implementation
- On file write, conventional disk file system will
  - Calls a\_op → begin\_write()
    - Modify disk metadata to allocate disk space for new data
    - Prepare VM pages for new data
  - Copy data from user buffer to VM pages
  - Calls a\_op → end\_write() to mark the touched pages dirty

```

struct address_space {
    struct inode          *host;                /* owning inode */
    struct radix_tree_root page_tree;           /* radix tree of all pages */
    spinlock_t           tree_lock;            /* page_tree lock */
    unsigned int          i_mmap_writable;      /* VM_SHARED ma count */
    struct prio_tree_root i_mmap;              /* list of all mappings */
    struct list_head      i_mmap_nonlinear;    /* VM_NONLINEAR ma list */
    spinlock_t           i_mmap_lock;         /* i_mmap lock */
    atomic_t             truncate_count;       /* truncate re count */
    unsigned long         nrpages;             /* total number of pages */
    pgoff_t              writeback_index;      /* writeback start offset */
    struct address_space_operations *a_ops;    /* operations table */
    unsigned long         flags;               /* gfp_mask and error flags */
    struct backing_dev_info *backing_dev_info; /* read-ahead information */
    spinlock_t           private_lock;         /* private lock */
    struct list_head      private_list;        /* private list */
    struct address_space  *assoc_mapping;      /* associated buffers */
};

```

Essential members:

- \*host
- page\_tree (like the radix tree in ramdisk)
- \*a\_ops



```

struct address_space_operations {
    int (*writepage)(struct page *, struct writeback_control *);
    int (*readpage)(struct file *, struct page *);
    int (*sync_page)(struct page *);
    int (*writepages)(struct address_space *,
                      struct writeback_control *);
    int (*set_page_dirty)(struct page *);
    int (*readpages)(struct file *, struct address_space *,
                     struct list_head *, unsigned);
    int (*write_begin)(struct file *, struct address_space *mapping,
                       loff_t pos, unsigned len, unsigned flags,
                       struct page **pagep, void **fsdata);
    int (*write_end)(struct file *, struct address_space *mapping,
                     loff_t pos, unsigned len, unsigned copied,
                     struct page *page, void *fsdata);
    sector_t (*bmap)(struct address_space *, sector_t);
    int (*invalidatepage)(struct page *, unsigned long);
    int (*releasepage)(struct page *, int);
    int (*direct_IO)(int, struct kiocb *, const struct iovec *,
                     loff_t, unsigned long);
    int (*get_xip_mem)(struct address_space *, pgoff_t, int,
                       void **, unsigned long *);
    int (*migratepage)(struct address_space *,
                       struct page *, struct page *);
    int (*launder_page)(struct page *);
    int (*is_partially_uptodate)(struct page *,
                                read_descriptor_t *,
                                unsigned long);
    int (*error_remove_page)(struct address_space *,
                             struct page *);
};

```

# Address Space Operations

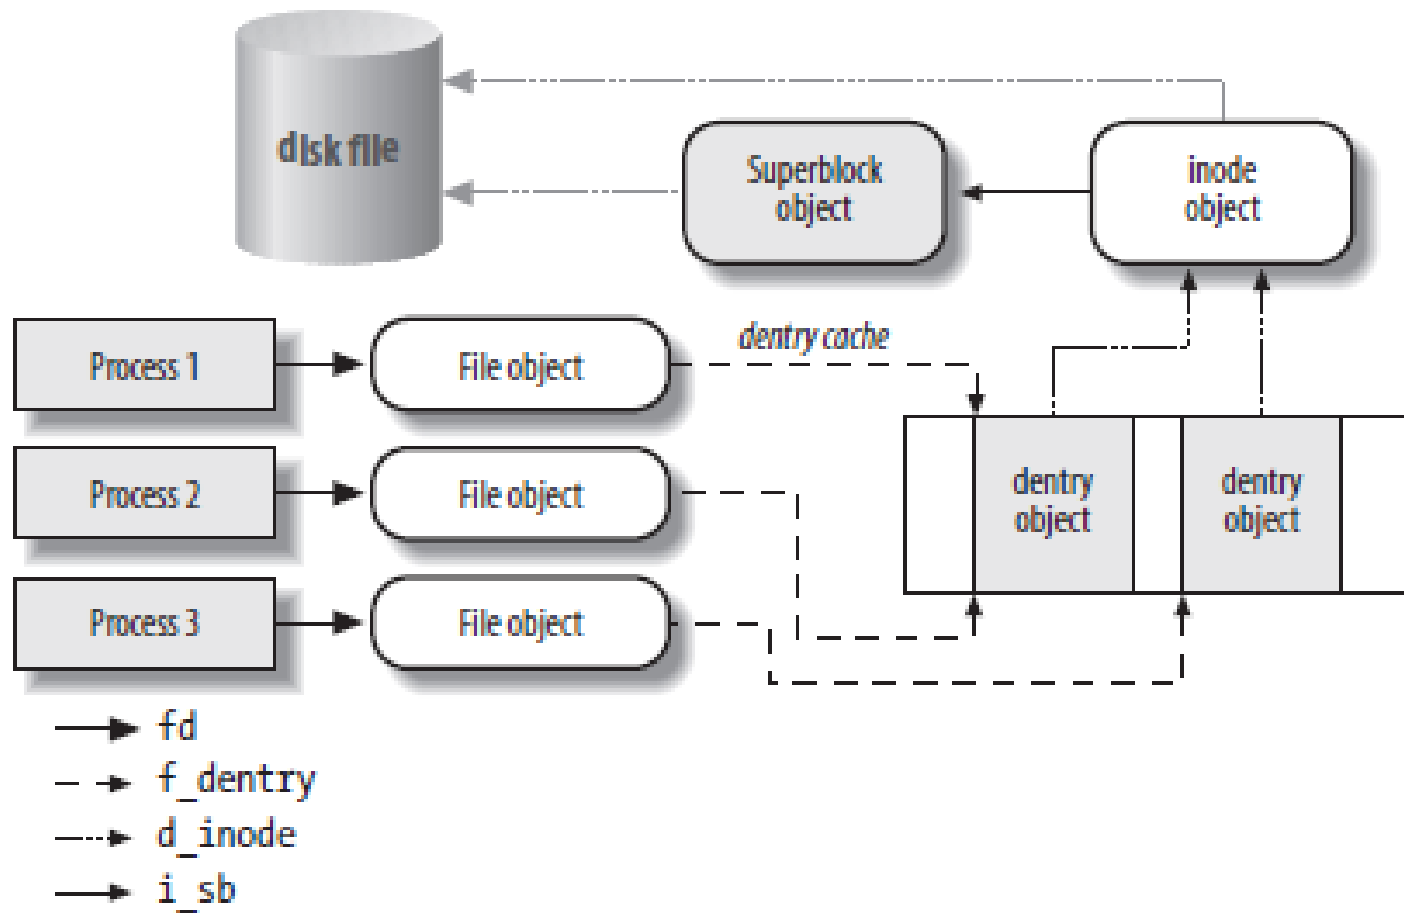
- `writepage()`
  - Write operation (from the page to the owner's disk image)
- `readpage()`
  - Read operation (from the owner's disk image to the page)
- `prepare_write()` or `write_begin()`
  - Prepare a write operation
  - Usually involves metadata modification for disk-based file systems
- `commit_write()` or `write_end()`
  - Complete a write operation
  - Mark the modified pages dirty

# ramfs Address Space Operations

```
const struct address_space_operations ramfs_aops = {  
    .readpage      = simple_readpage,  
    .write_begin   = simple_write_begin,  
    .write_end     = simple_write_end,  
    .set_page_dirty = __set_page_dirty_no_writeback,  
};  
  
int simple_readpage(struct file *file, struct page *page)  
{  
    clear_highpage(page);  
    flush_dcache_page(page);  
    SetPageUptodate(page);  
    unlock_page(page);  
    return 0;  
}
```

- ramfs does nothing useful in address space operations because there is no need to copy/map data between cached pages and disk blocks

# Review



# Lab 10: Simple File Encryption

# File Encryption

- Implement XOR-based file encryption in ramfs
- Files are encrypted and decrypted by XOR'ing file bytes with a key
- If encryption is on
  - XOR bytes after read
  - XOR bytes before write
- By turning off encryption, you get garbage when reading an encrypted file

# /proc file system

- An alternative way to communicate with kernel code
  - File system drivers are not “devices”, and they cannot accept ioctl as ramdisks do
- Register an /proc file entry for user-land access, as well as the entry’s read/write handlers
  - The handlers will be called when user programs access the registered /proc entry
- Use this feature to implement turning on/off encryption
  - `cat 1 > /proc/ramfs_flag`

# ramfs Modification

- You need to register a /proc file so that you can turn on/off encryption in user space
- You need to register your own file operation callbacks to process encryption/decryption



# References

- Robert love, “Linux Kernel Development 3<sup>rd</sup> Edition,” 2010
- Daniel B. Bovet et al., “Understanding the Linux Kernel 3<sup>rd</sup> Edition,” 2005
- Linux kernel source tree 2.6.34